




Llama hex L31: power per Laplacian mode



cumulative variance vs #modes




power: eigenmode x layer




Llama hex: Laplacian eigenmode DIVIDERS (% = variance share at L31)


mode 1 $\lambda=0.78$ (26%)




mode 2 $\lambda=1.02$ (17%)




mode 3 $\lambda=1.46$ (7%)




mode 4 $\lambda=2.69$ (6%)




mode 5 $\lambda=2.80$ (4%)




mode 6 $\lambda=3.40$ (3%)




mode 7 $\lambda=4.01$ (3%)




mode 8 $\lambda=4.73$ (3%)




mode 9 $\lambda=4.81$ (4%)




mode 10 $\lambda=5.30$ (4%)




mode 11 $\lambda=5.90$ (4%)




mode 12 $\lambda=6.42$ (4%)




mode 13 $\lambda=6.99$ (4%)



mode 14 $\lambda=7.77$ (5%)




mode 15 $\lambda=7.91$ (5%)




Llama hex: named-cut basis (left) & model's own top axes (right)


named cuts, greedy orthogonal (cumulative Σ)




PC0 (30%)
 $\approx$ x (coord 0)



PC1 (15%)
 $\approx$ y (coord 1)



PC2 (10%)
 $\approx$ centre|edge



PC3 (8%)
 $\approx$ centre|edge

