




Llama_r2 random L31: power per Laplacian mode



cumulative variance vs #modes




power: eigenmode x layer




Llama_r2 random: Laplacian eigenmode DIVIDERS (% = variance share at L31)


mode 1 $\lambda=1.30$ (14%)




mode 2 $\lambda=1.53$ (13%)




mode 3 $\lambda=2.13$ (7%)




mode 4 $\lambda=2.45$ (4%)




mode 5 $\lambda=2.83$ (4%)




mode 6 $\lambda=3.34$ (4%)




mode 7 $\lambda=3.56$ (4%)




mode 8 $\lambda=3.80$ (3%)




mode 9 $\lambda=4.40$ (4%)




mode 10 $\lambda=4.71$ (5%)




mode 11 $\lambda=5.87$ (8%)




mode 12 $\lambda=6.05$ (5%)




mode 13 $\lambda=6.63$ (10%)



mode 14 $\lambda=7.64$ (6%)



mode 15 $\lambda=7.76$ (8%)



Llama_r2 random: named-cut basis (left) & model's own top axes (right)

named cuts, greedy orthogonal (cumulative Σ)

