




Llama ring L31: power per Laplacian mode



cumulative variance vs #modes




power: eigenmode × layer




Llama ring: Laplacian eigenmode DIVIDERS (% = variance share at L31)


mode 1 $\lambda=0.15$ (21%)




mode 2 $\lambda=0.15$ (27%)




mode 3 $\lambda=0.59$ (7%)




mode 4 $\lambda=0.59$ (6%)




mode 5 $\lambda=1.23$ (2%)




mode 6 $\lambda=1.23$ (2%)




mode 7 $\lambda=2.00$ (1%)




mode 8 $\lambda=2.00$ (1%)




mode 9 $\lambda=2.77$ (2%)




mode 10 $\lambda=2.77$ (2%)




mode 11 $\lambda=3.41$ (5%)




mode 12 $\lambda=3.41$ (4%)




mode 13 $\lambda=3.85$ (8%)



mode 14 $\lambda=3.85$ (5%)




mode 15 $\lambda=4.00$ (5%)




Llama ring: named-cut basis (left) & model's own top axes (right)


named cuts, greedy orthogonal (cumulative Σ)




PC0 (29%)
 $\approx$ y (coord 1)



PC1 (21%)
 $\approx$ x (coord 0)



PC2 (12%)
 $\approx$ checkerboard (parity)



PC3 (8%)
 $\approx$ checkerboard (parity)

