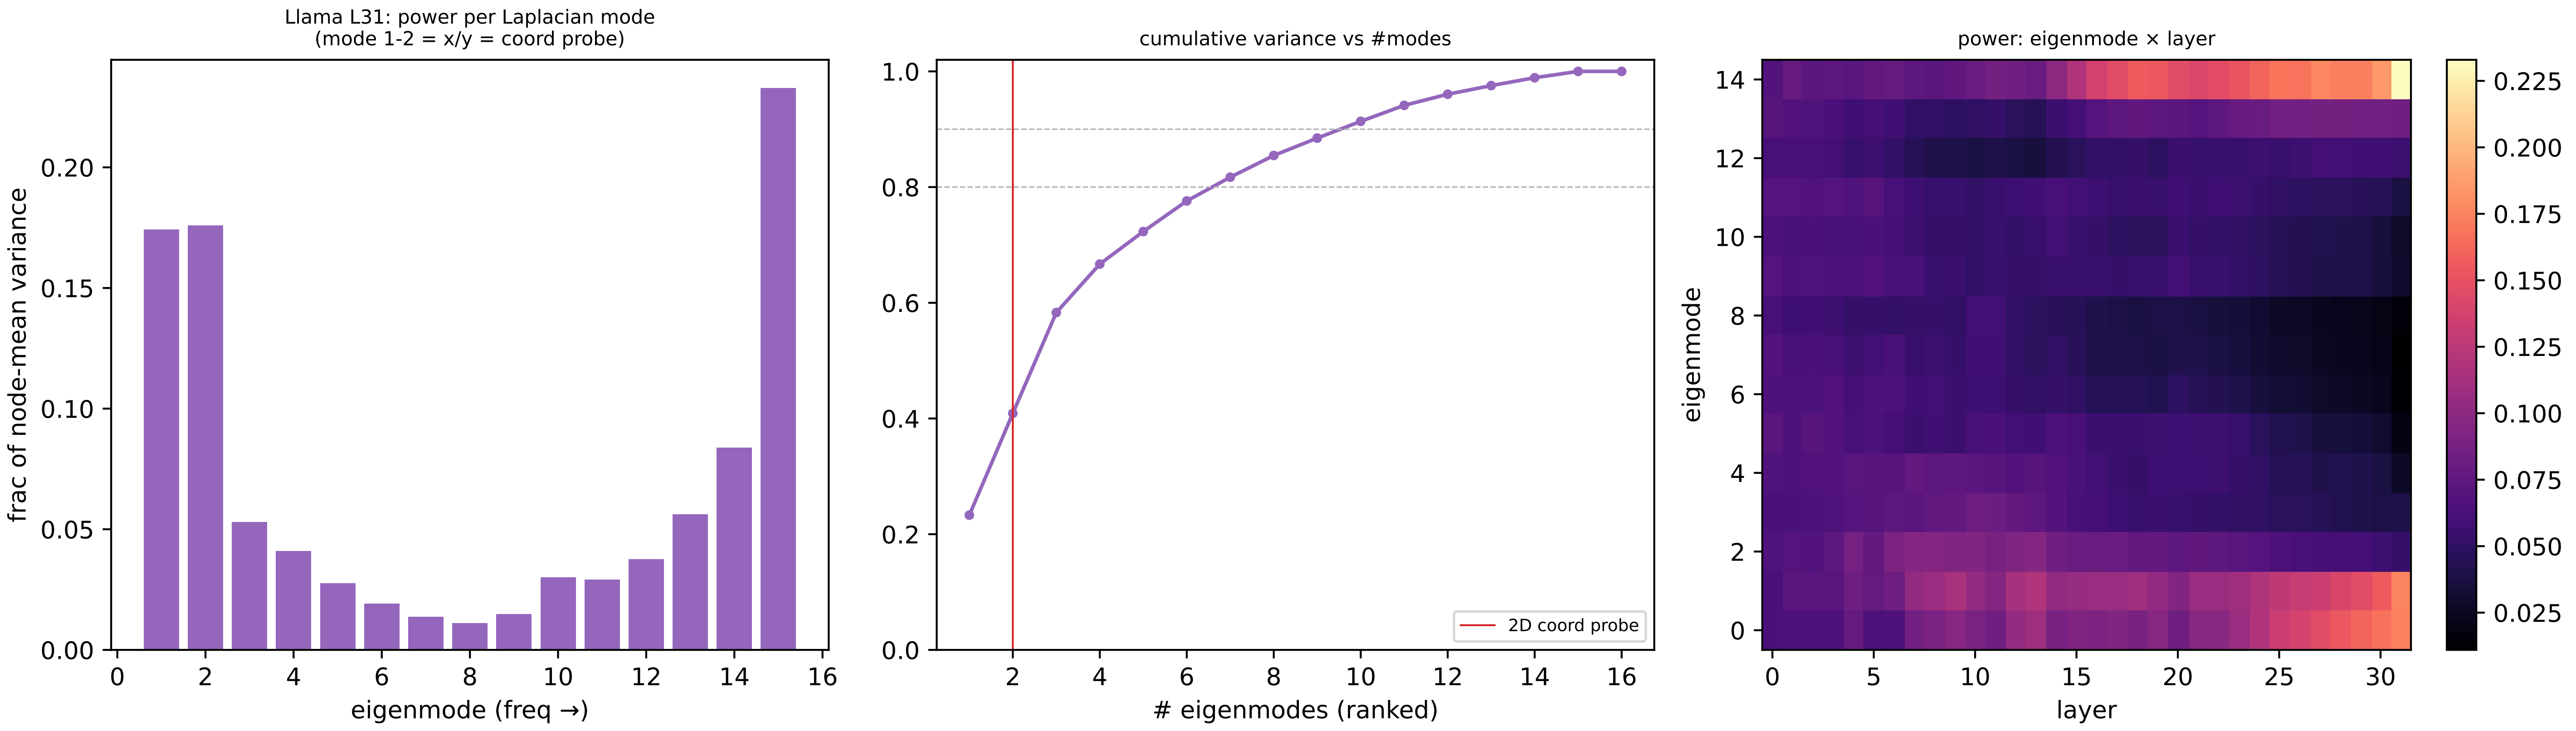

Llama L31: power per Laplacian mode
(mode 1-2 = x/y = coord probe)
frac of node-mean variance
eigenmode (freq →)
cumulative variance vs #modes
2D coord probe
eigenmodes (ranked)
power: eigenmode × layer
eigenmode
layer

Llama: Laplacian eigenmode DIVIDERS (blue|red = the two sides of the cut)
% = share of node-mean variance at L31

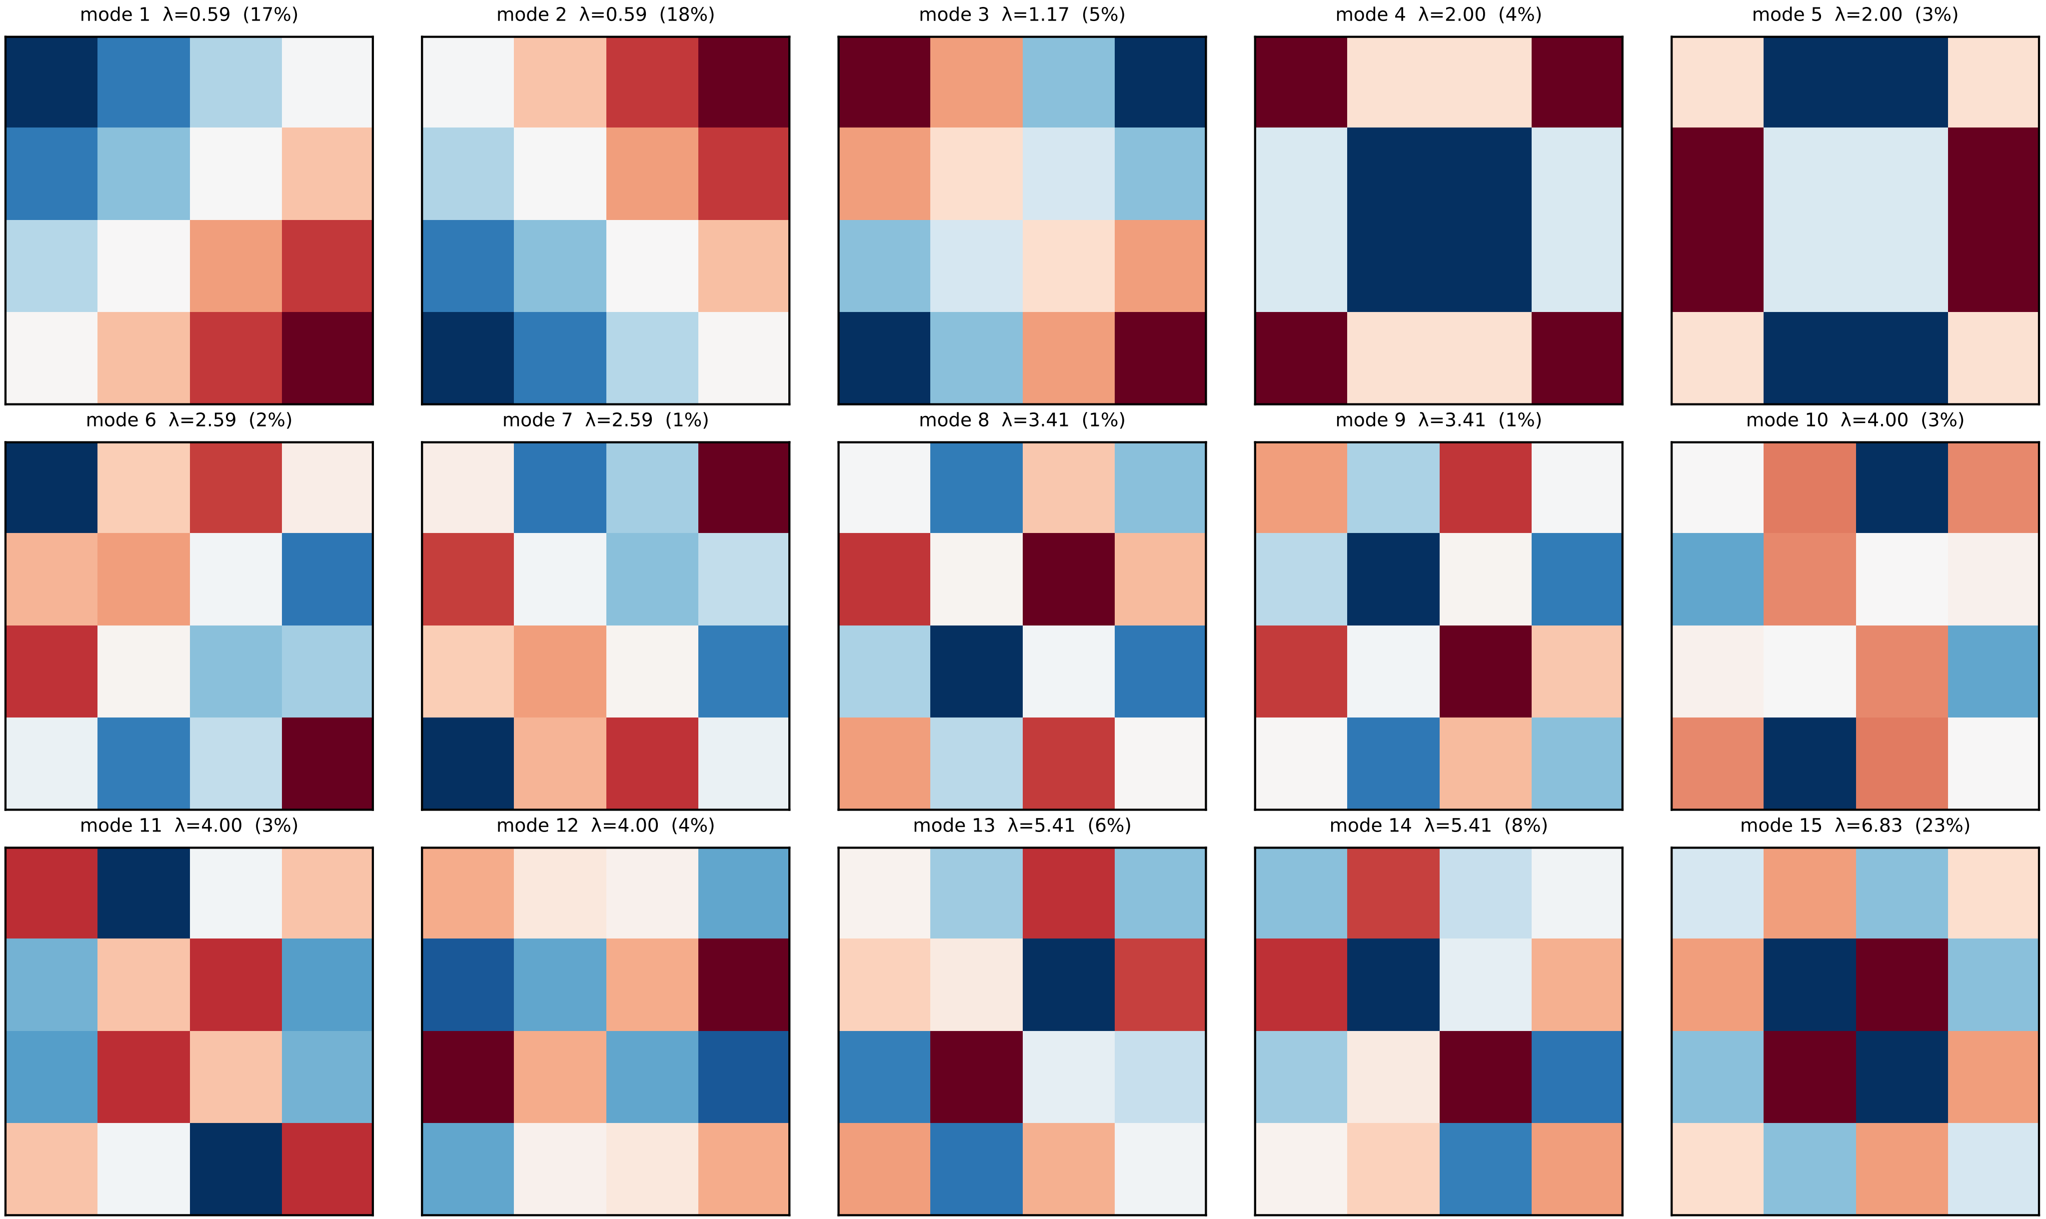

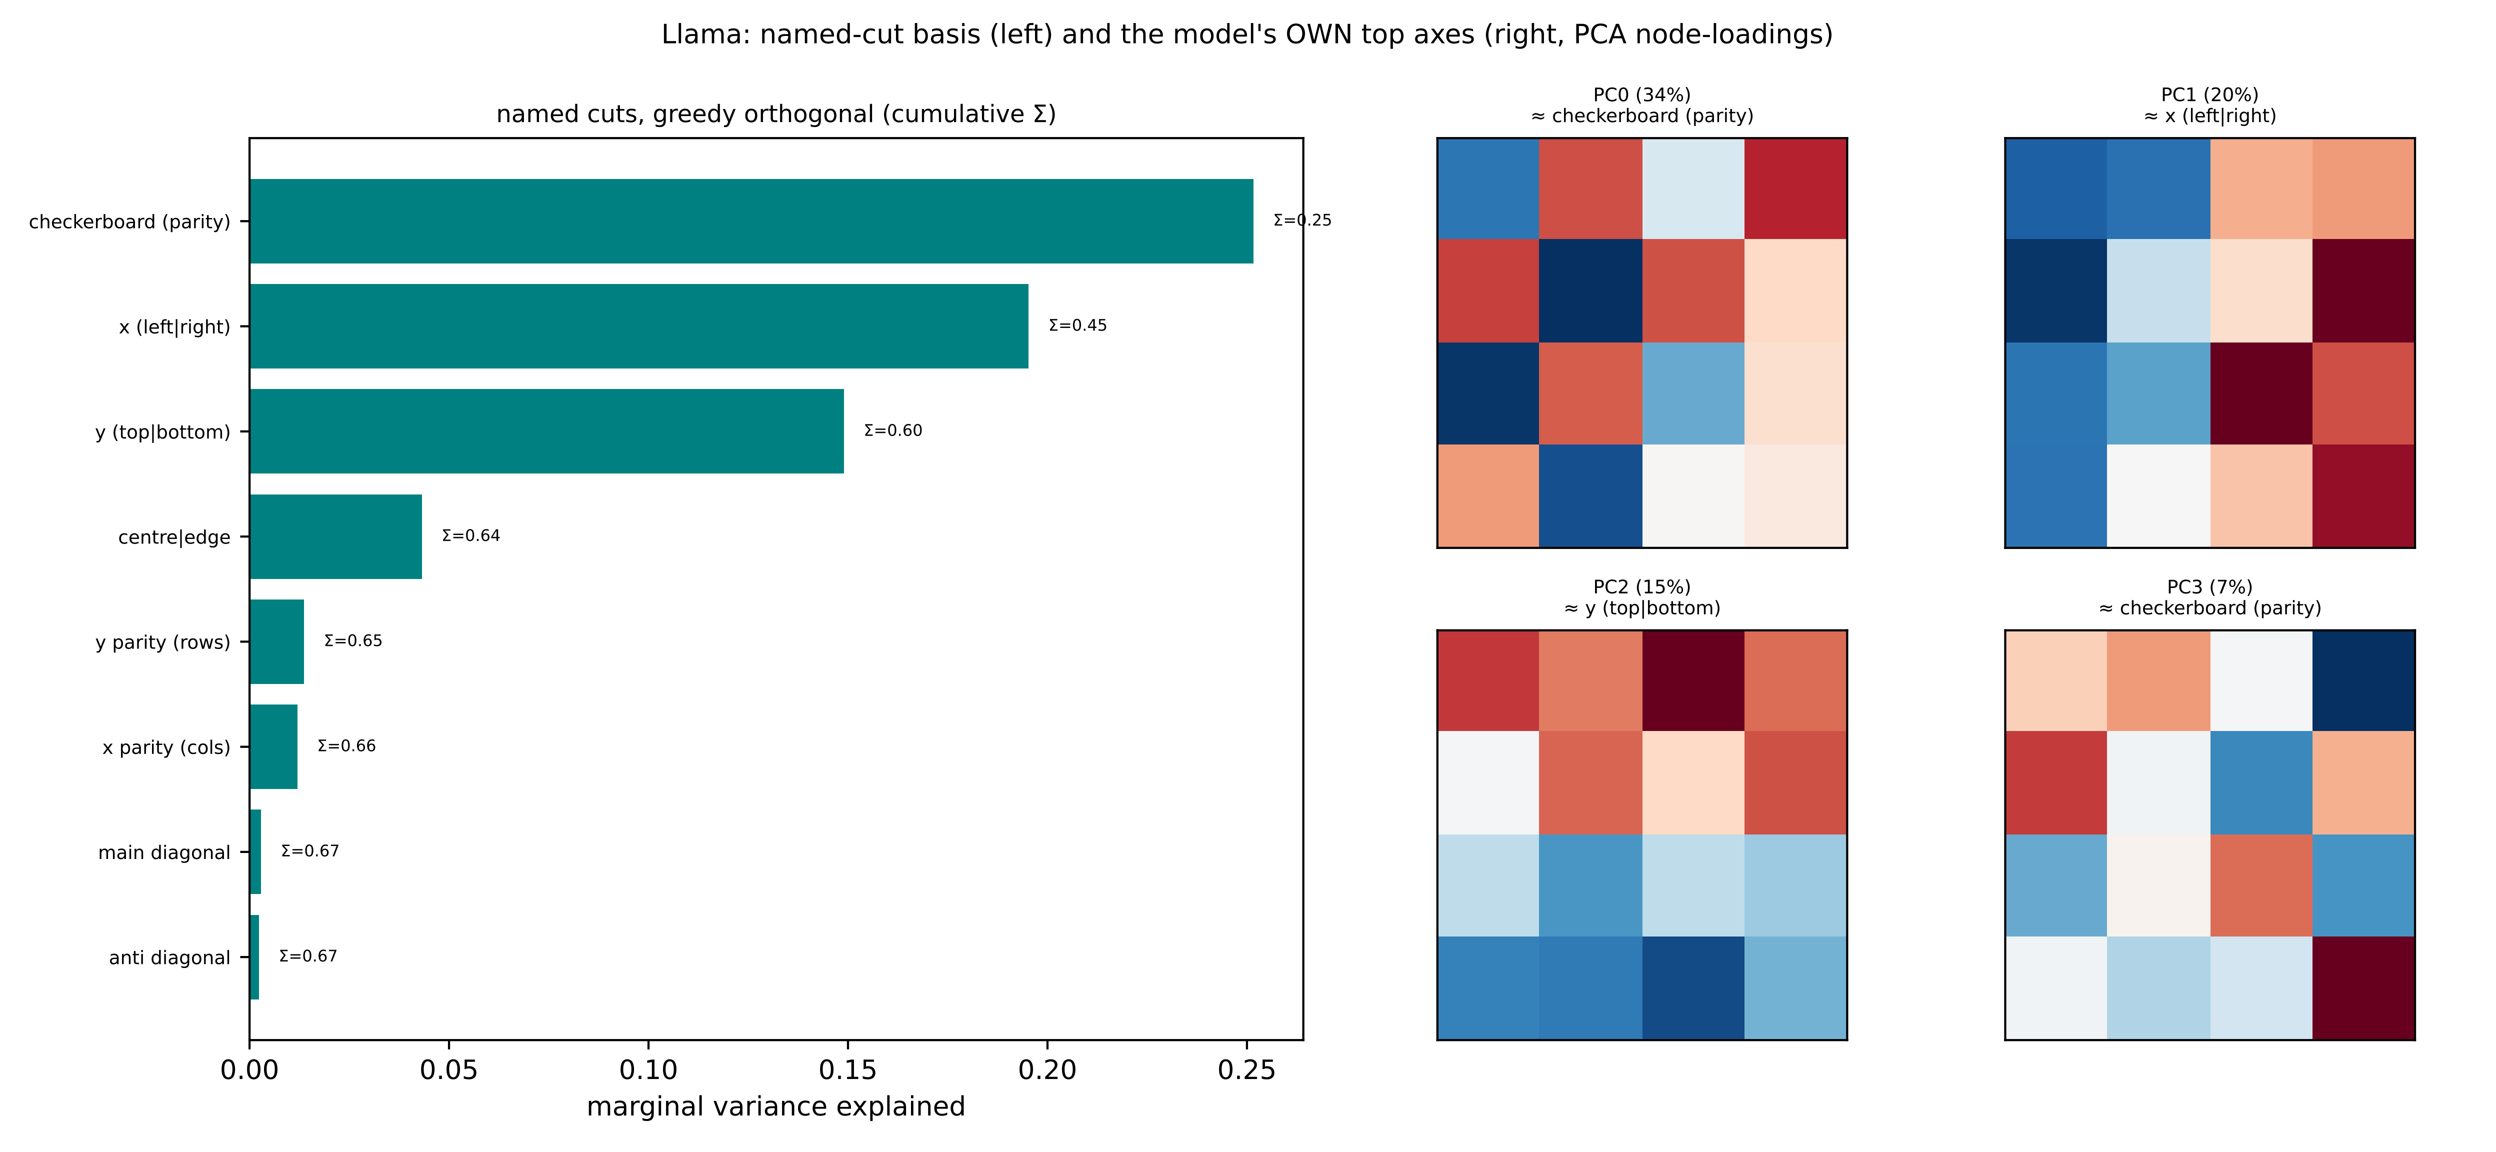

Llama: named-cut basis (left) and the model's OWN top axes (right, PCA node-loadings)
named cuts, greedy orthogonal (cumulative Σ)
checkerboard (parity)
x (left|right)
y (top|bottom)
centre|edge
y parity (rows)
x parity (cols)
main diagonal
anti diagonal
Σ=0.25
Σ=0.45
Σ=0.60
Σ=0.64
Σ=0.65
Σ=0.66
Σ=0.67
Σ=0.67
0.00
0.05
0.10
0.15
0.20
0.25
marginal variance explained
PC0 (34%)
≈ checkerboard (parity)
PC1 (20%)
≈ x (left|right)
PC2 (15%)
≈ y (top|bottom)
PC3 (7%)
≈ checkerboard (parity)