



[square\_grid] Llama — remove x / y / parity direction (rank-1, all layers) during generation.  
red=parity. Does the biggest descriptive cut carry the most causal weight?

neighbour mass (behaviour)



validity (behaviour)



linear decodability after removal

