

Next-node prediction is a DISTRIBUTED circuit: 2 heads recover ~0%, need ~24-48 heads for half, ~200+ for most (non-monotonic dips = mean-ablation off-distribution artifact)

