






[square_grid] DYNAMIC swap switch @ step 150 — swapped nodes: does the map relearn (orig R^2 falls, swapped R^2 rises)?




Llama L3 PRE




Gemma L4 PRE




Qwen L4 PRE




Llama L3 POST




Gemma L4 POST




Qwen L4 POST




Llama L5 PRE




Gemma L7 PRE




Qwen L6 PRE




Llama L5 POST









Gemma L7 POST



Qwen L6 POST




[square_grid] swap @ depth 0.26 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred




[square_grid] swap @ depth 0.33 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L10 PRE




Gemma L14 PRE




Qwen L12 PRE




Llama L10 POST



Gemma L14 POST




Qwen L12 POST




[square_grid] swap @ depth 0.41 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L13 PRE




Gemma L17 PRE




Qwen L14 PRE




Llama L13 POST



Gemma L17 POST




Qwen L14 POST




[square_grid] swap @ depth 0.49 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L15 PRE




Gemma L20 PRE




Qwen L17 PRE




Llama L15 POST




Gemma L20 POST




Qwen L17 POST




Llama L17 PRE




Gemma L23 PRE




Qwen L20 PRE




Llama L17 POST



Gemma L23 POST




Qwen L20 POST




[square_grid] swap @ depth 0.64 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L20 PRE




Gemma L26 PRE




Qwen L22 PRE




Llama L20 POST



Gemma L26 POST




Qwen L22 POST




[square_grid] swap @ depth 0.72 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L22 PRE




Gemma L29 PRE




Qwen L25 PRE




Llama L22 POST



Gemma L29 POST




Qwen L25 POST




[square_grid] swap @ depth 0.80 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L25 PRE




Gemma L33 PRE




Qwen L28 PRE




Llama L25 POST




Gemma L33 POST




Qwen L28 POST




Llama L27 PRE




Gemma L36 PRE




Qwen L31 PRE




Llama L27 POST



Gemma L36 POST




Qwen L31 POST




[square_grid] swap @ depth 0.95 — probe-axis activations (top=PRE-switch, bottom=POST); colour = node; swapped words starred


Llama L29 PRE




Gemma L39 PRE




Qwen L33 PRE



Llama L29 POST



Gemma L39 POST



Qwen L33 POST

