





Llama [square_grid]: direct logit attribution of the next-step prediction

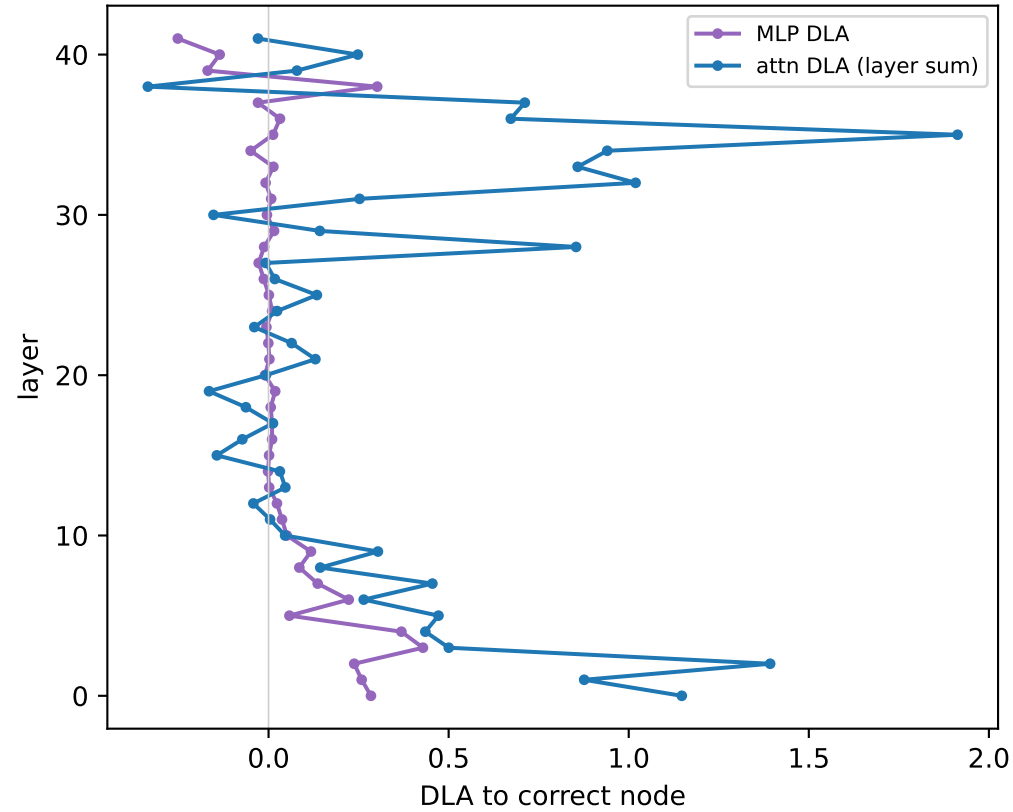


Gemma [square_grid]: direct logit attribution of the next-step prediction

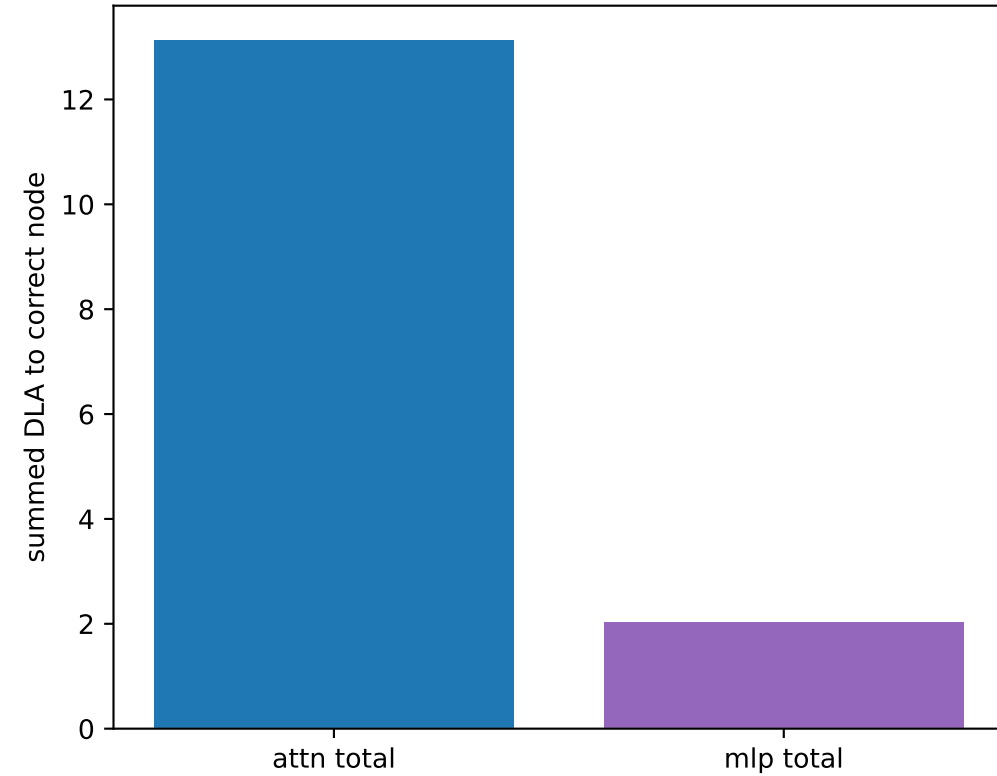
Gemma: per-head DLA to correct next node



Gemma: per-layer attention vs MLP DLA



Gemma: who writes the answer? (attn +13.1 vs mlp +2.0)



Qwen [square_grid]: direct logit attribution of the next-step prediction

