



Per-head direct-logit-attribution to the correct next node (log y) — bulk near 0, thin tail of 'writer' heads (red)


Llama (1024 heads; mean +0.004, max +0.32, min -0.17)



Gemma (672 heads; mean +0.020, max +1.04, min -0.41)



Qwen (1152 heads; mean +0.098, max +10.86, min -2.42)



DLA distributions, overlaid (density, log y)

