


Where the attractor lives: mid-stack self-history encoding, late-stack routing & suppression

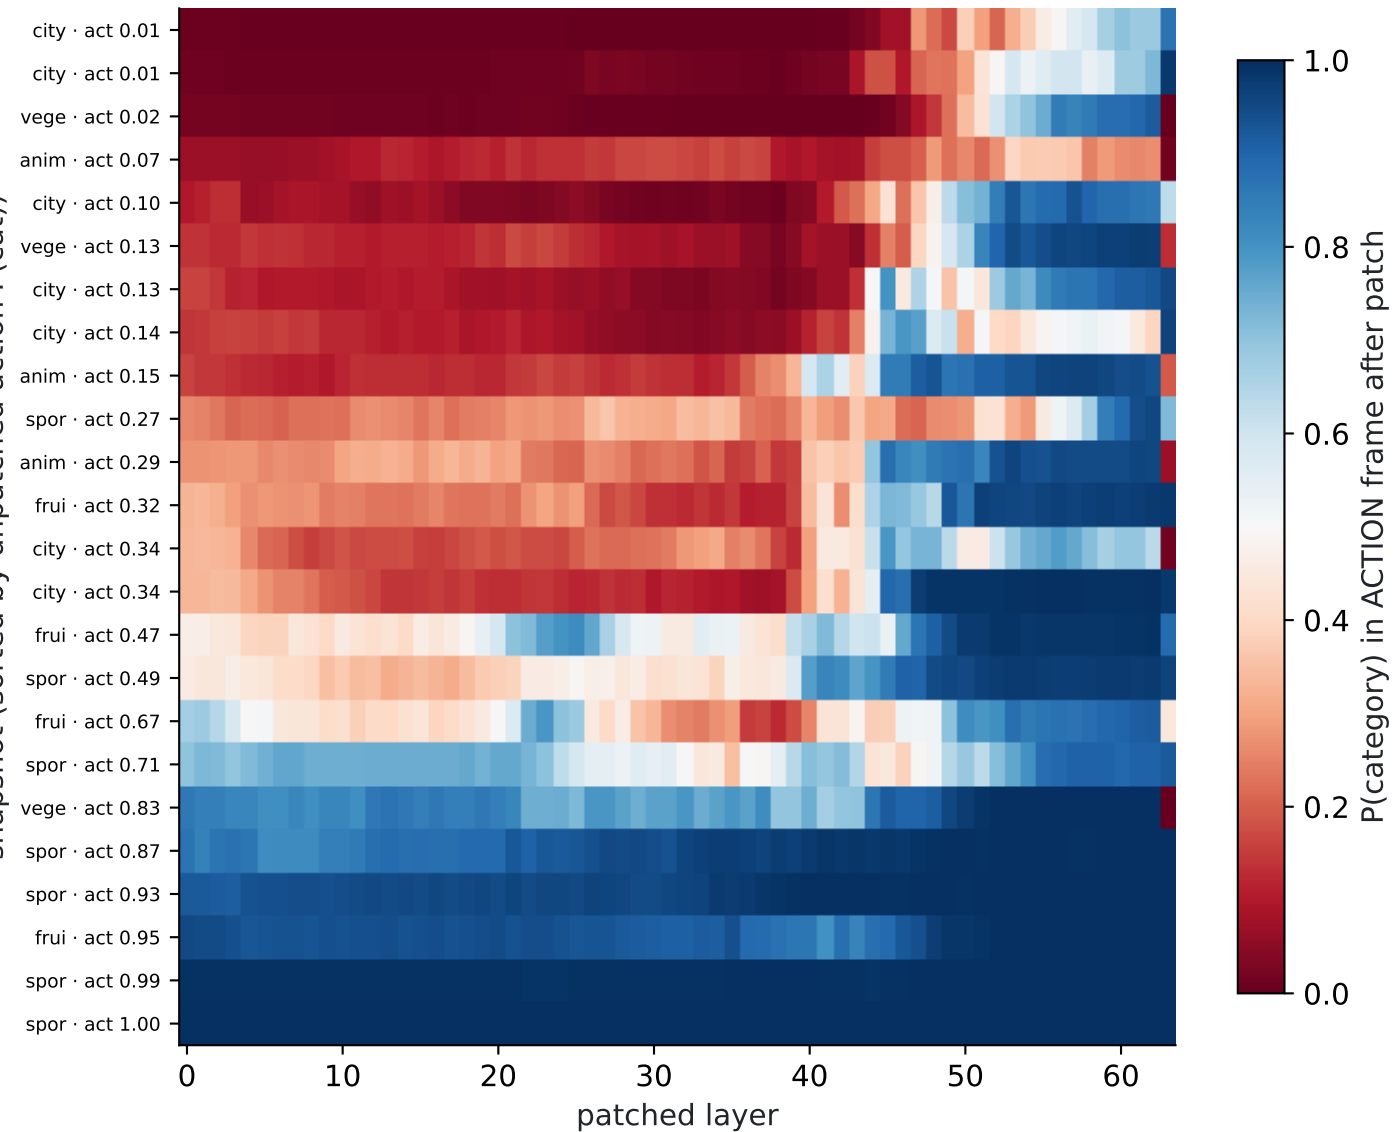


final-layer point (L63) excluded from the final-position curve: replacing the last residual is trivially control-like (overshoot 2.5). Span patches mean-pool when donor token lengths differ — early-layer span values

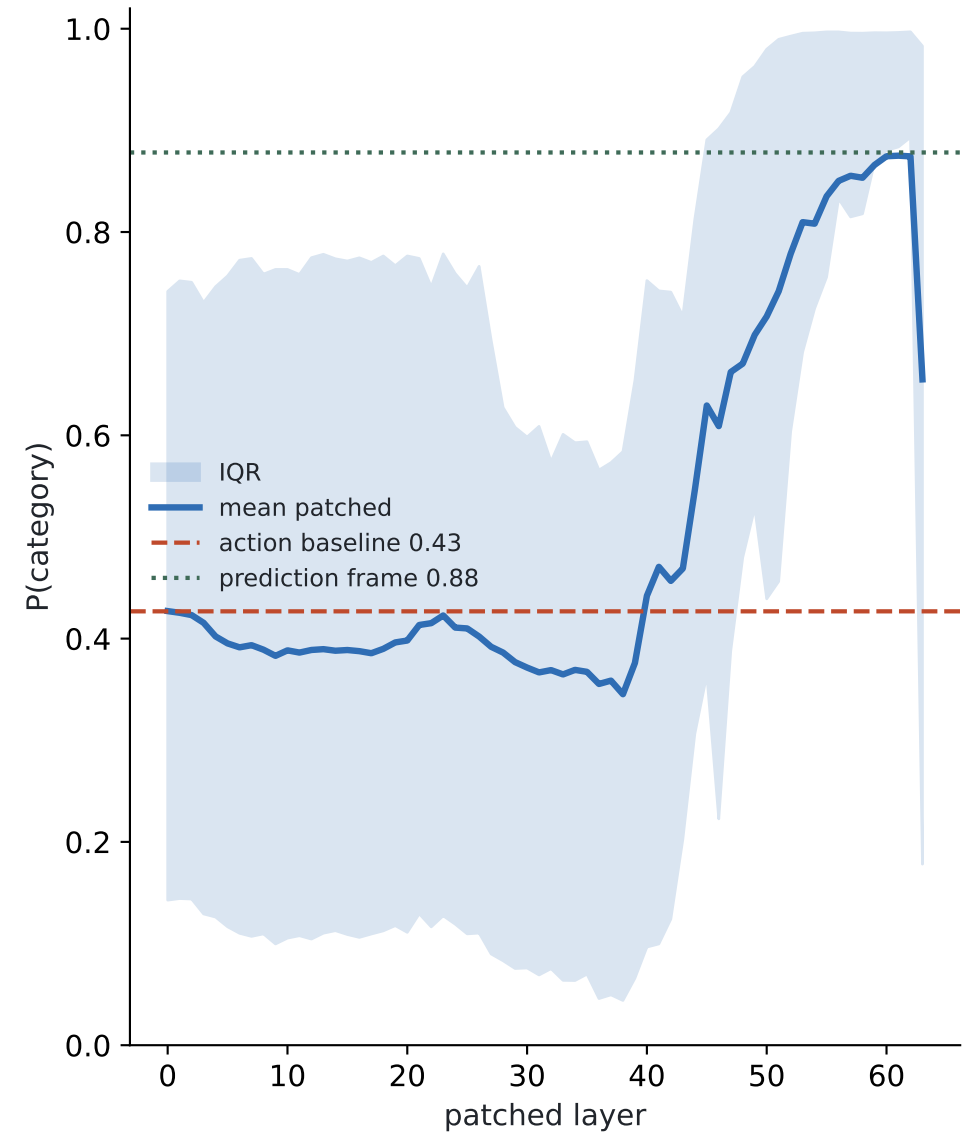
Prediction → action frame patching

same reconstructed loop state, same shuffled 8-candidate listing; prediction-frame final-position residual patched into the action frame, one layer at a time (24 snapshot)

Per-snapshot patched action-frame category mass



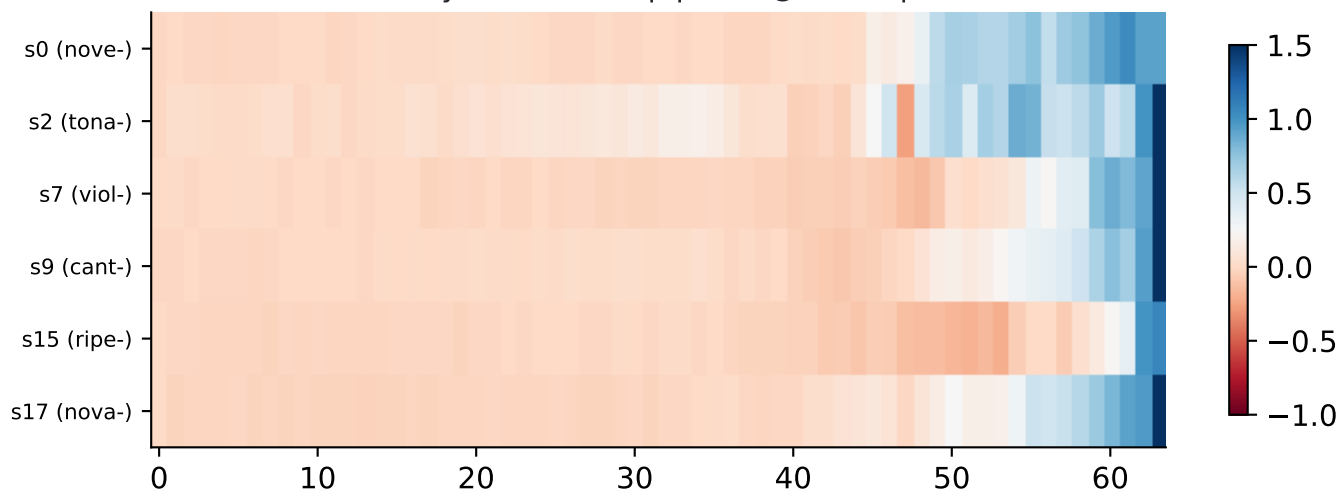
Mean transfer curve



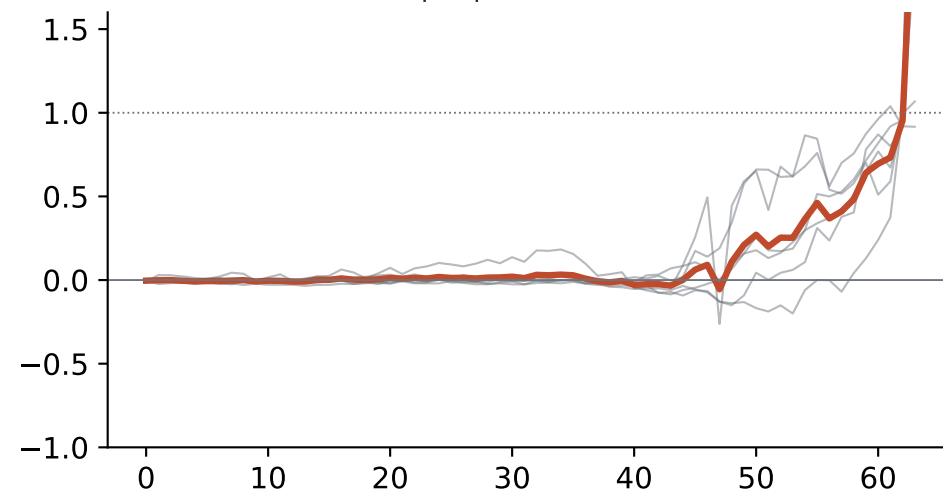
Paired transplant patching (loop-seed vs neutral-seed, identical partner stream)

$\Delta = \text{logsumexp}(\text{self-stem first-token logits}) - \text{logsumexp}(\text{category first-token logits})$; recovery 1 = patched loop run reaches the control Δ .

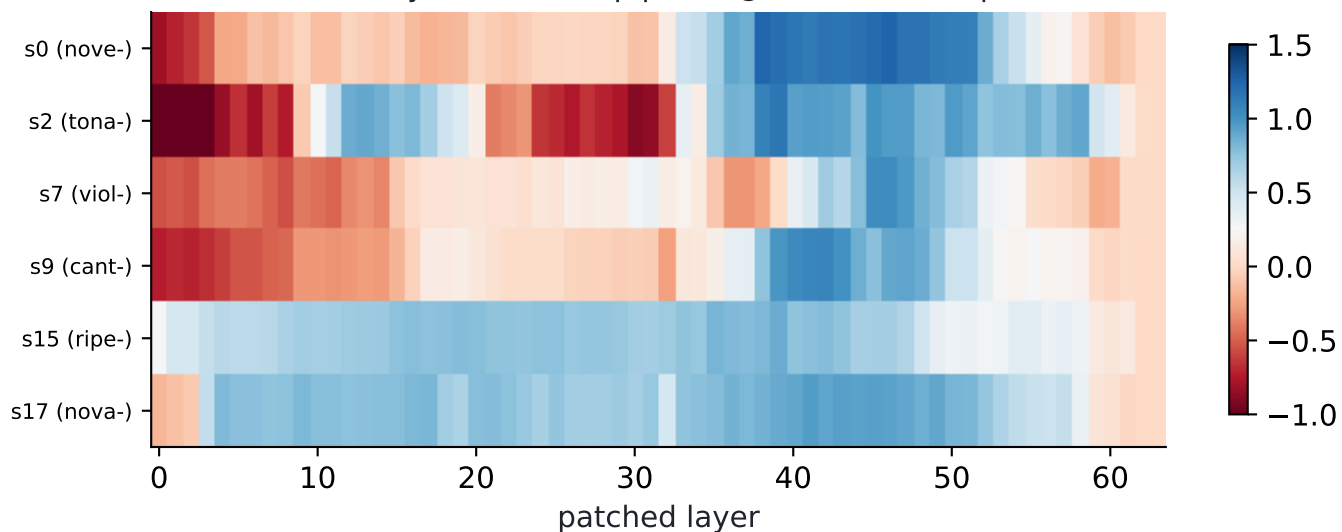
Δ recovery, control→loop patch @ FINAL position



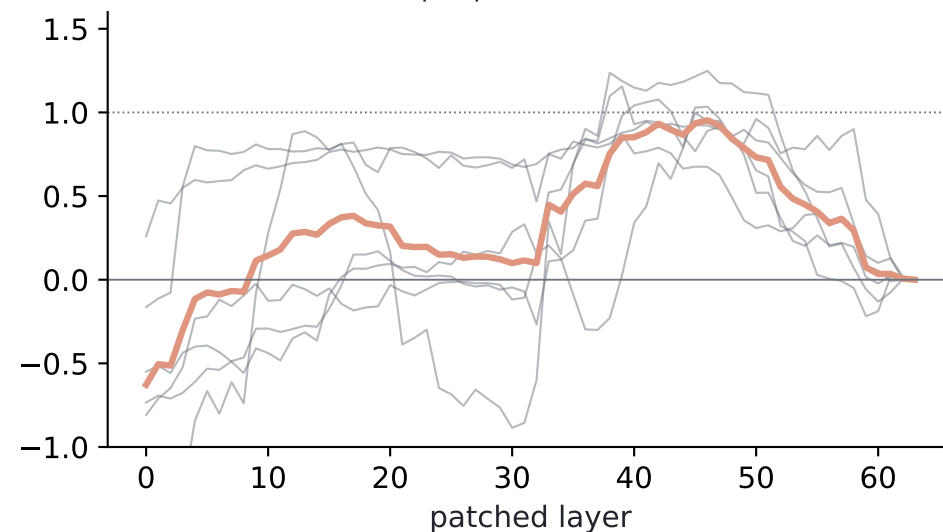
per-pair + mean



Δ recovery, control→loop patch @ SEED-WORD spans



per-pair + mean



Baselines and sublayer decomposition

