



Semantic run detector — embedding-based (text-embedding-3-small)

run = consecutive words with $\cos(\text{word}, \text{run centroid}) \geq \tau = 0.357$ (95th pct of cross-category similarity; within-category pairs exceed τ 65% of the time)

Semantic runs split by alignment with B's words



Longest run per model



Longest SELF-DRIFT semantic runs (not aligned with partner's words):

Qwen3-32B (strict)	echoingest → echoedest → unlockedest → unlockeded → unlockedly → echoedly → echoingly → unlockedness → echoing
Llama-70B (cap100)	same → same → same → same → same → same → same → same → same → same → same → same...
gpt-4o-mini	craftsmanship → craftsmanship → craftsmanship → craftsmanship → craftsmanship → craftsmanship → craftsmanship
gpt-4.1	lingonberry → goji → seaweed → kelp → dulse → kombu
gpt-5	unscripted → improv → improvised → improvisation → improvise → improviser
Haiku 4.5	—
Sonnet 4.5	—

caveats: 'aligned' = run centroid near EARLY-B centroid (coarse — gpt-4.1's geographic wandering counts as aligned); drift runs include repeat-locks (identical words are trivially similar).