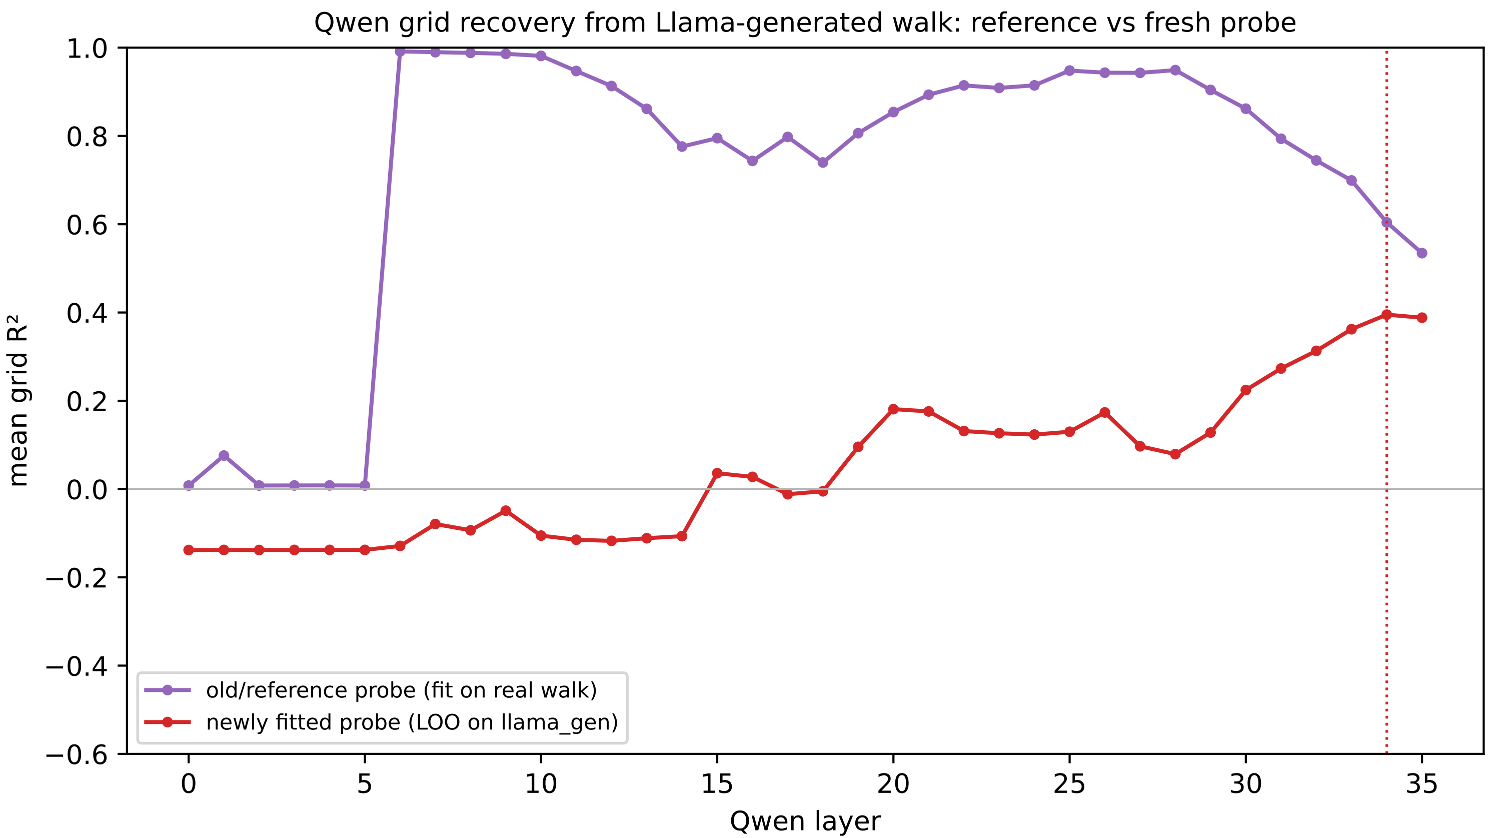

Qwen grid recovery from Llama-generated walk: reference vs fresh probe
1.0
0.8
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
mean grid R²
0
5
10
15
20
25
30
35
Qwen layer
old/reference probe (fit on real walk)
newly fitted probe (LOO on llama_gen)

Qwen LAYER 0/35 — 16 node-means from Llama's generated walk, decoded into grid coords

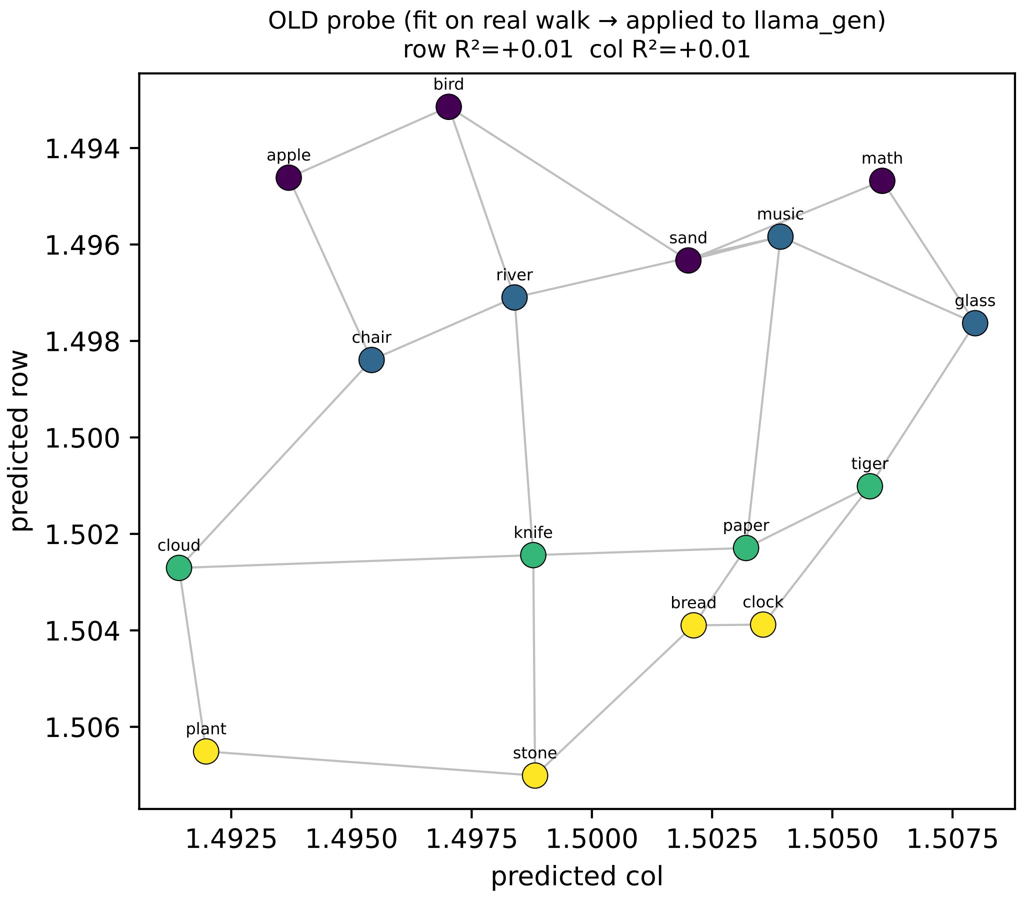

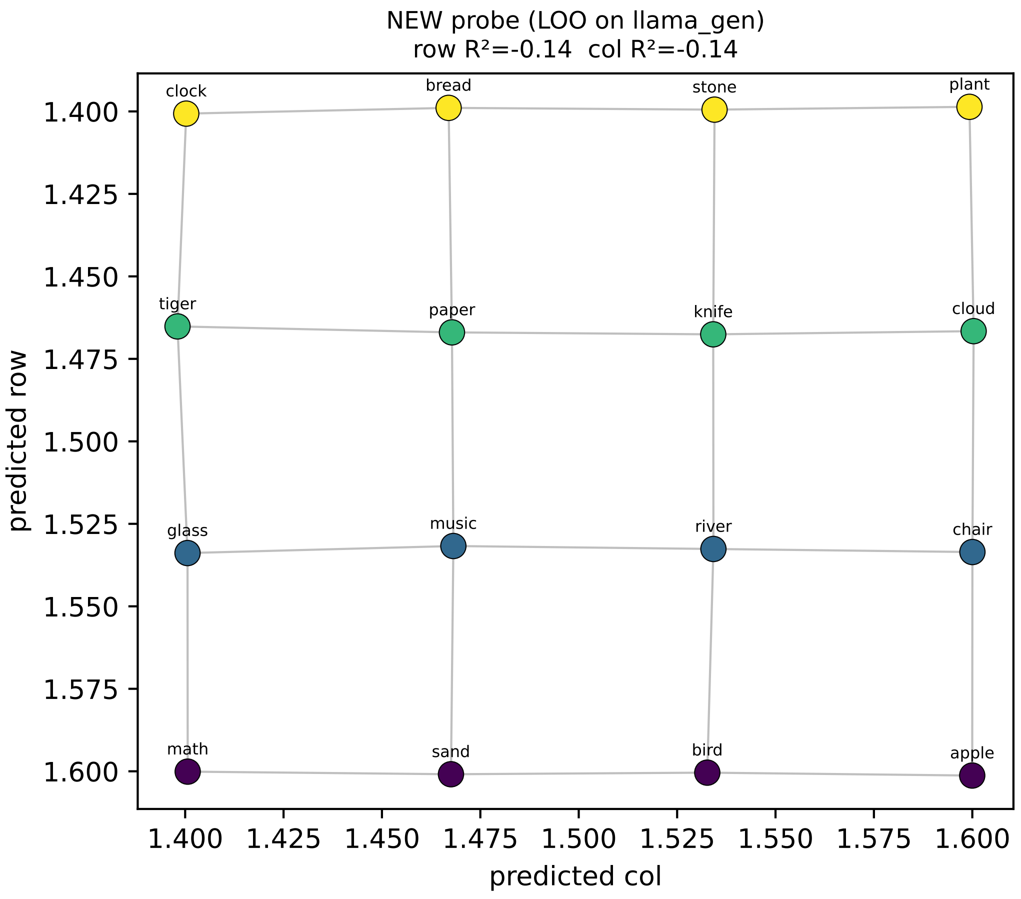

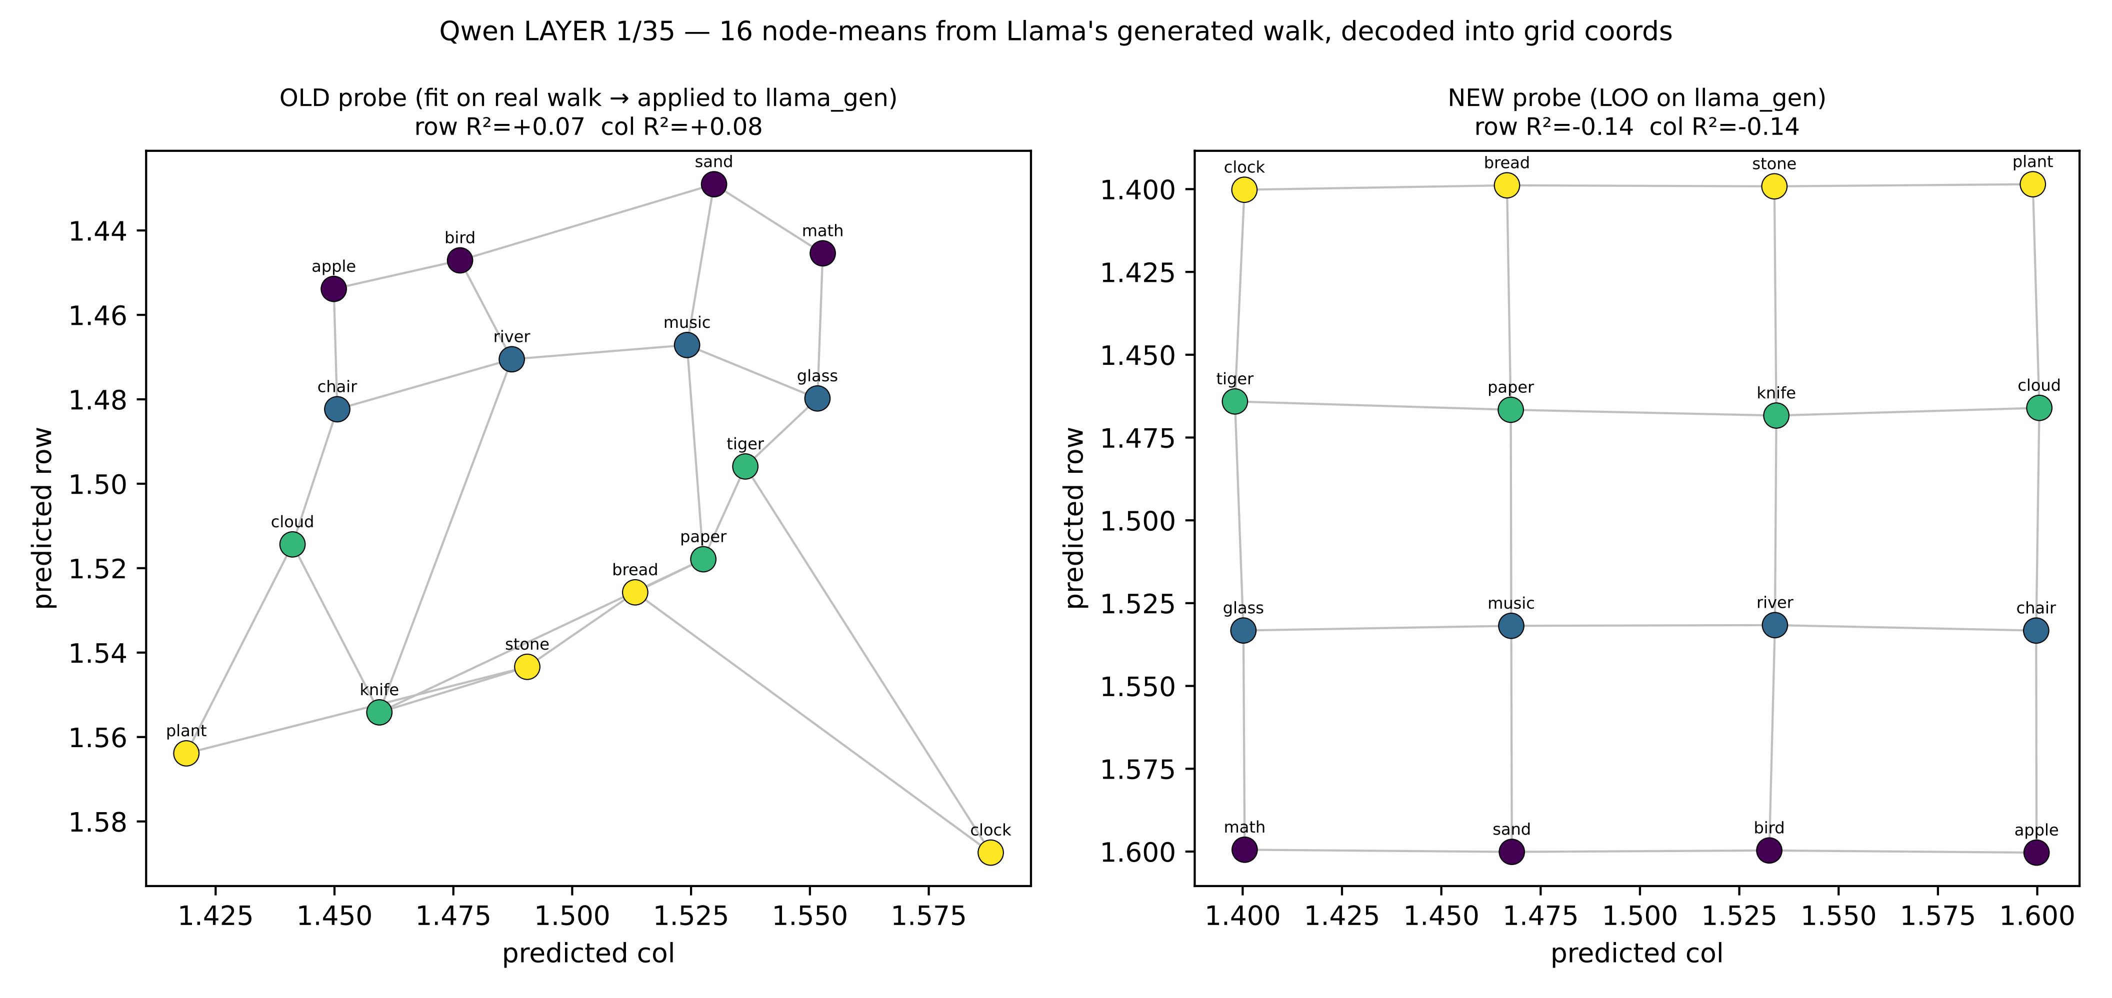
Qwen LAYER 1/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.07 col R²=+0.08
NEW probe (LOO on llama_gen)
row R²=-0.14 col R²=-0.14
predicted row
predicted col
sand
math
bird
apple
river
music
glass
chair
tiger
cloud
paper
bread
stone
knife
plant
clock
1.44
1.46
1.48
1.50
1.52
1.54
1.56
1.58
1.425
1.450
1.475
1.500
1.525
1.550
1.575
1.400
1.600

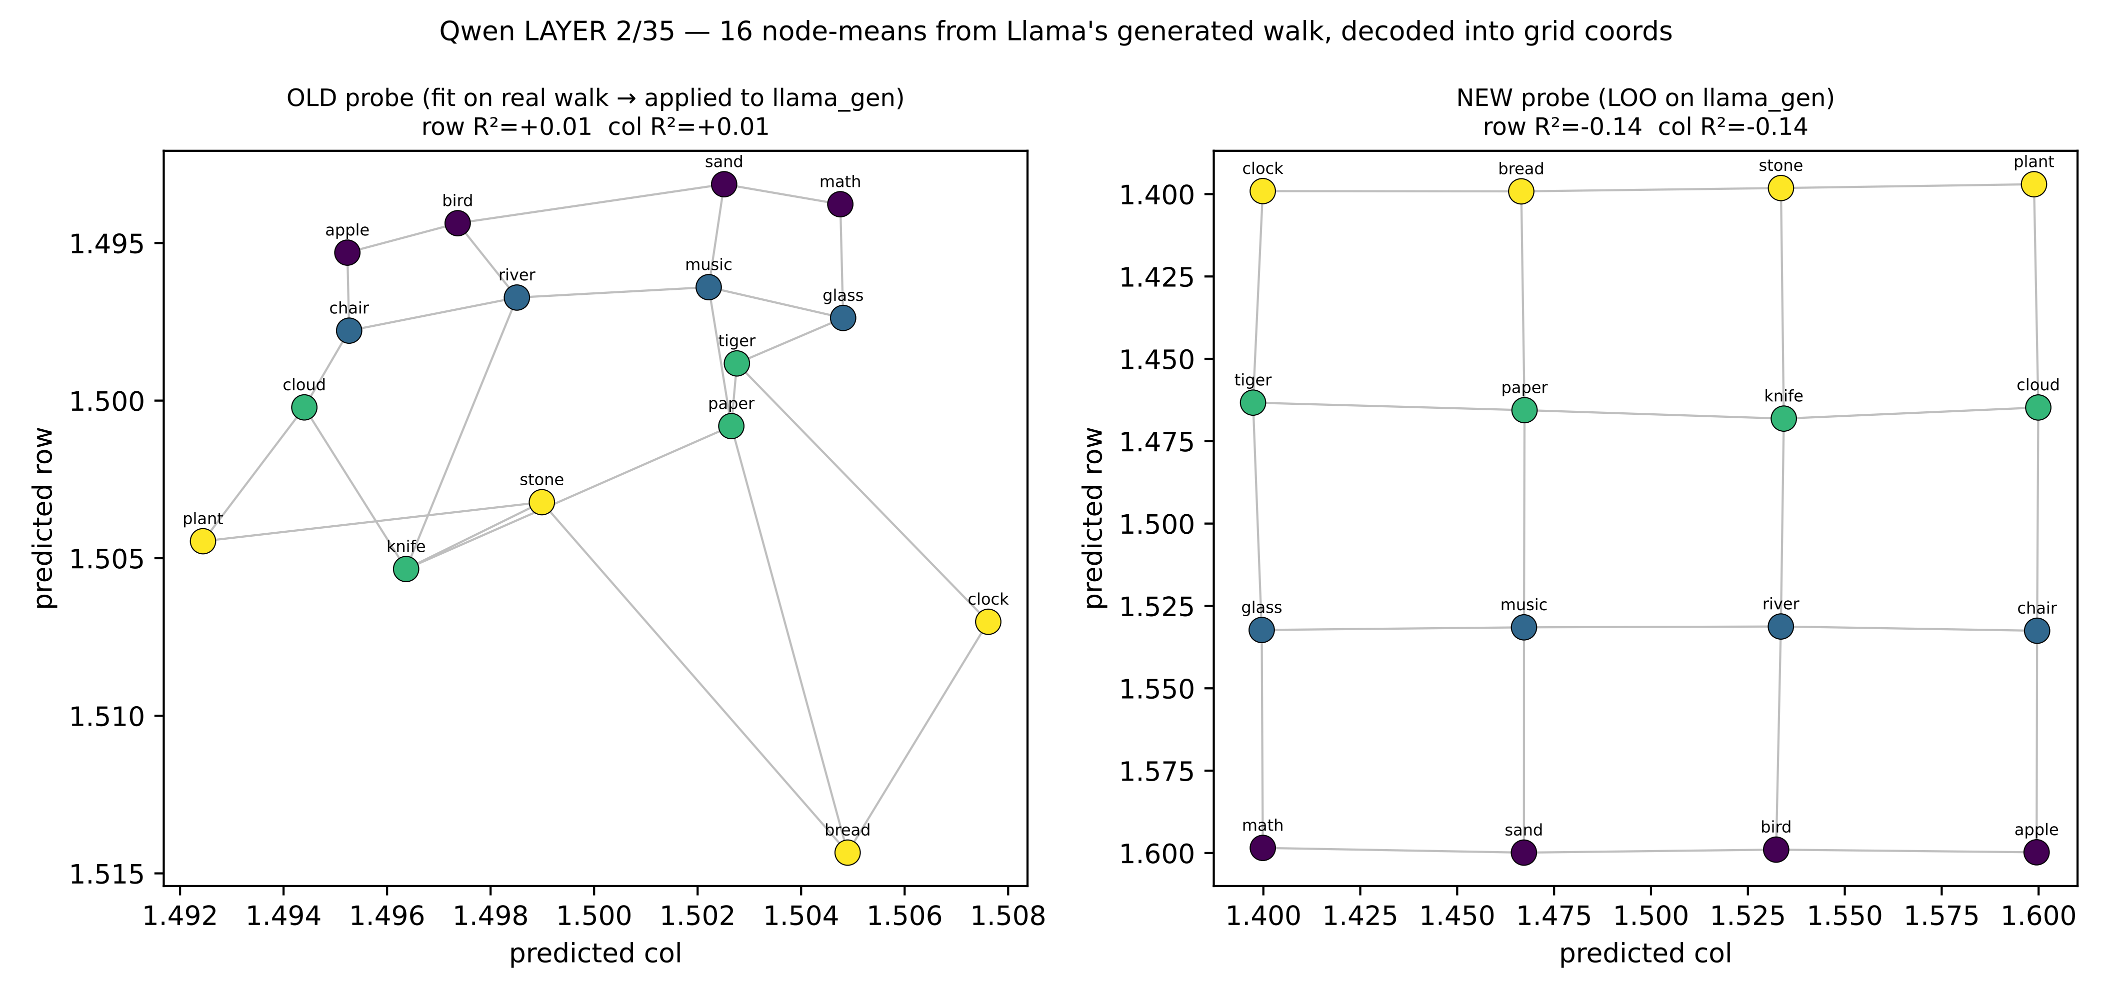
Qwen LAYER 2/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.01 col R²=+0.01
predicted row
predicted col
sand
math
bird
apple
river
music
glass
chair
tiger
cloud
paper
stone
plant
knife
clock
bread
1.495
1.500
1.505
1.510
1.515
1.492
1.494
1.496
1.498
1.500
1.502
1.504
1.506
1.508
NEW probe (LOO on llama_gen)
row R²=-0.14 col R²=-0.14
predicted row
predicted col
clock
bread
stone
plant
tiger
paper
knife
cloud
glass
music
river
chair
math
sand
bird
apple
1.400
1.425
1.450
1.475
1.500
1.525
1.550
1.575
1.600
1.400
1.425
1.450
1.475
1.500
1.525
1.550
1.575
1.600

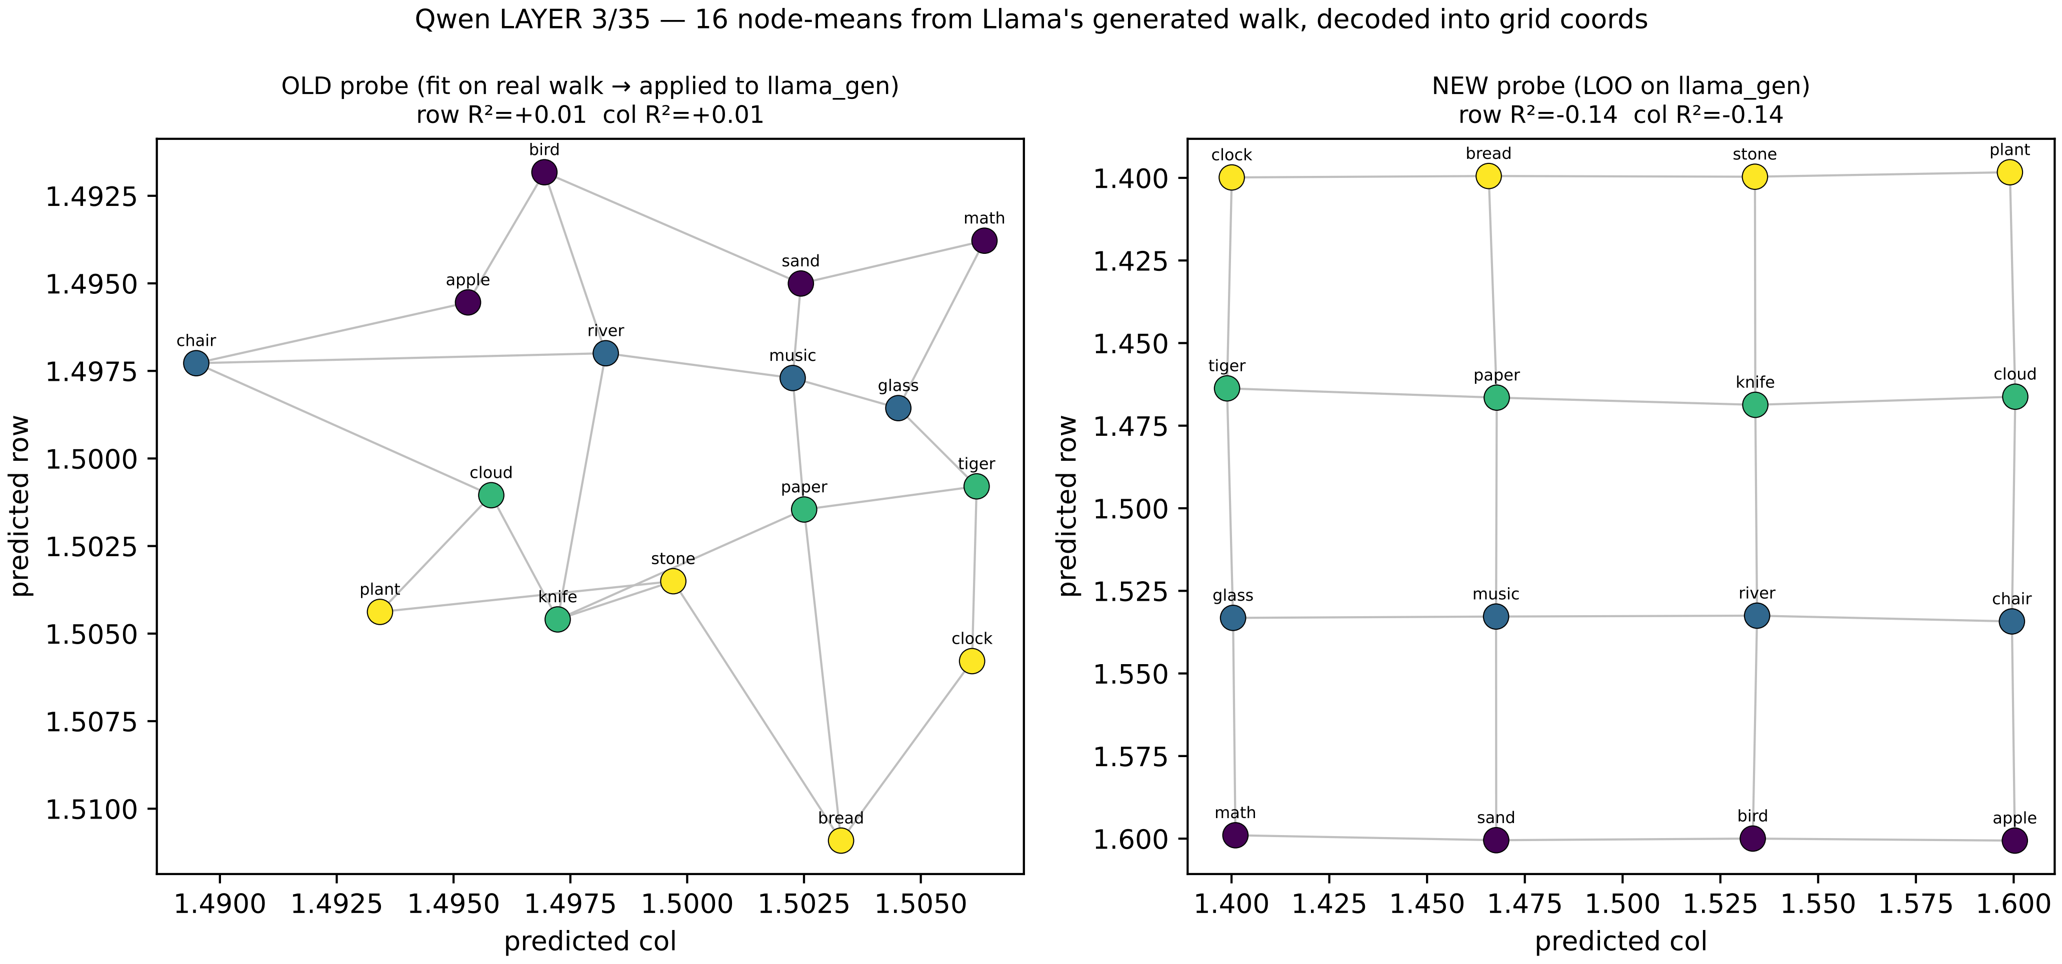
Qwen LAYER 3/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.01 col R²=+0.01
NEW probe (LOO on llama_gen)
row R²=-0.14 col R²=-0.14
predicted row
predicted col
bird
apple
sand
math
chair
river
music
glass
cloud
paper
tiger
stone
plant
knife
clock
bread
tiger
cloud
glass
music
river
chair
1.4900
1.4925
1.4950
1.4975
1.5000
1.5025
1.5050
1.5075
1.5100
1.400
1.425
1.450
1.475
1.500
1.525
1.550
1.575
1.600

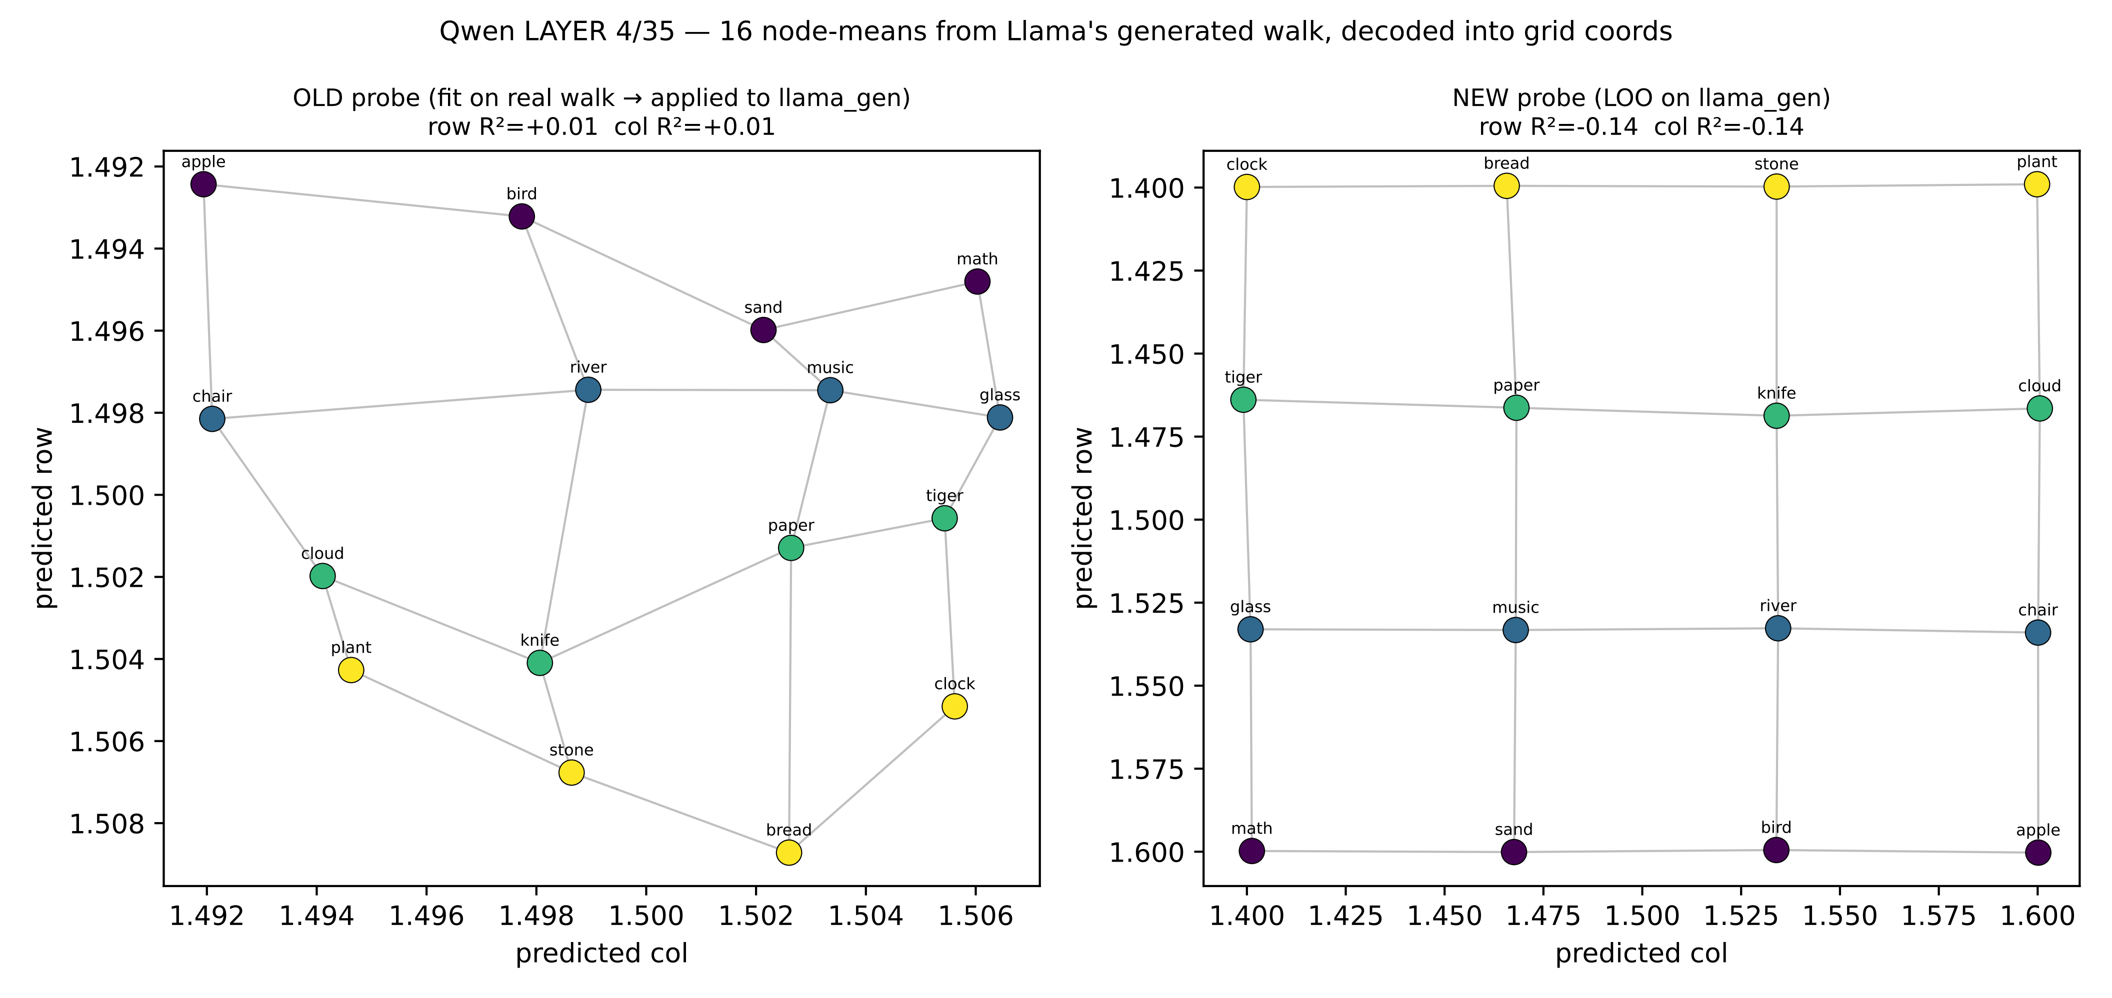
Qwen LAYER 4/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.01 col R²=+0.01
NEW probe (LOO on llama_gen)
row R²=-0.14 col R²=-0.14
predicted row
predicted col
apple
bird
math
sand
river
music
chair
glass
tiger
paper
cloud
knife
plant
clock
stone
bread
1.492
1.494
1.496
1.498
1.500
1.502
1.504
1.506
1.508
1.400
1.425
1.450
1.475
1.500
1.525
1.550
1.575
1.600

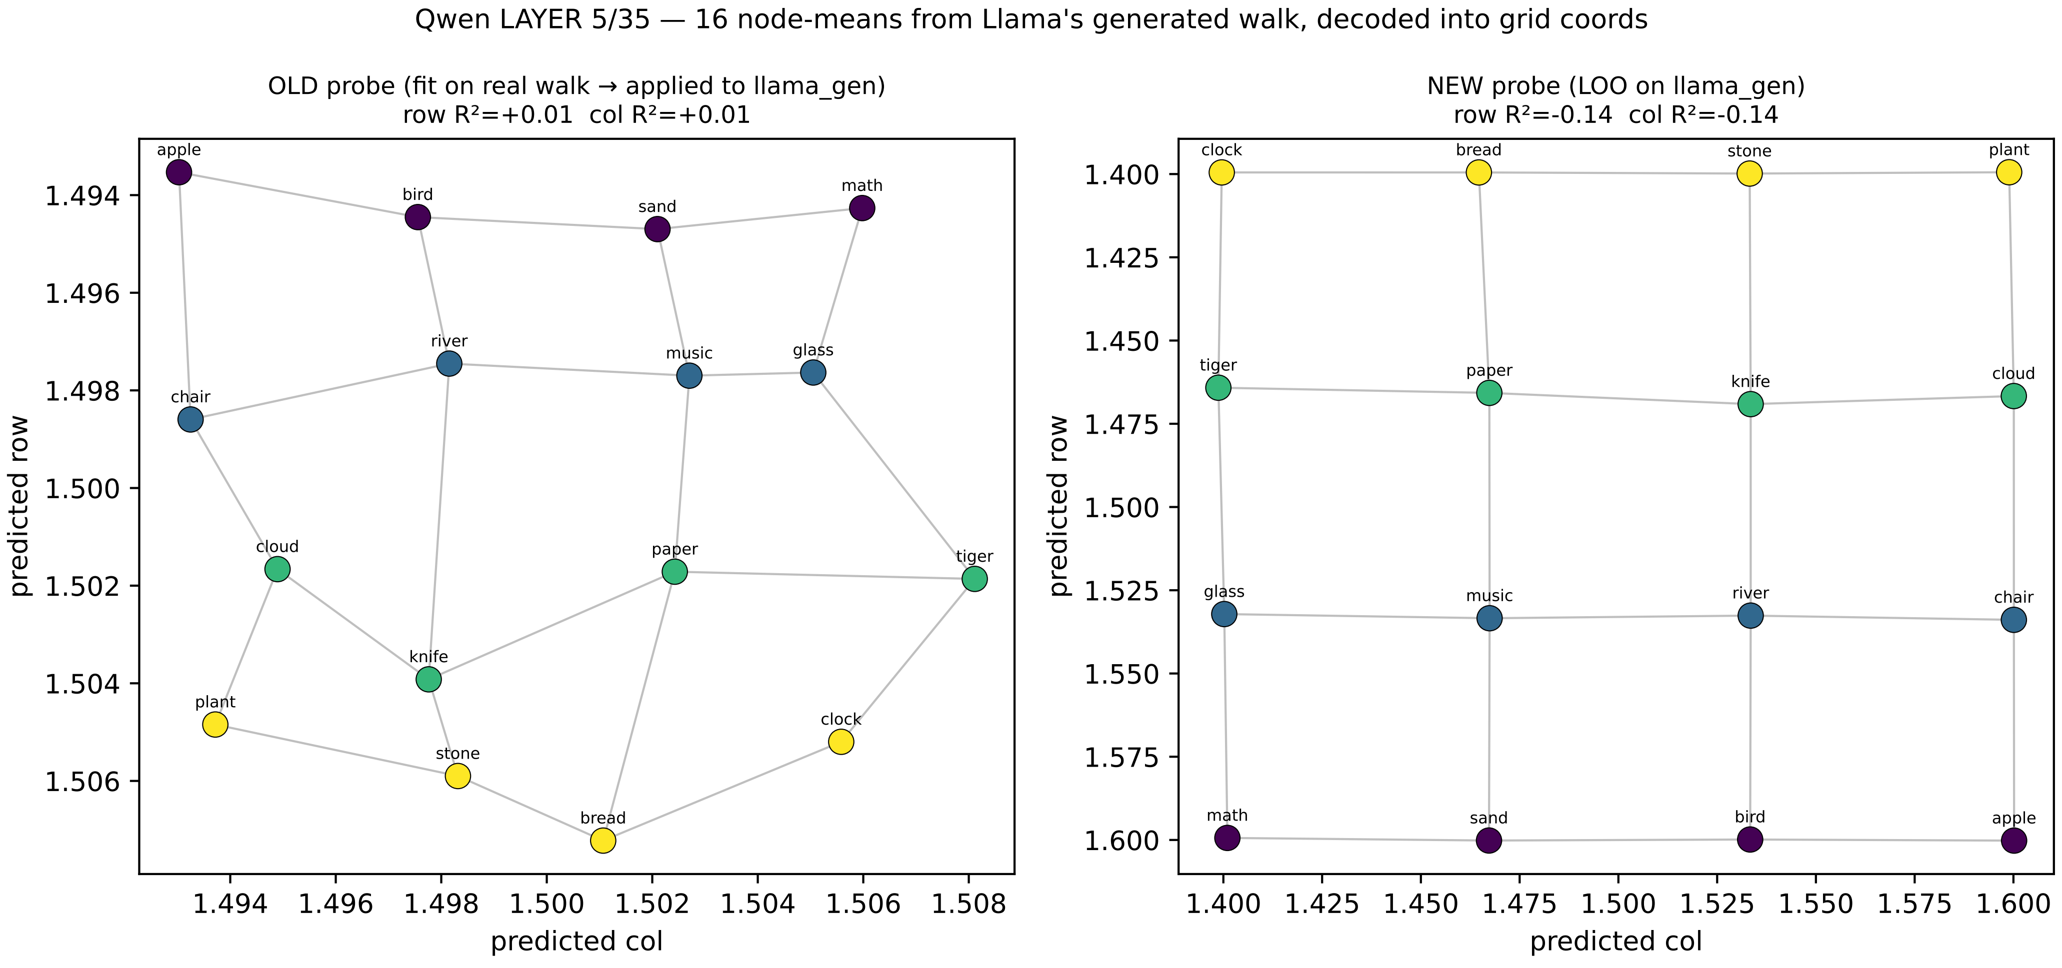

Qwen LAYER 5/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.01 col R²=+0.01
apple
bird
sand
math
river
music
glass
chair
cloud
paper
tiger
knife
plant
stone
clock
bread
predicted row
predicted col
1.494
1.496
1.498
1.500
1.502
1.504
1.506
1.508
NEW probe (LOO on llama_gen)
row R²=-0.14 col R²=-0.14
clock
bread
stone
plant
tiger
paper
knife
cloud
glass
music
river
chair
math
sand
bird
apple
predicted row
predicted col
1.400
1.425
1.450
1.475
1.500
1.525
1.550
1.575
1.600

Qwen LAYER 6/35 — 16 node-means from Llama's generated walk, decoded into grid coords

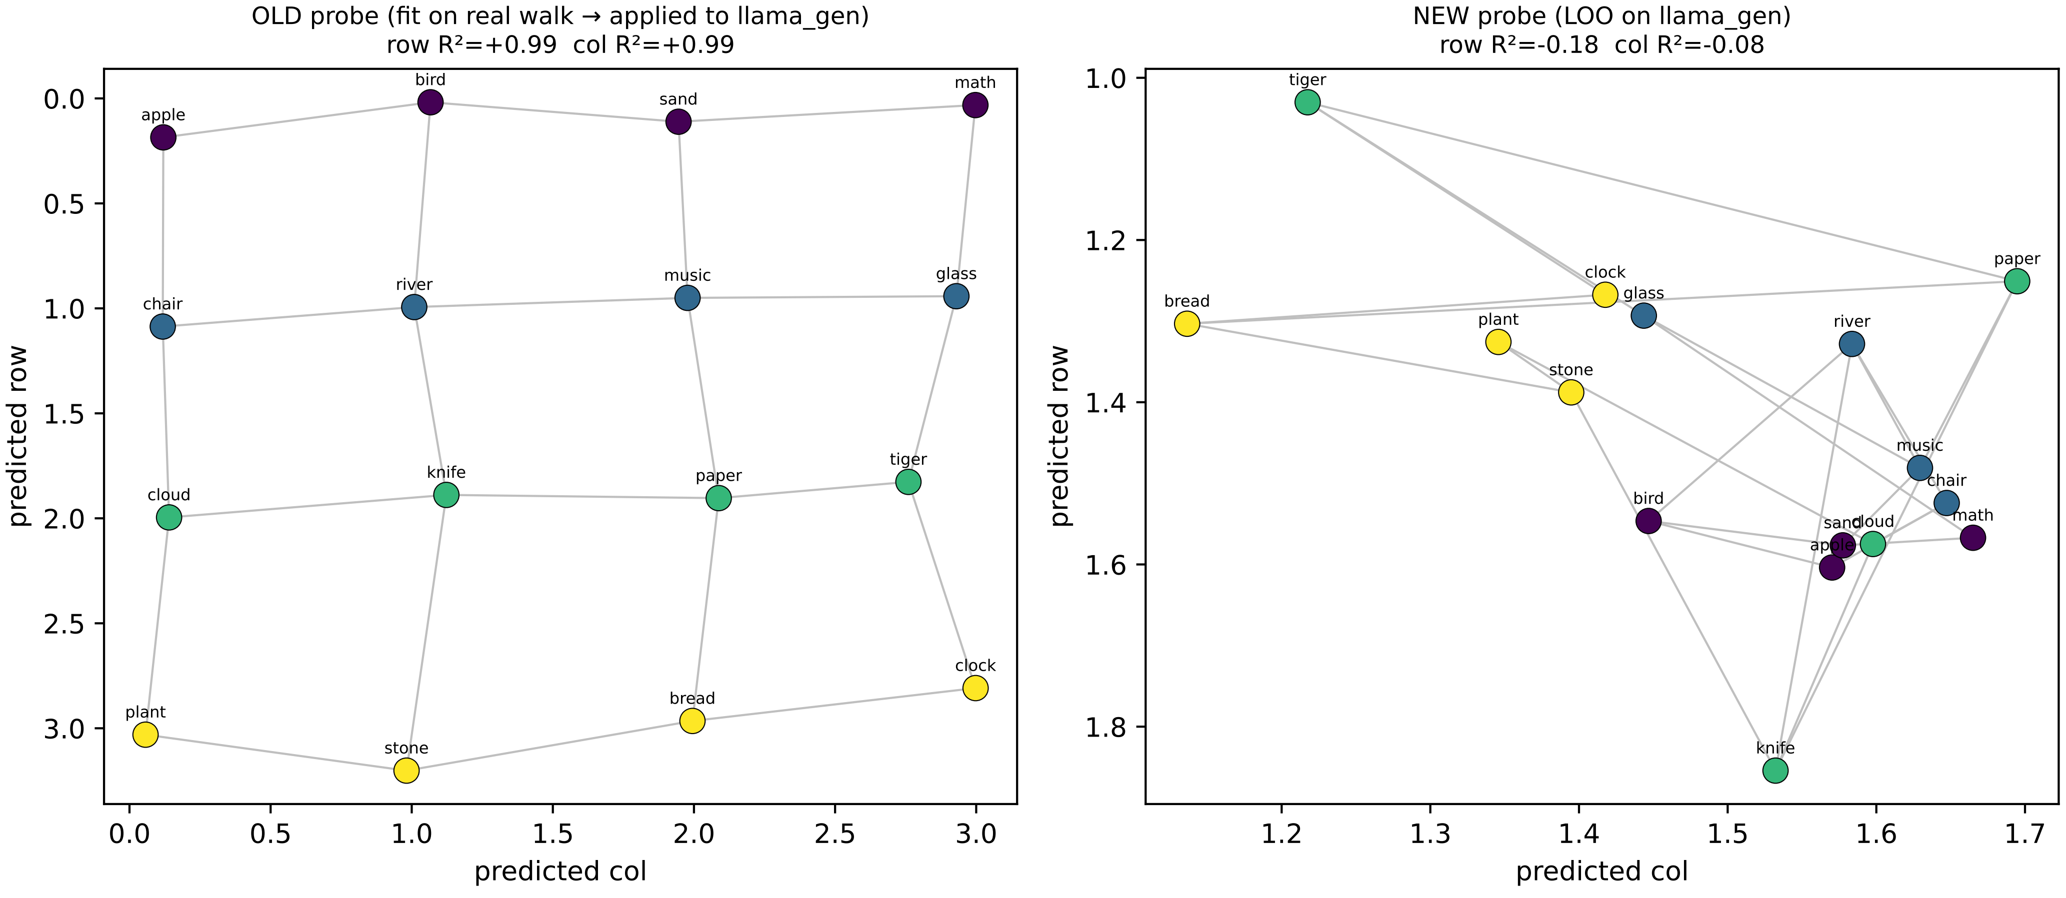

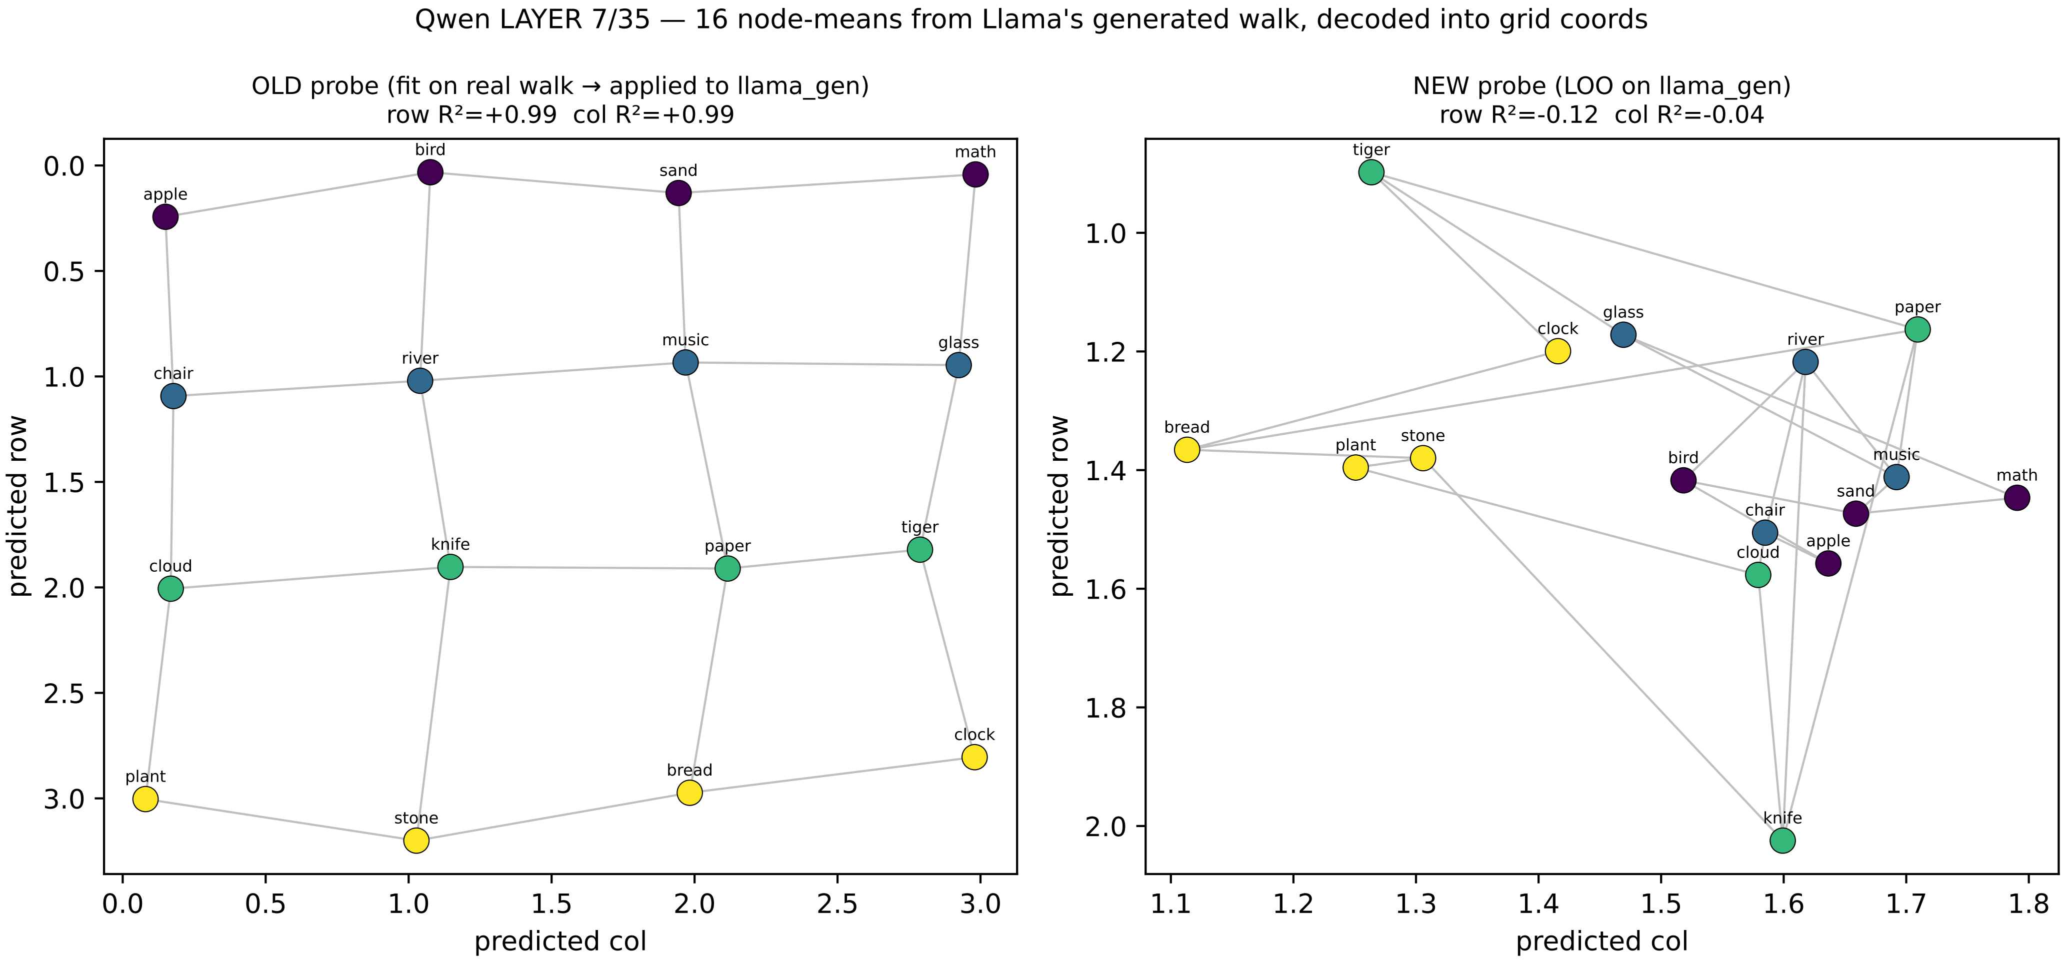
Qwen LAYER 7/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.99 col R²=+0.99
NEW probe (LOO on llama_gen)
row R²=-0.12 col R²=-0.04
predicted row
predicted col
apple
bird
sand
math
chair
river
music
glass
cloud
knife
paper
tiger
plant
stone
bread
clock

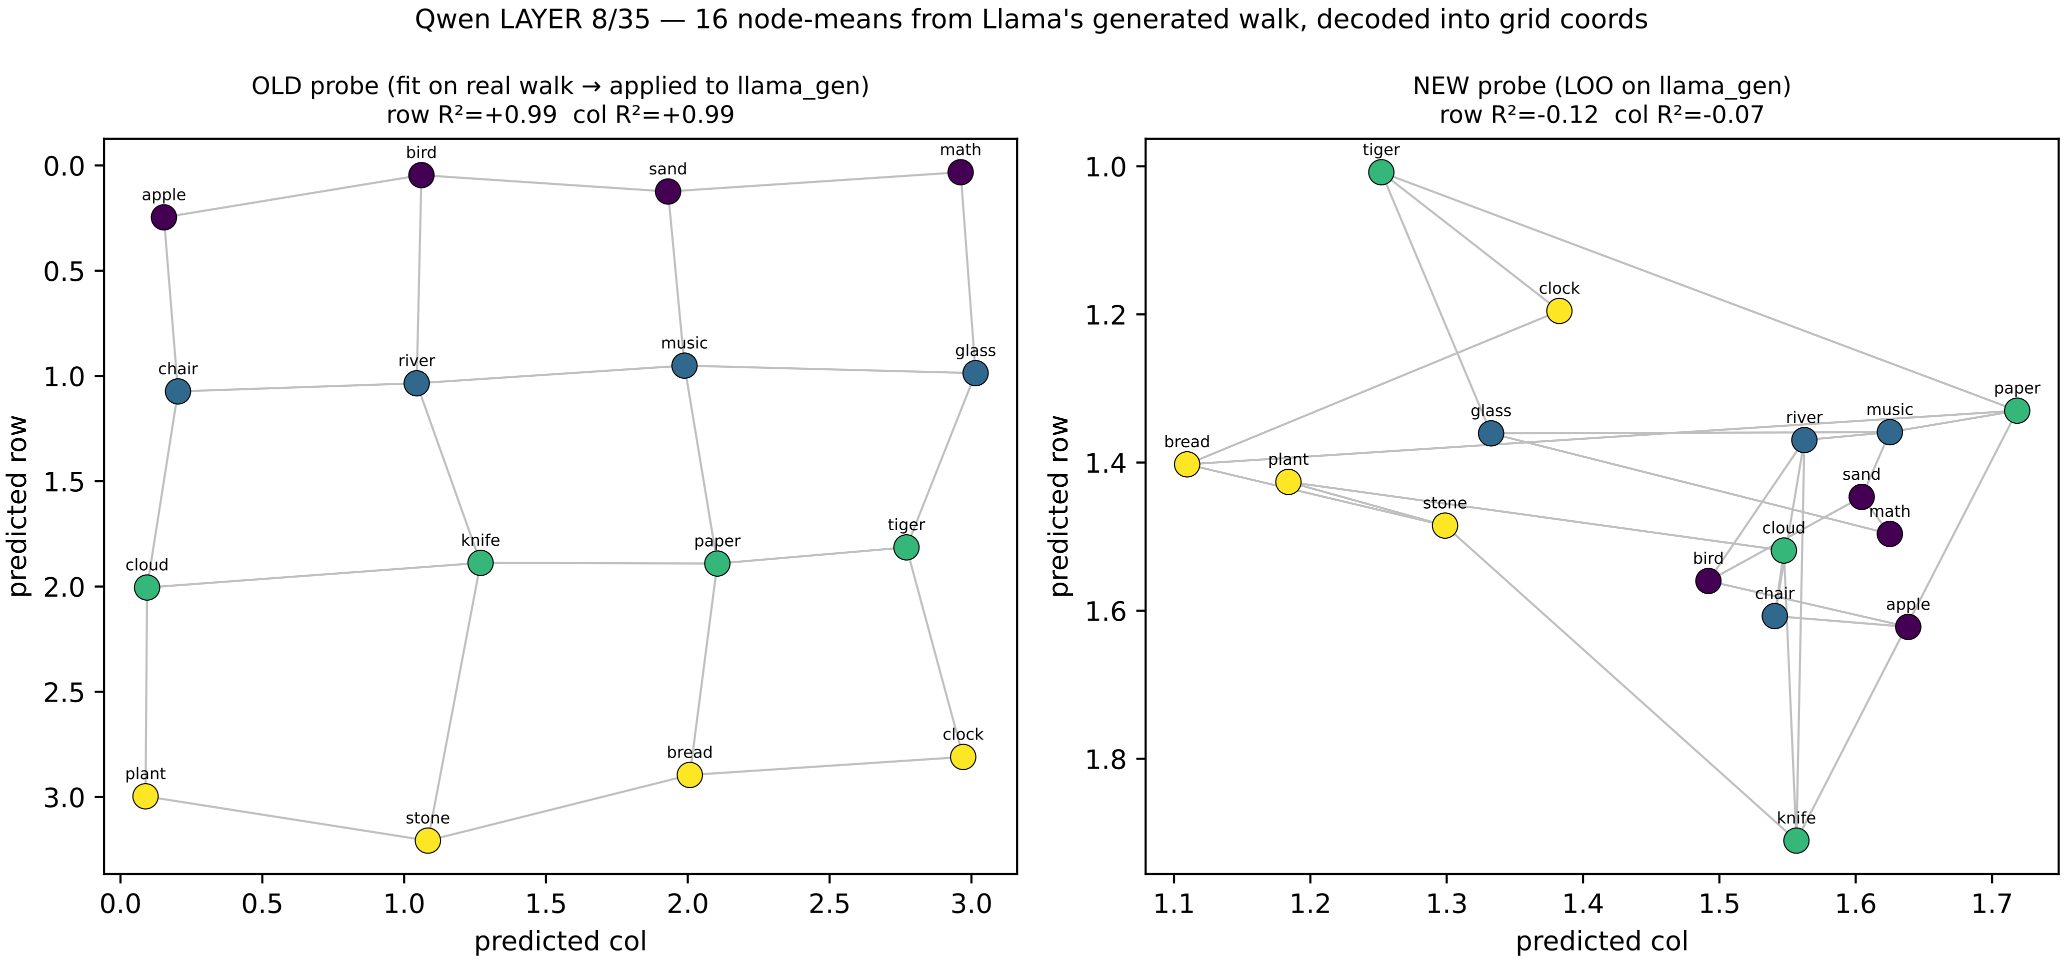
Qwen LAYER 8/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.99 col R²=+0.99
predicted row
predicted col
apple
bird
sand
math
chair
river
music
glass
cloud
knife
paper
tiger
plant
stone
bread
clock
NEW probe (LOO on llama_gen)
row R²=-0.12 col R²=-0.07
predicted row
predicted col
tiger
clock
paper
glass
river
music
bread
plant
sand
stone
math
cloud
bird
chair
apple
knife

Qwen LAYER 9/35 — 16 node-means from Llama's generated walk, decoded into grid coords

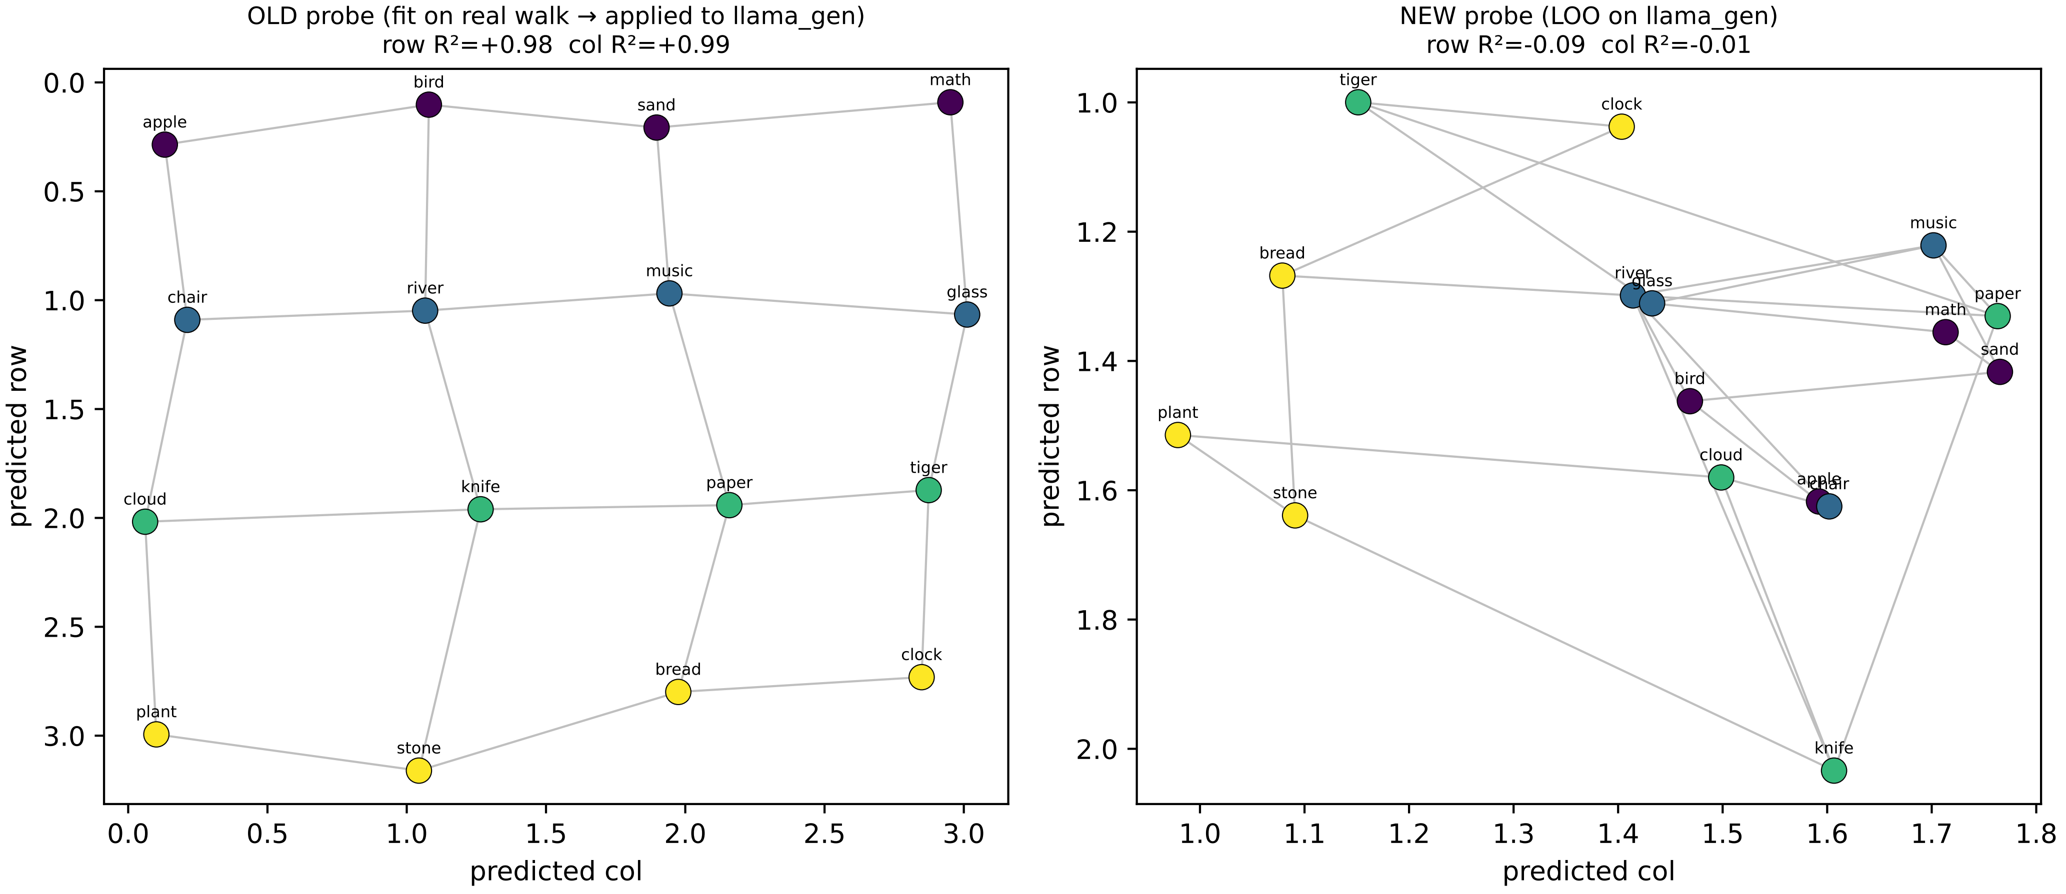

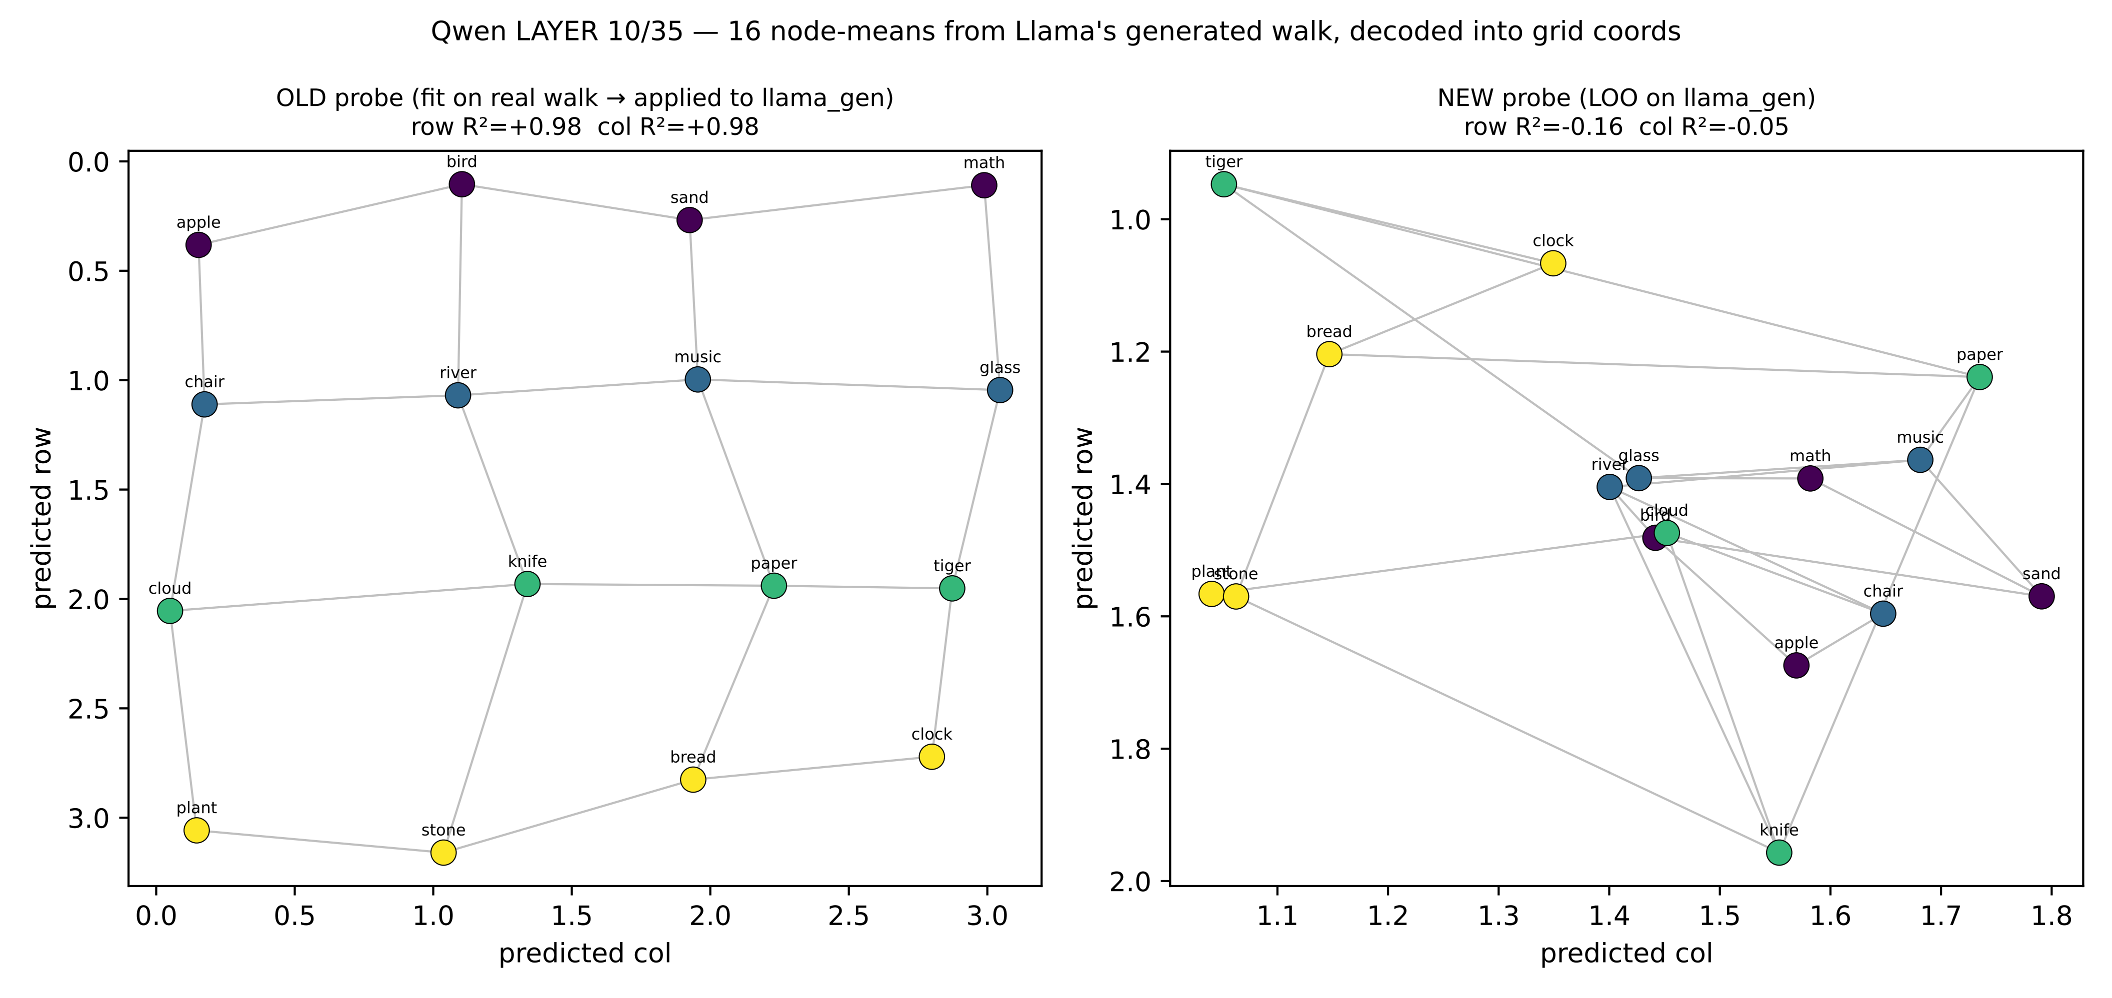
Qwen LAYER 10/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.98 col R²=+0.98
predicted row
predicted col
0.0
0.5
1.0
1.5
2.0
2.5
3.0
bird
math
sand
apple
chair
river
music
glass
knife
paper
tiger
cloud
clock
bread
plant
stone
NEW probe (LOO on llama_gen)
row R²=-0.16 col R²=-0.05
predicted row
predicted col
1.0
1.2
1.4
1.6
1.8
2.0
1.1
1.3
1.5
1.7
tiger
clock
bread
paper
music
glass
math
river
cloud
bird
plant
stone
sand
chair
apple
knife

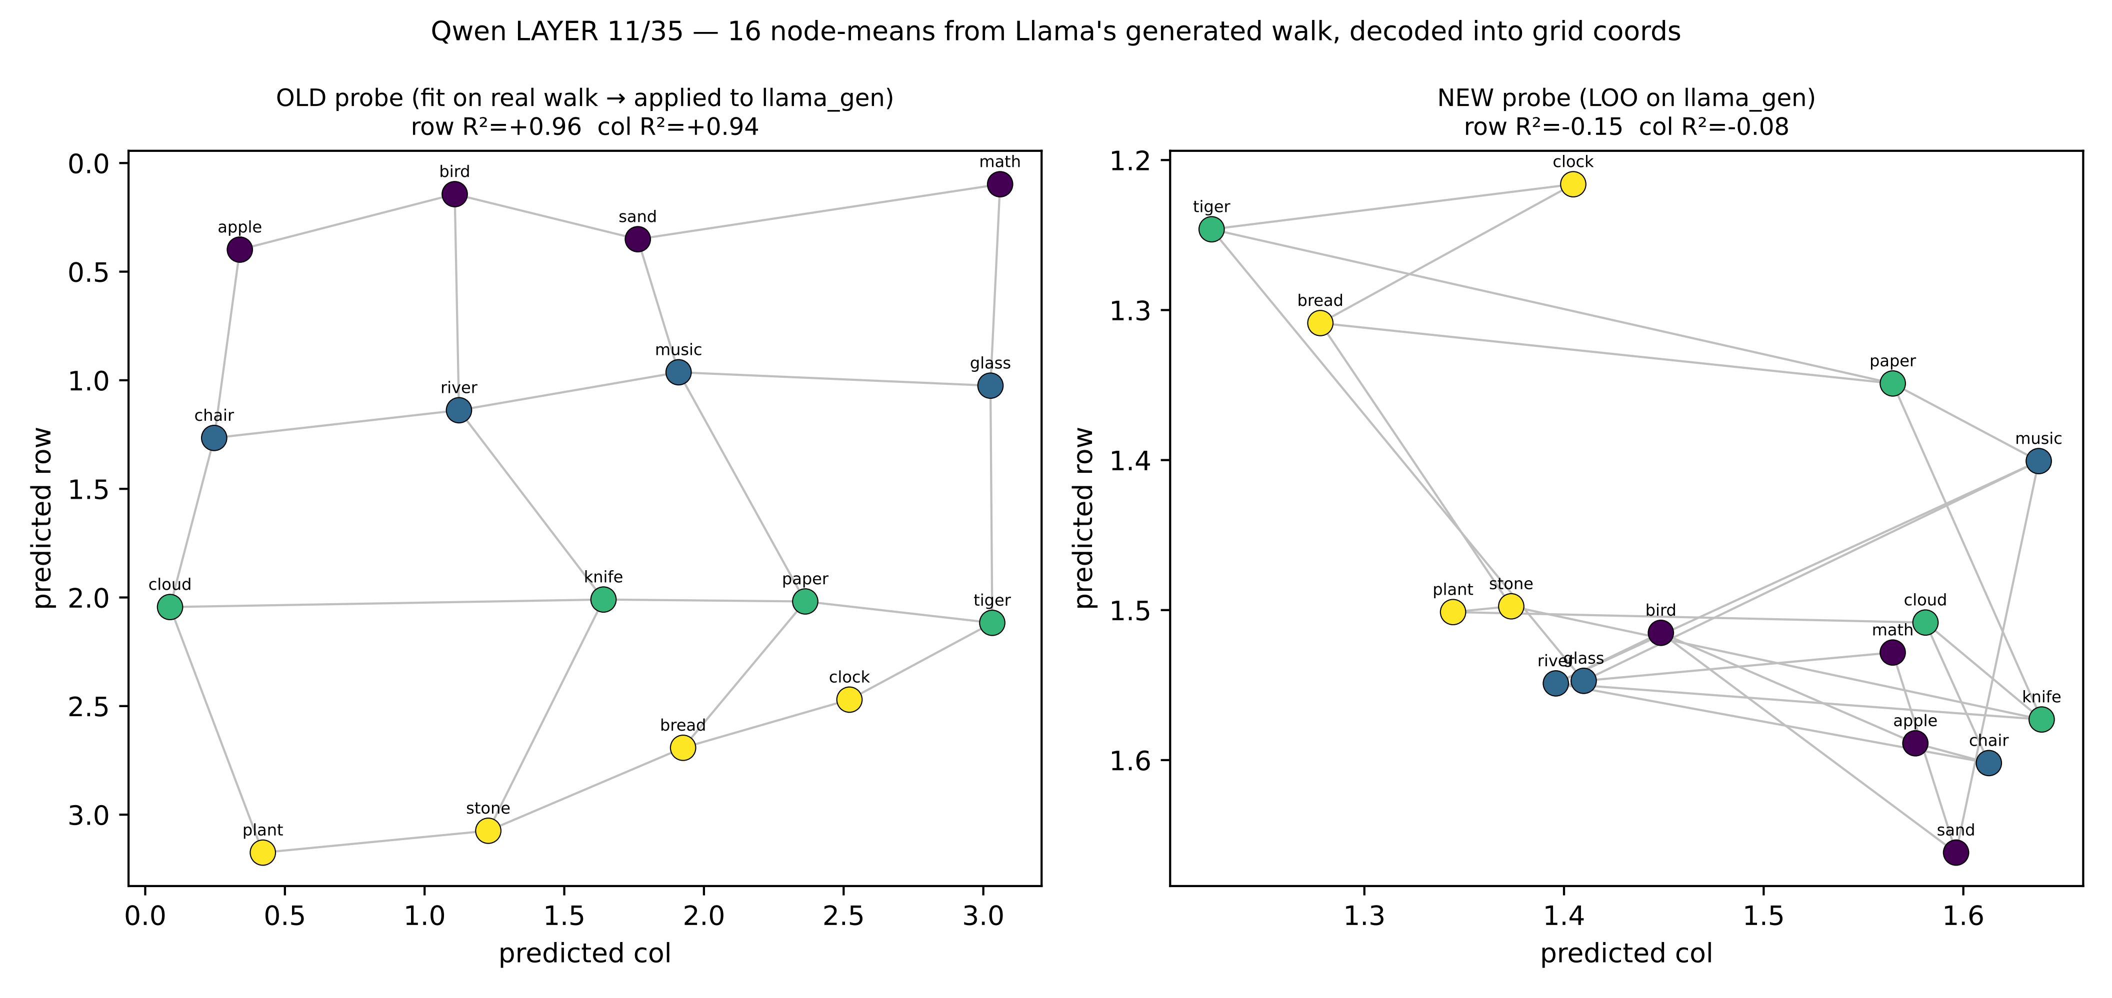
Qwen LAYER 11/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.96 col R²=+0.94
NEW probe (LOO on llama_gen)
row R²=-0.15 col R²=-0.08
predicted row
predicted col
bird
math
apple
sand
music
glass
river
chair
cloud
knife
paper
tiger
clock
bread
stone
plant

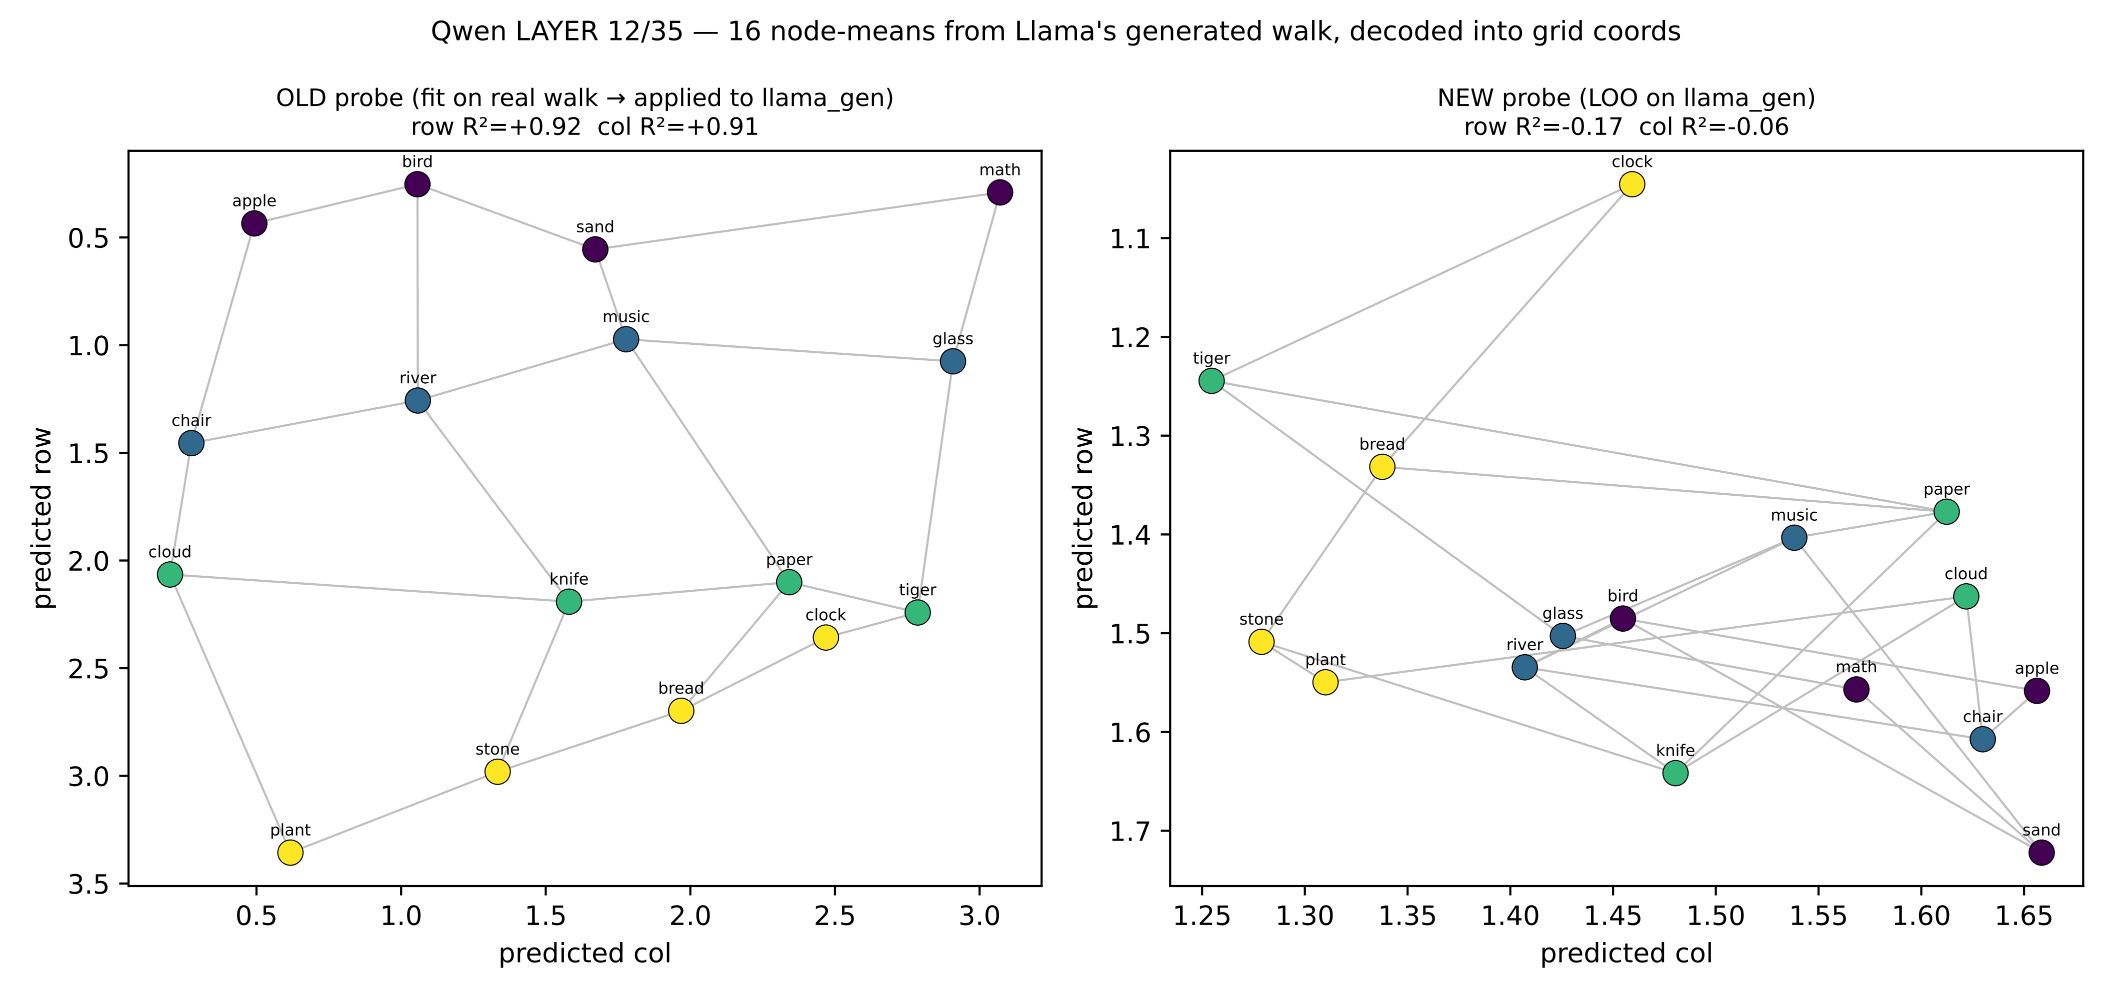

Qwen LAYER 12/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.92 col R²=+0.91
NEW probe (LOO on llama_gen)
row R²=-0.17 col R²=-0.06
predicted row
predicted col
bird
apple
math
sand
music
glass
river
chair
cloud
knife
paper
tiger
clock
bread
stone
plant

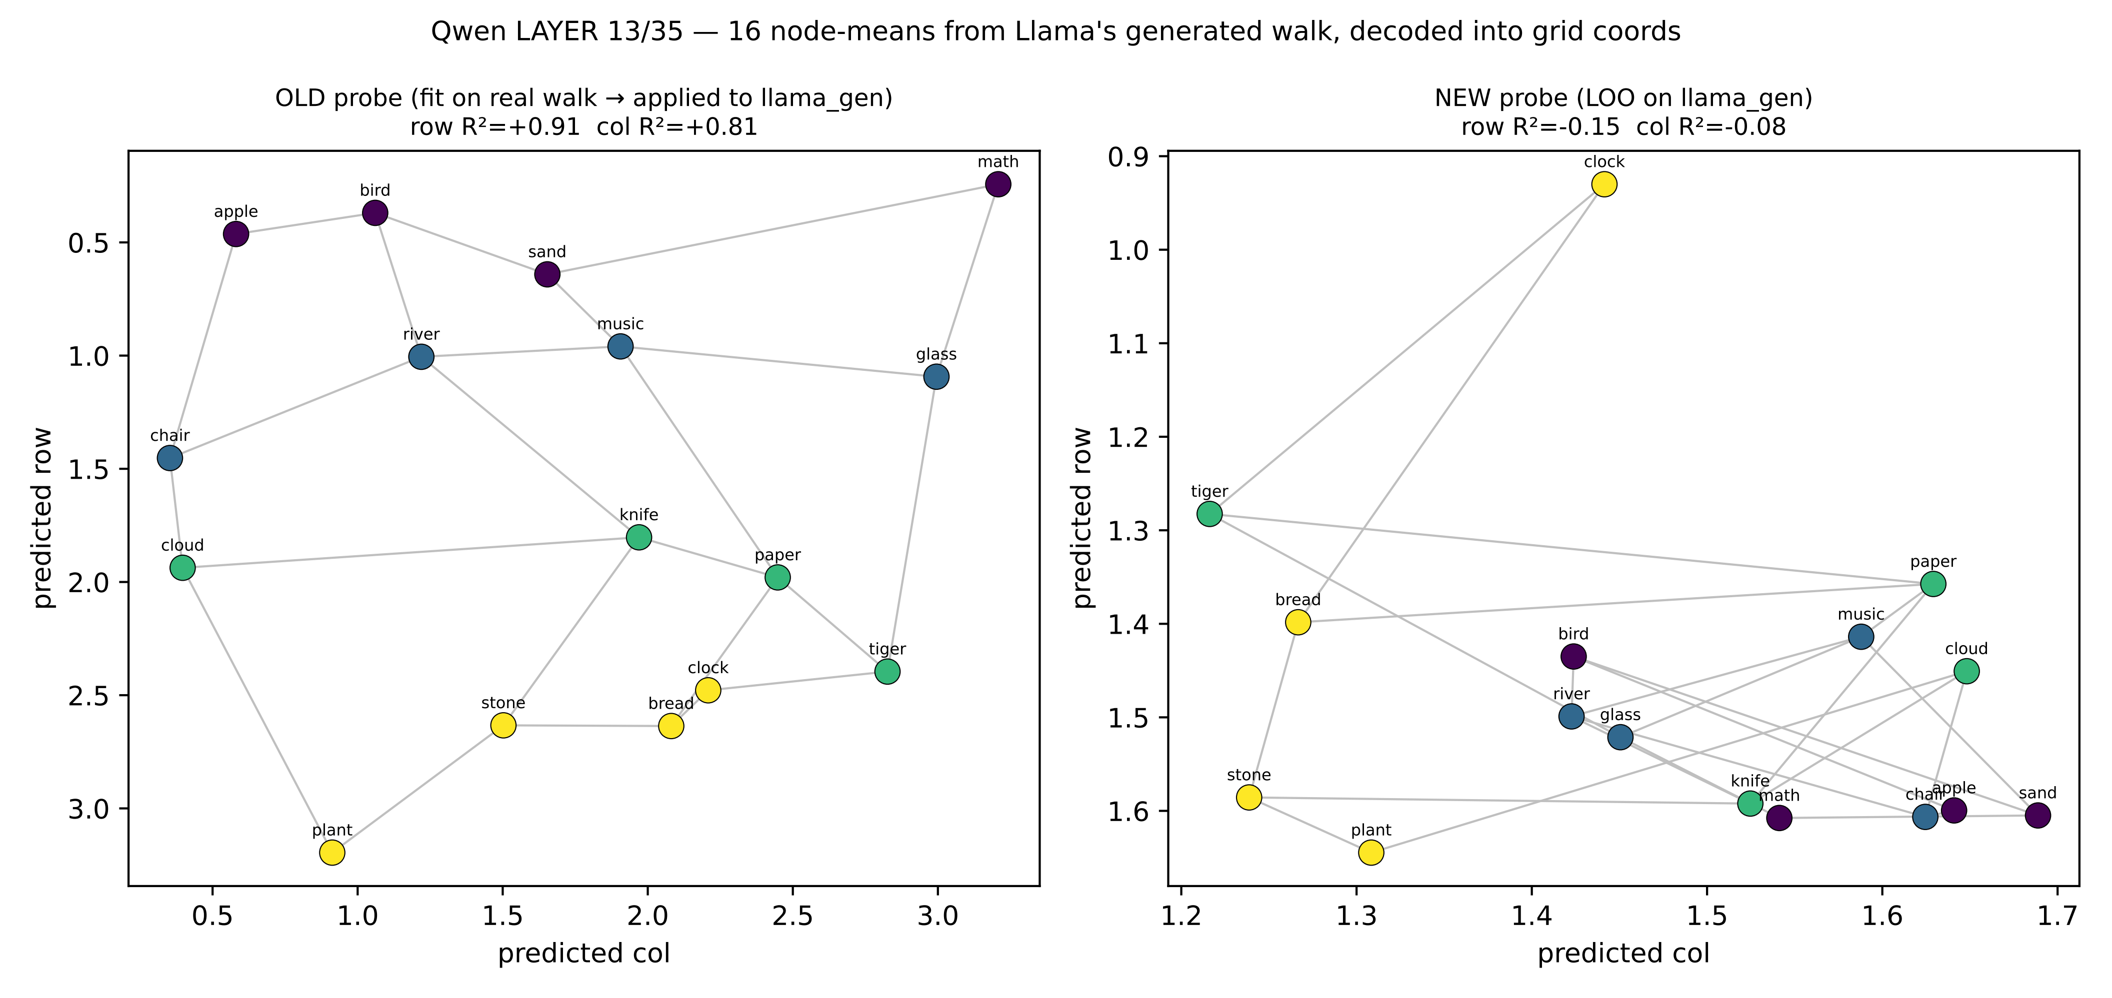
Qwen LAYER 13/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.91 col R²=+0.81
NEW probe (LOO on llama_gen)
row R²=-0.15 col R²=-0.08
predicted row
predicted col
math
bird
apple
sand
music
river
glass
chair
knife
cloud
paper
tiger
clock
stone
bread
plant
0.5
1.0
1.5
2.0
2.5
3.0
0.9
1.1
1.2
1.3
1.4
1.6
1.7

Qwen LAYER 14/35 — 16 node-means from Llama's generated walk, decoded into grid coords

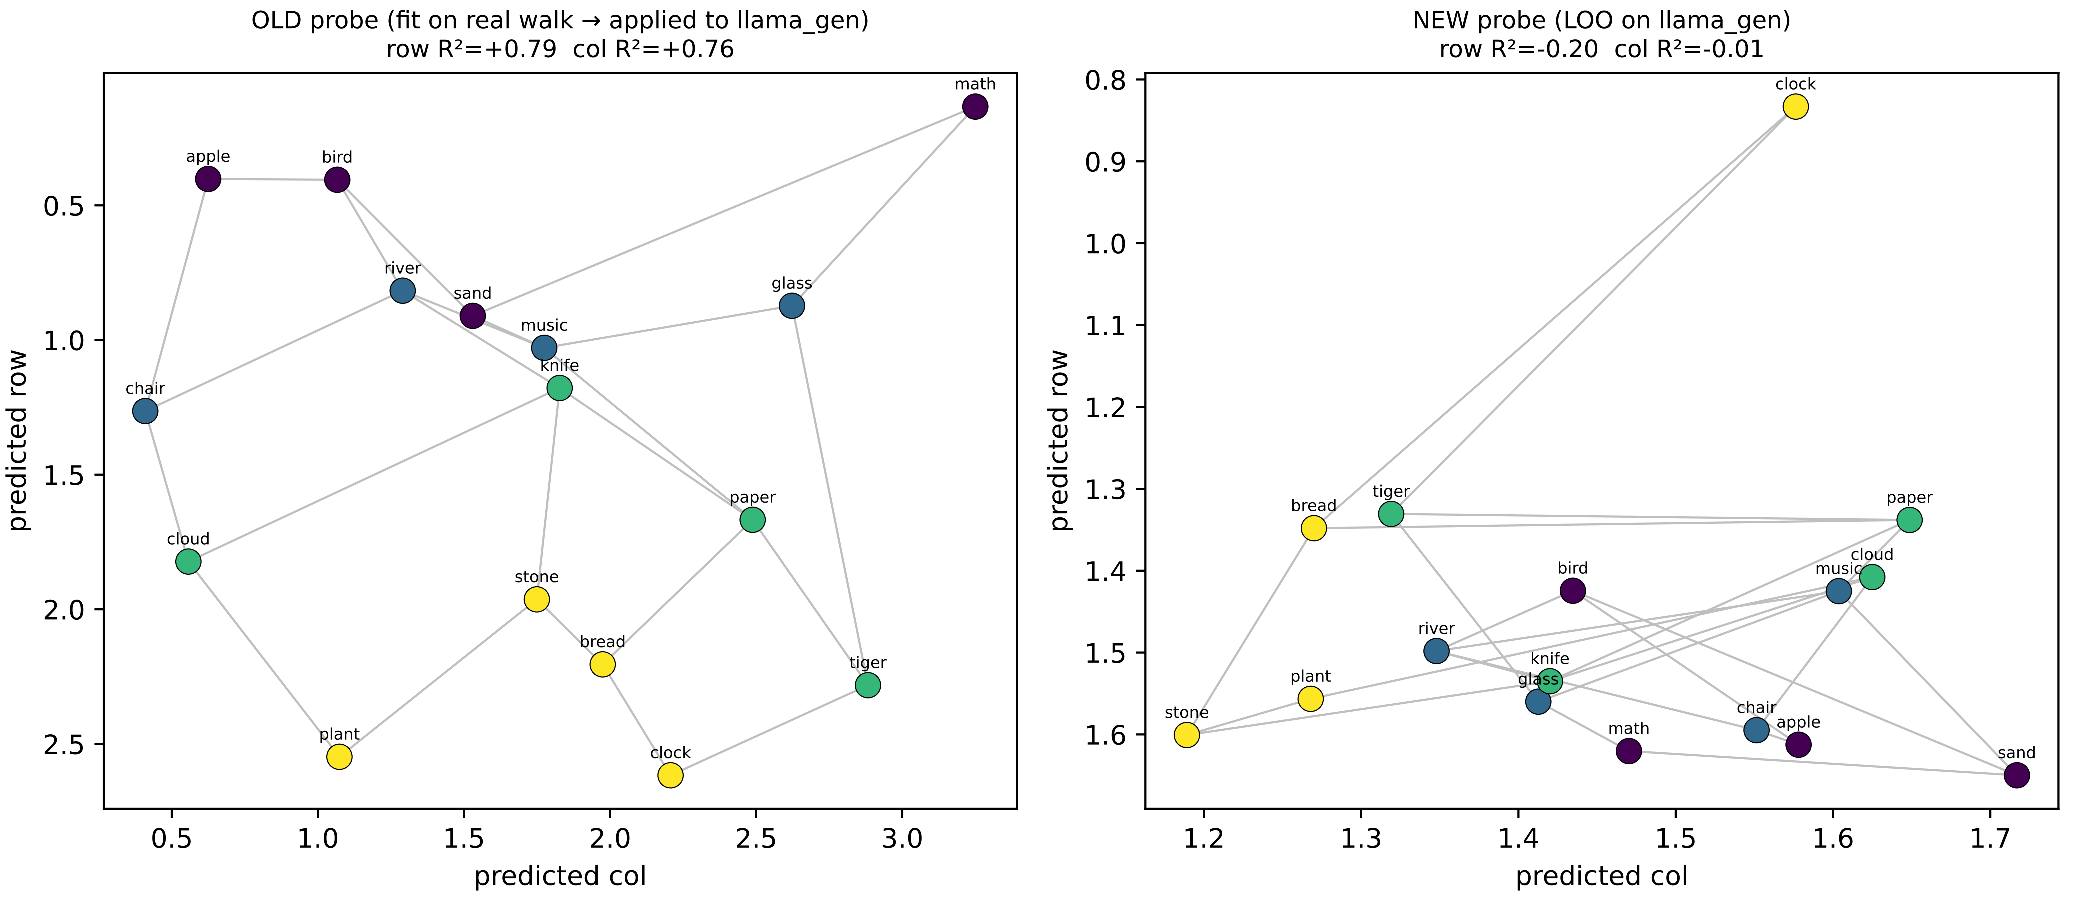

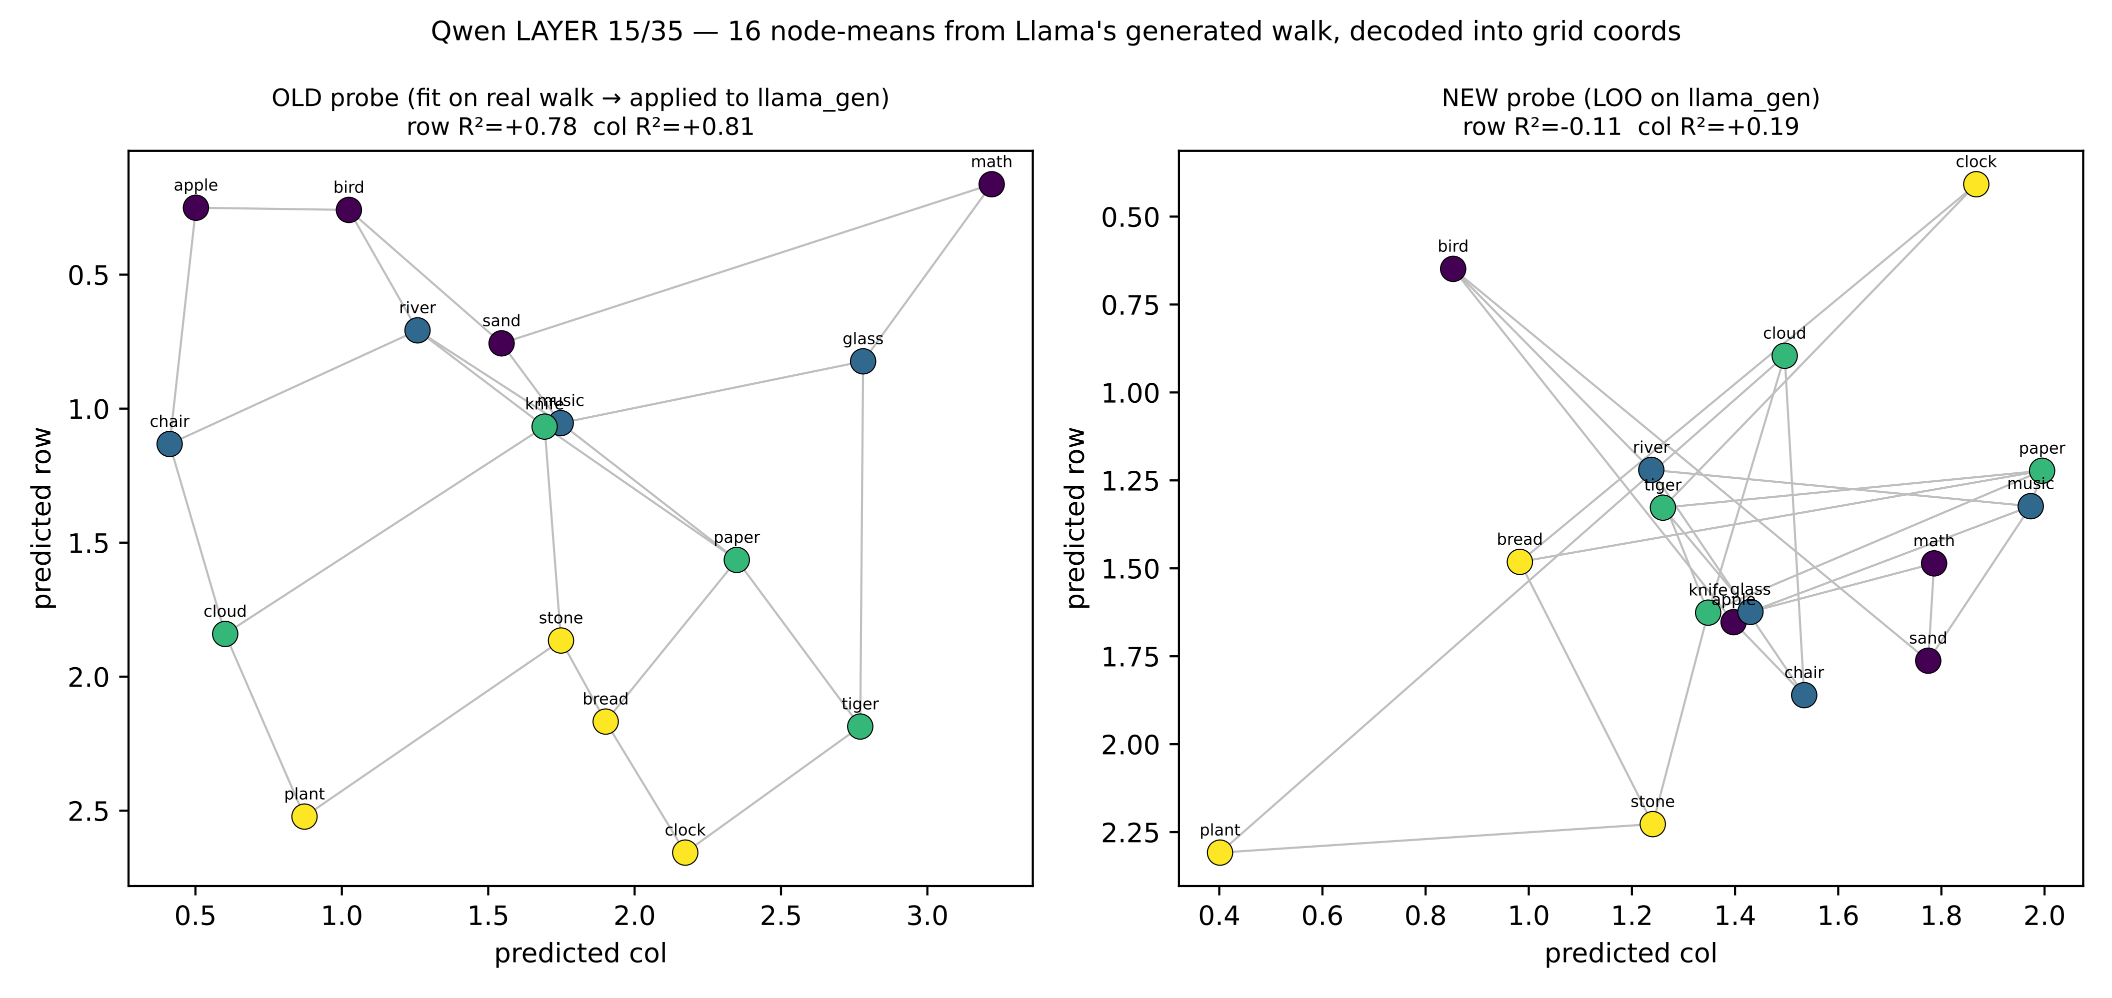
Qwen LAYER 15/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.78 col R²=+0.81
NEW probe (LOO on llama_gen)
row R²=-0.11 col R²=+0.19
predicted row
predicted col
apple
bird
math
river
sand
glass
chair
knife
music
paper
cloud
stone
bread
tiger
plant
clock

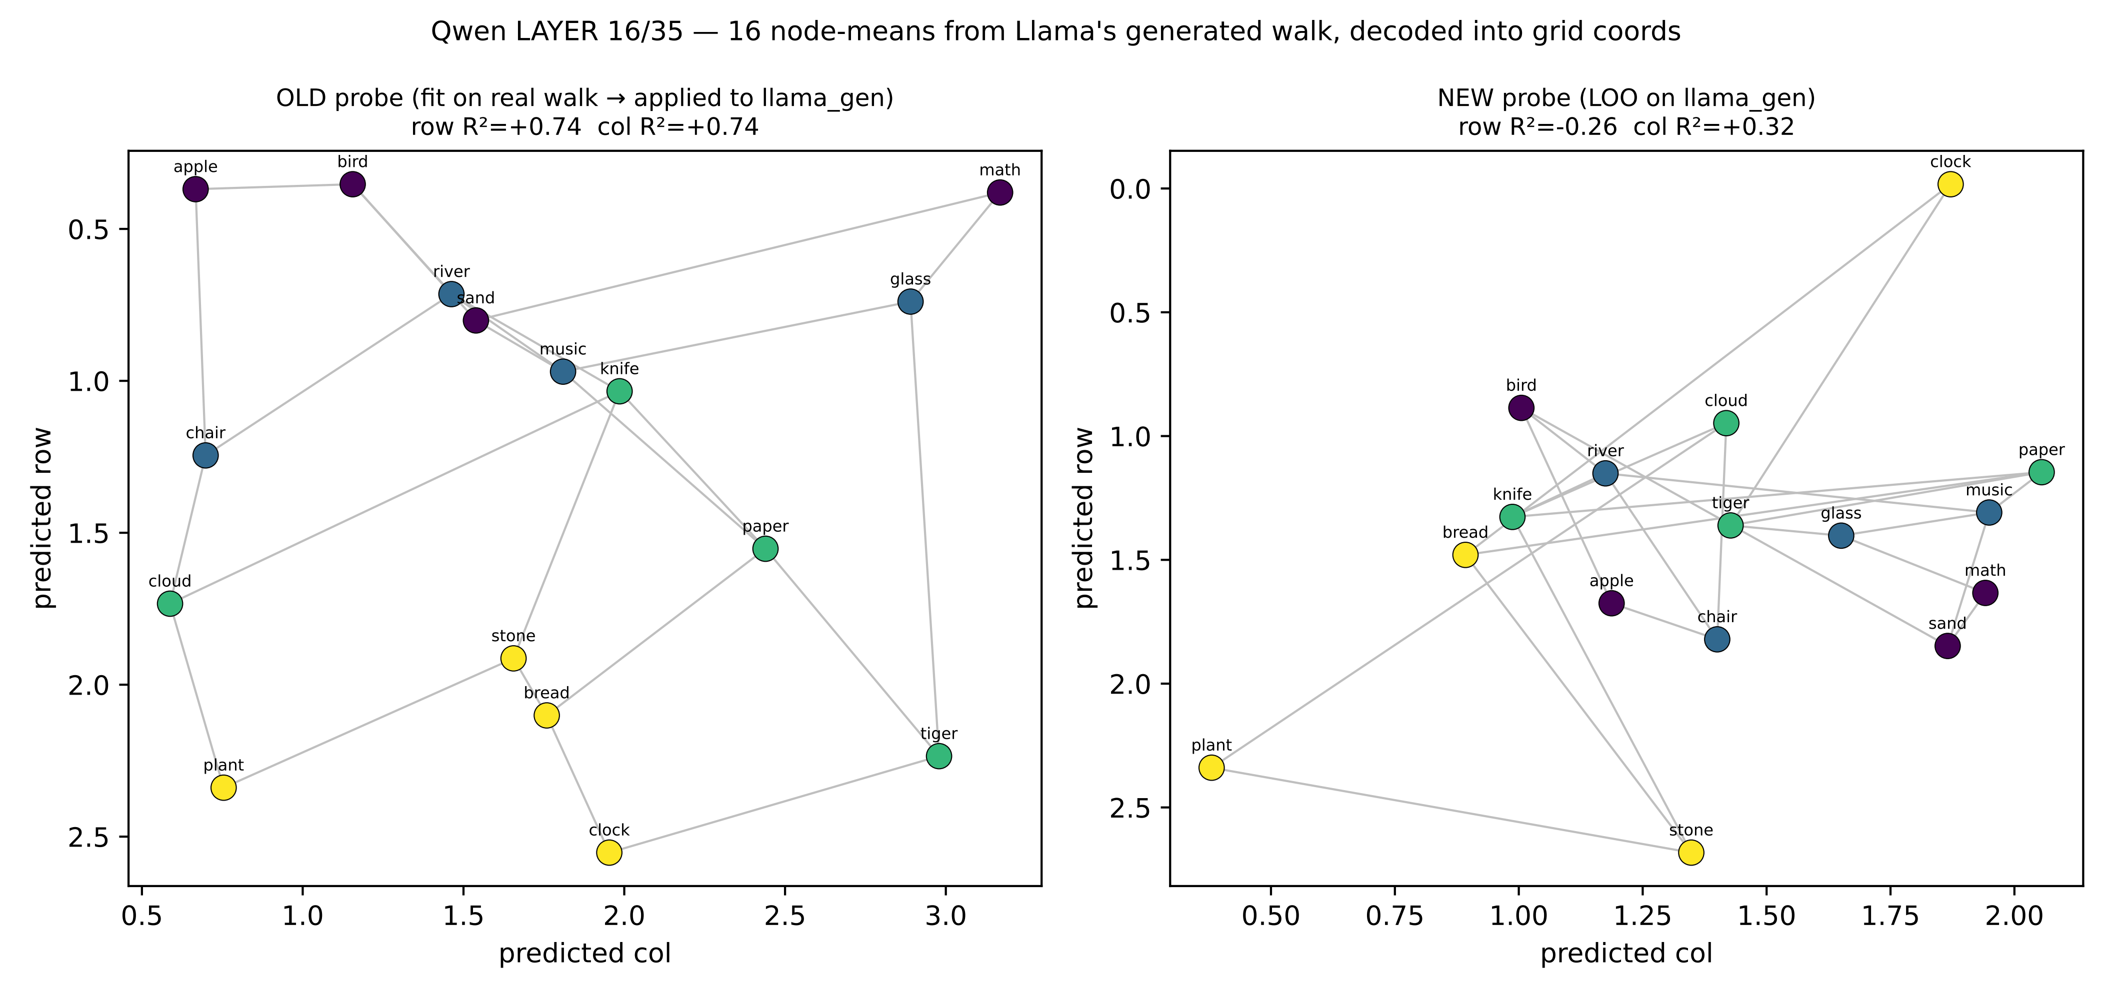

Qwen LAYER 16/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.74 col R²=+0.74
NEW probe (LOO on llama_gen)
row R²=-0.26 col R²=+0.32
predicted row
predicted col
apple
bird
math
river
sand
glass
music
knife
chair
paper
cloud
stone
bread
tiger
plant
clock

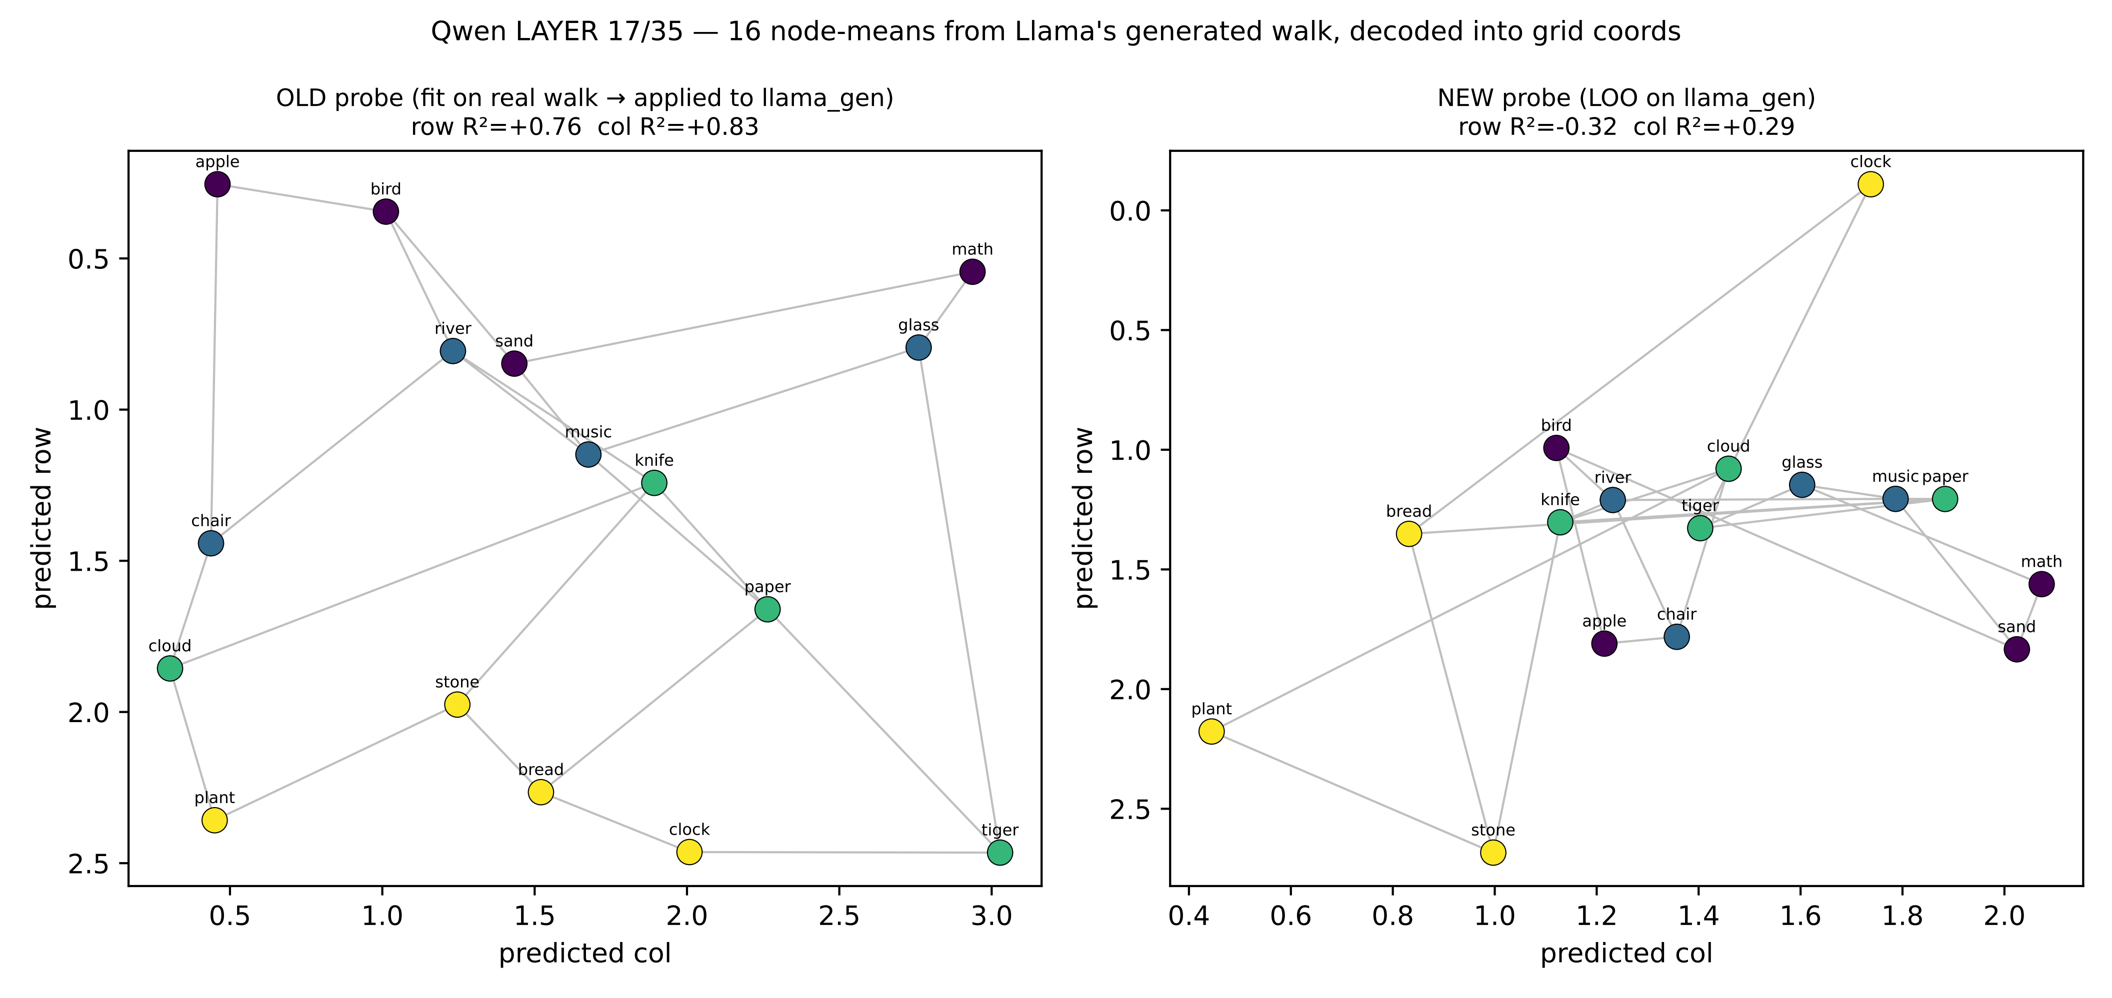

Qwen LAYER 17/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.76 col R²=+0.83
apple
bird
math
river
sand
glass
music
knife
chair
paper
cloud
stone
bread
plant
clock
tiger
predicted row
predicted col
0.5
1.0
1.5
2.0
2.5
3.0
NEW probe (LOO on llama_gen)
row R²=-0.32 col R²=+0.29
clock
bird
cloud
river
glass
music
paper
knife
tiger
bread
math
apple
chair
sand
plant
stone
predicted row
predicted col
0.0
0.5
1.0
1.5
2.0
2.5
0.4
0.6
0.8
1.0
1.2
1.4
1.6
1.8
2.0

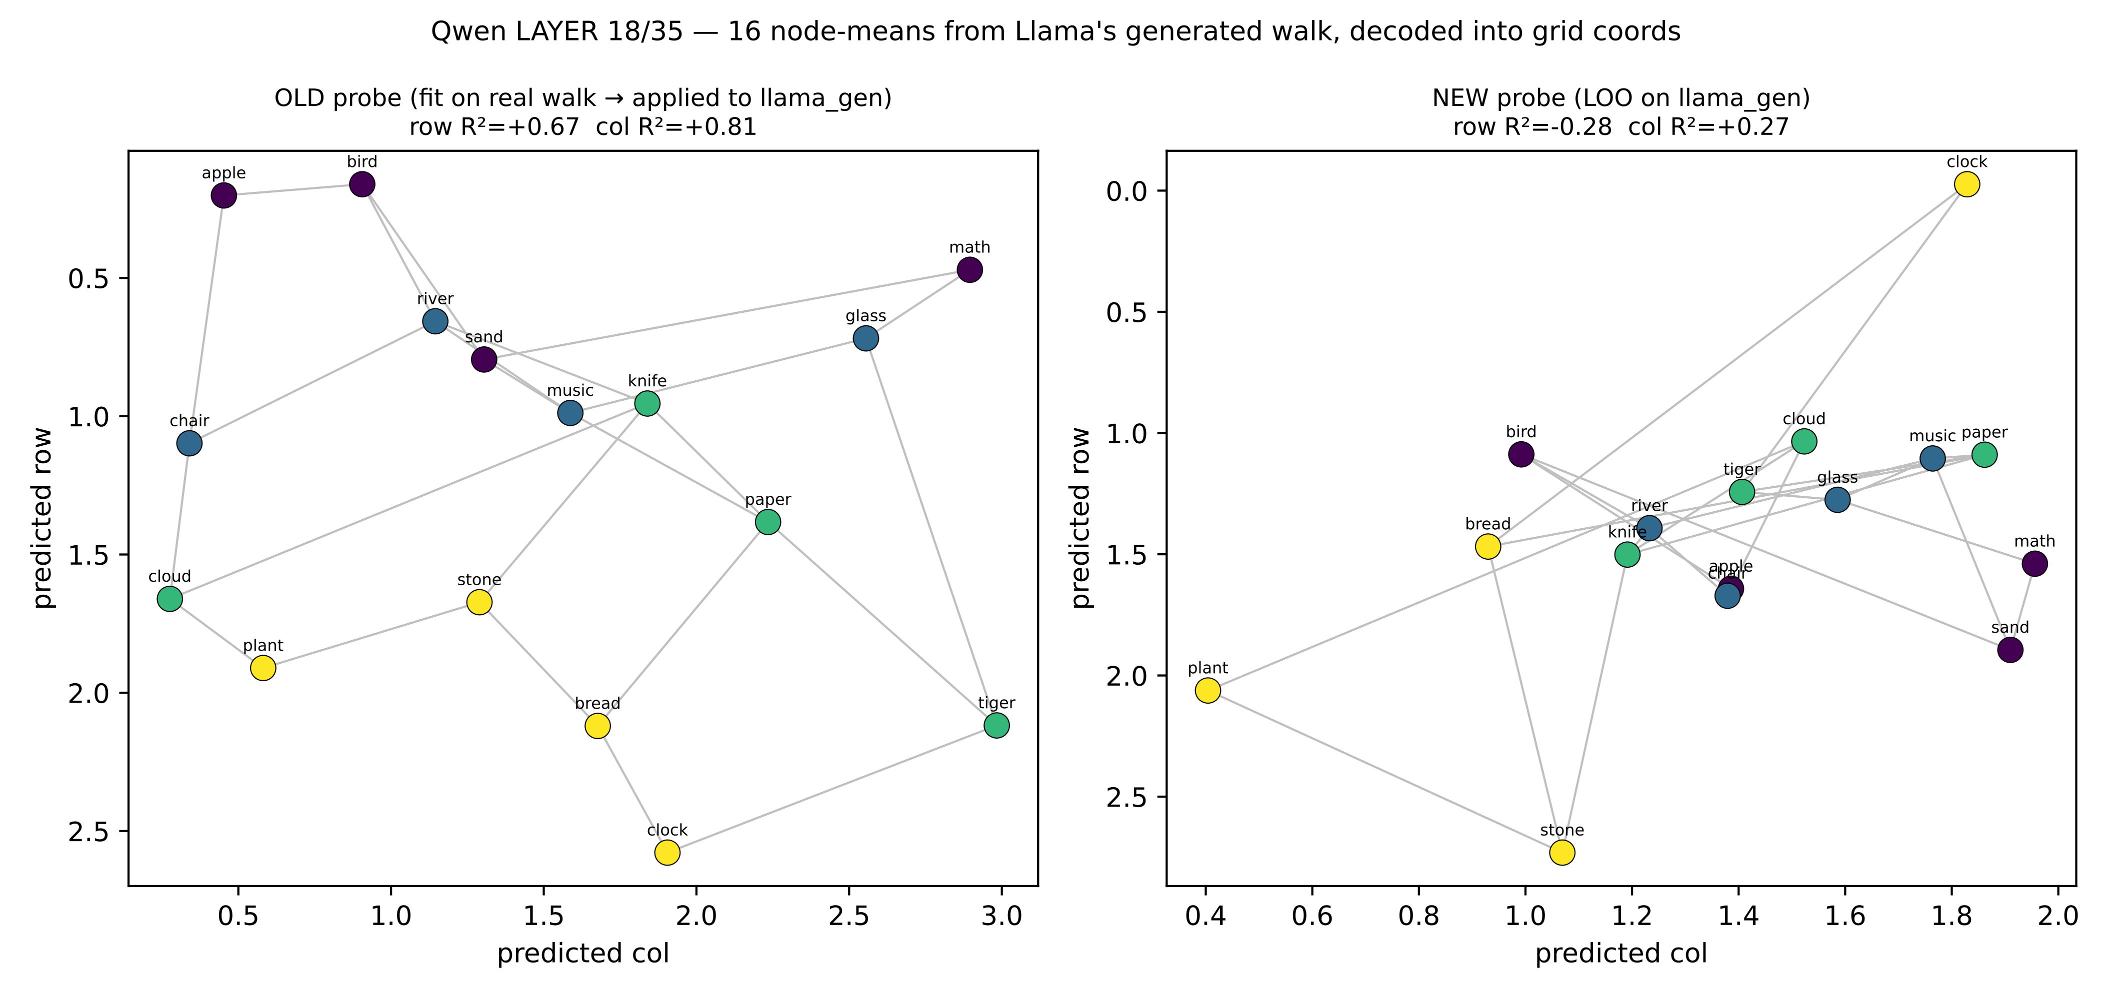

Qwen LAYER 18/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.67 col R²=+0.81
NEW probe (LOO on llama_gen)
row R²=-0.28 col R²=+0.27
predicted row
predicted col
apple
bird
math
river
sand
glass
knife
music
chair
paper
cloud
stone
plant
bread
tiger
clock

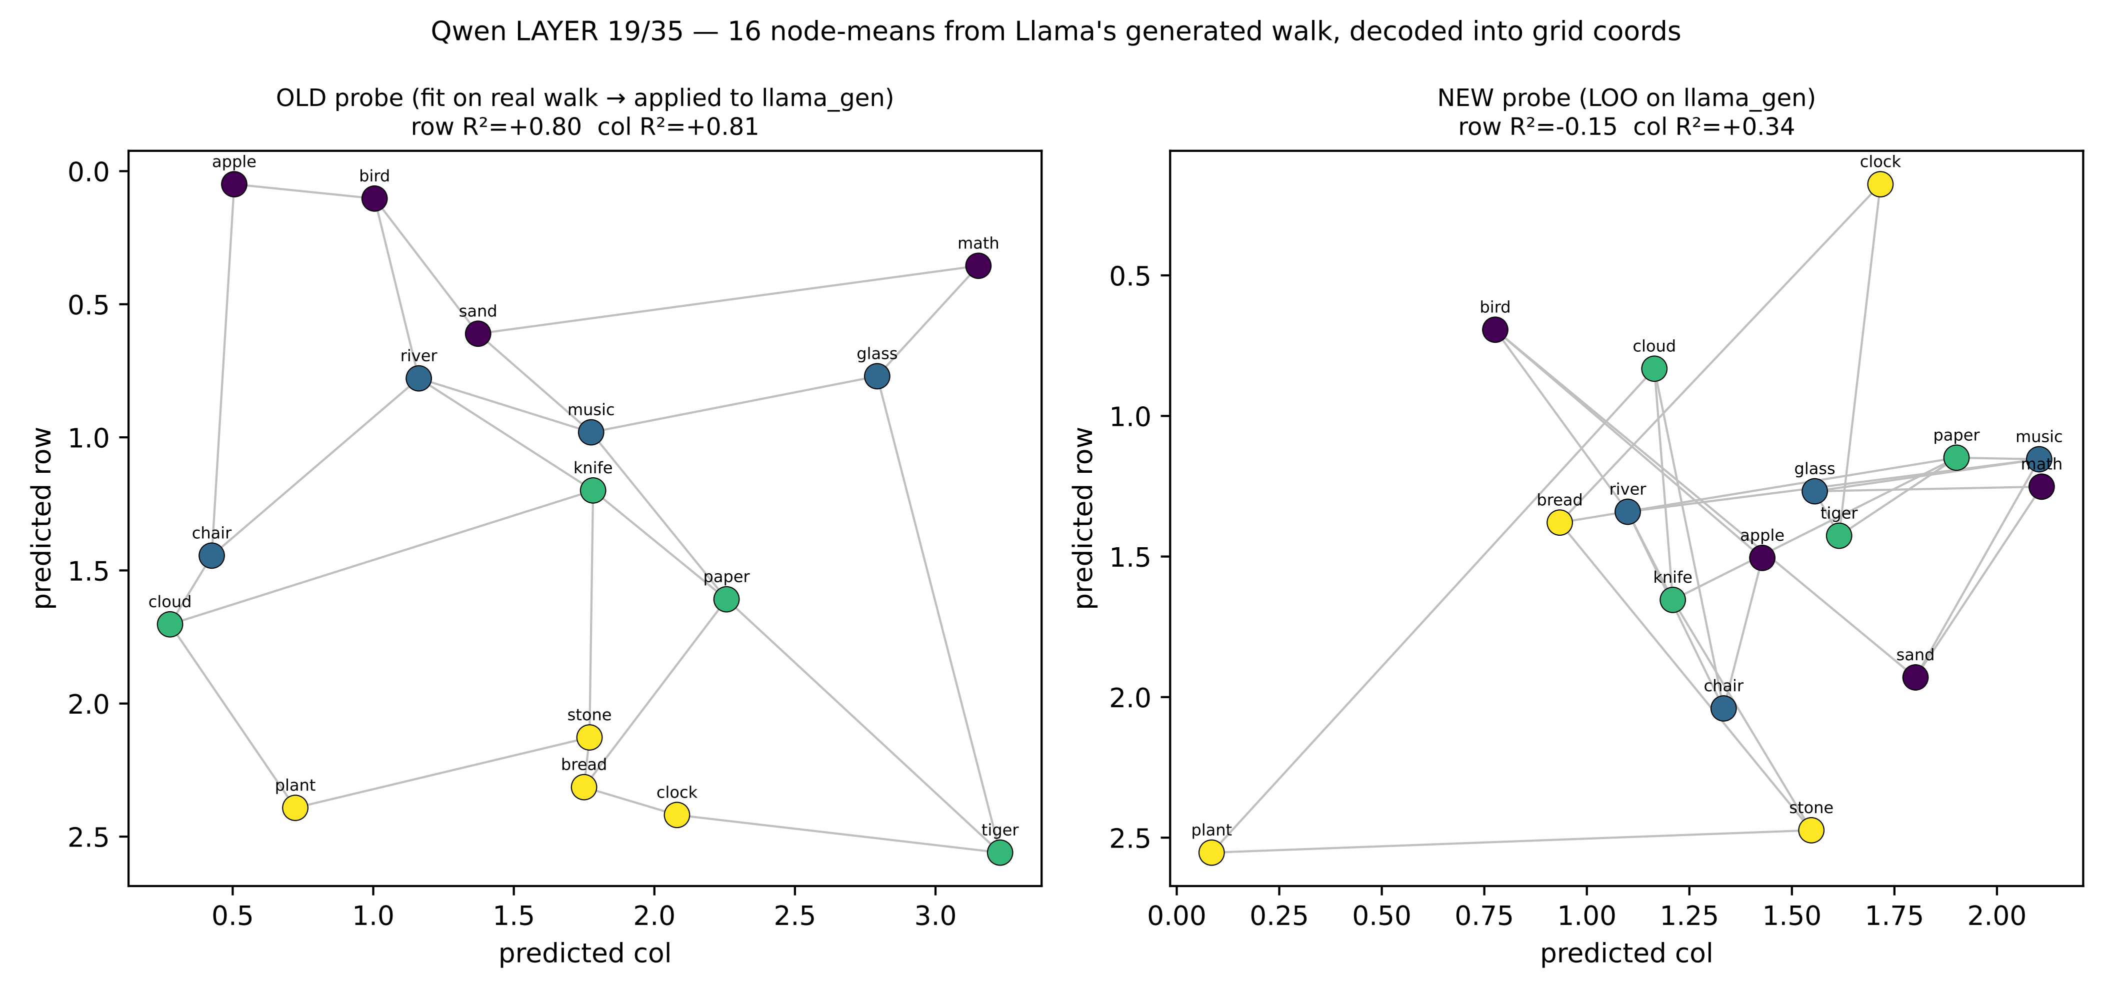
Qwen LAYER 19/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.80 col R²=+0.81
predicted row
predicted col
apple
bird
math
sand
river
glass
music
knife
chair
cloud
paper
stone
bread
clock
plant
tiger
NEW probe (LOO on llama_gen)
row R²=-0.15 col R²=+0.34
predicted row
predicted col
clock
bird
cloud
paper
music
math
glass
bread
river
tiger
apple
knife
sand
chair
stone
plant

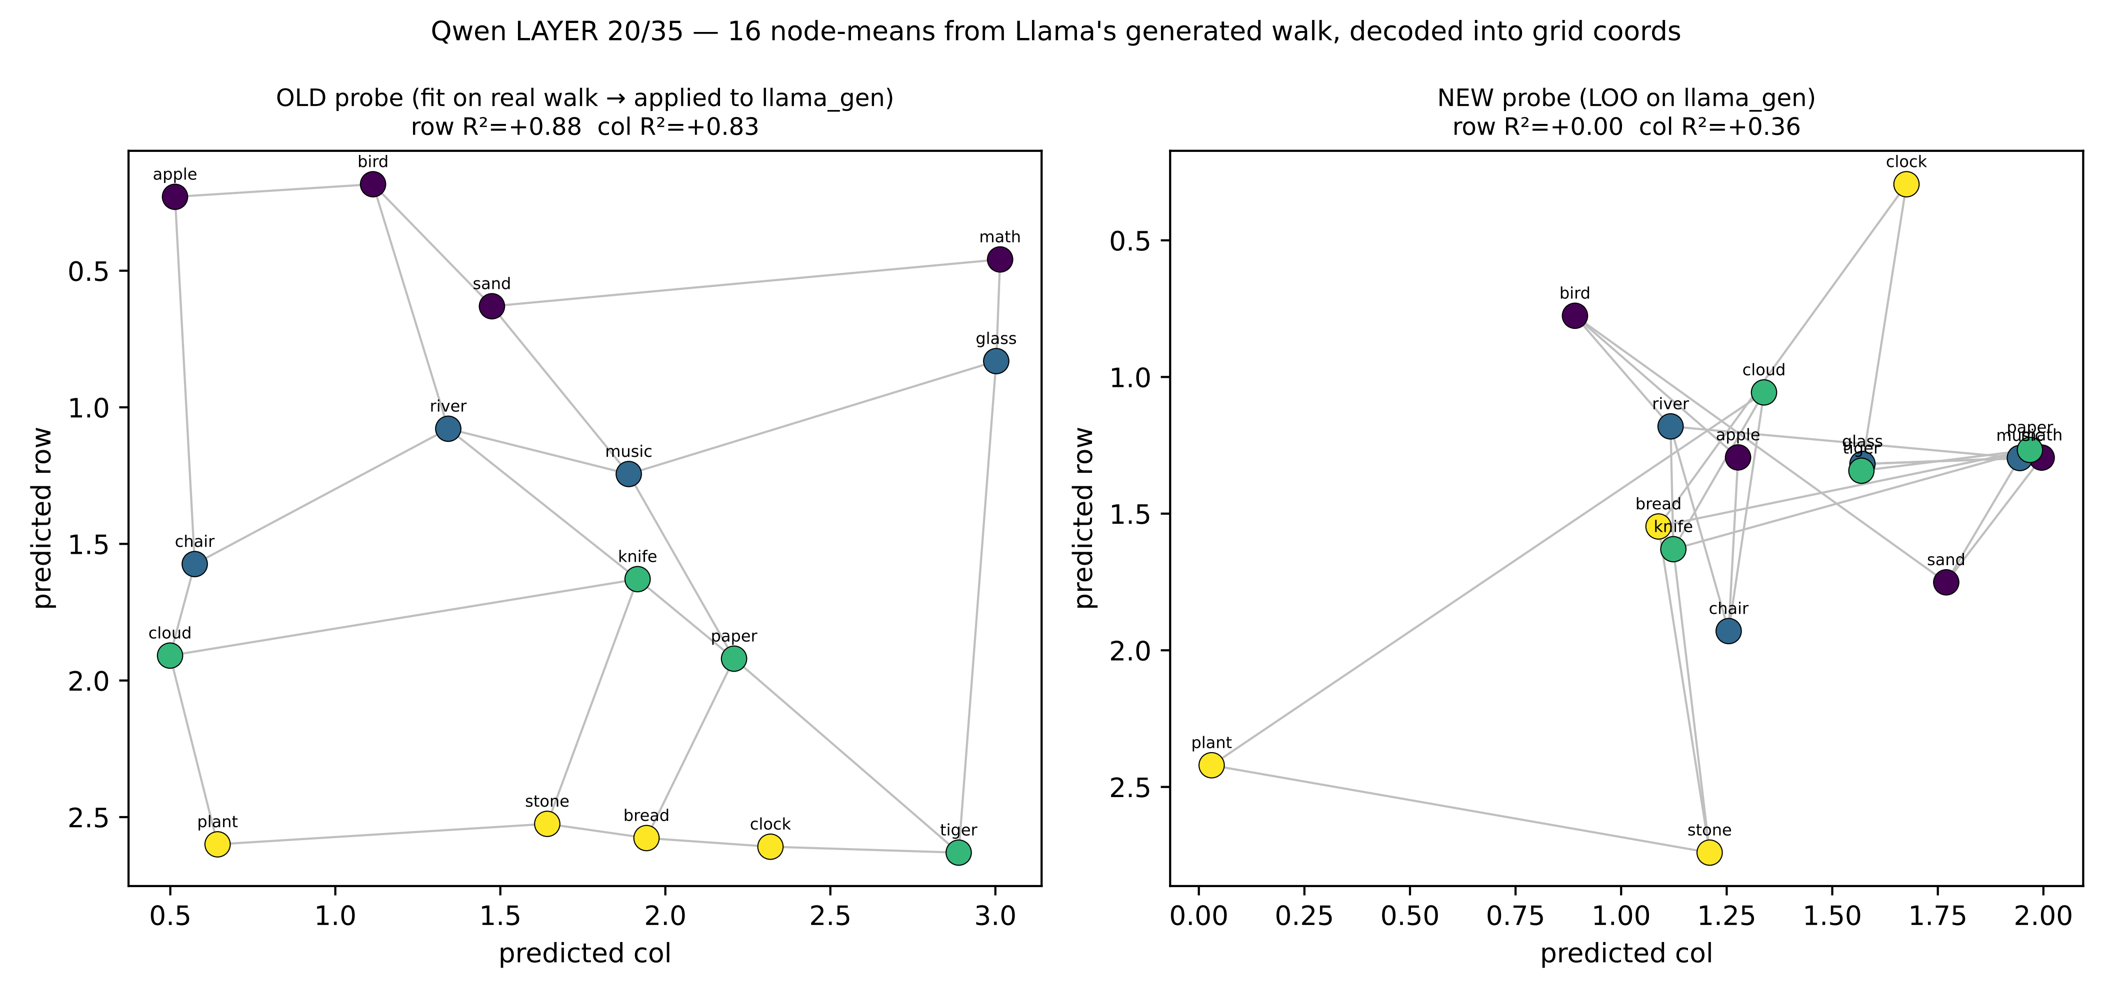
Qwen LAYER 20/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.88 col R²=+0.83
NEW probe (LOO on llama_gen)
row R²=+0.00 col R²=+0.36
predicted row
predicted col
apple
bird
math
sand
glass
river
music
chair
knife
cloud
paper
stone
bread
clock
plant
tiger
glass
tiger
paper
music
math
sand
bread
knife
chair
stone
plant

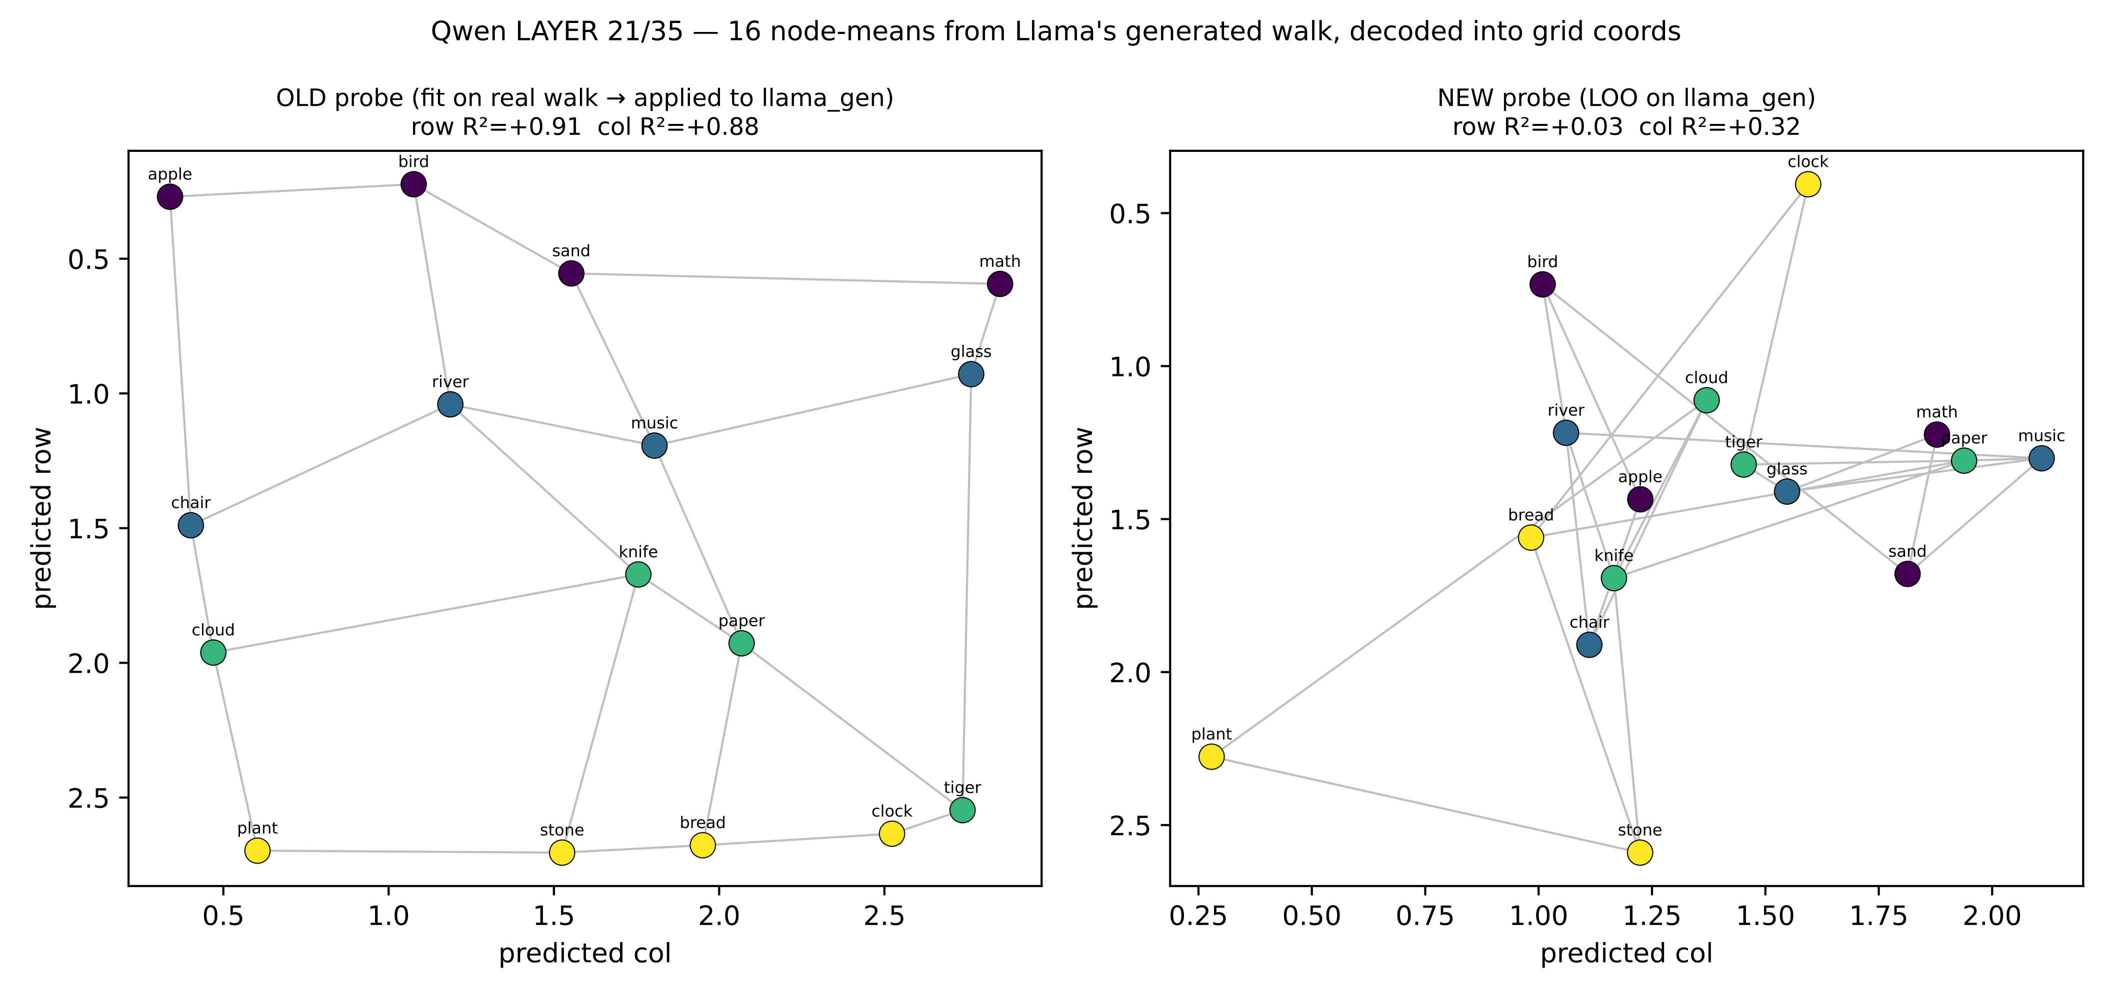

Qwen LAYER 21/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.91 col R²=+0.88
NEW probe (LOO on llama_gen)
row R²=+0.03 col R²=+0.32
predicted row
predicted col
apple
bird
sand
math
glass
river
music
chair
knife
cloud
paper
tiger
plant
stone
bread
clock

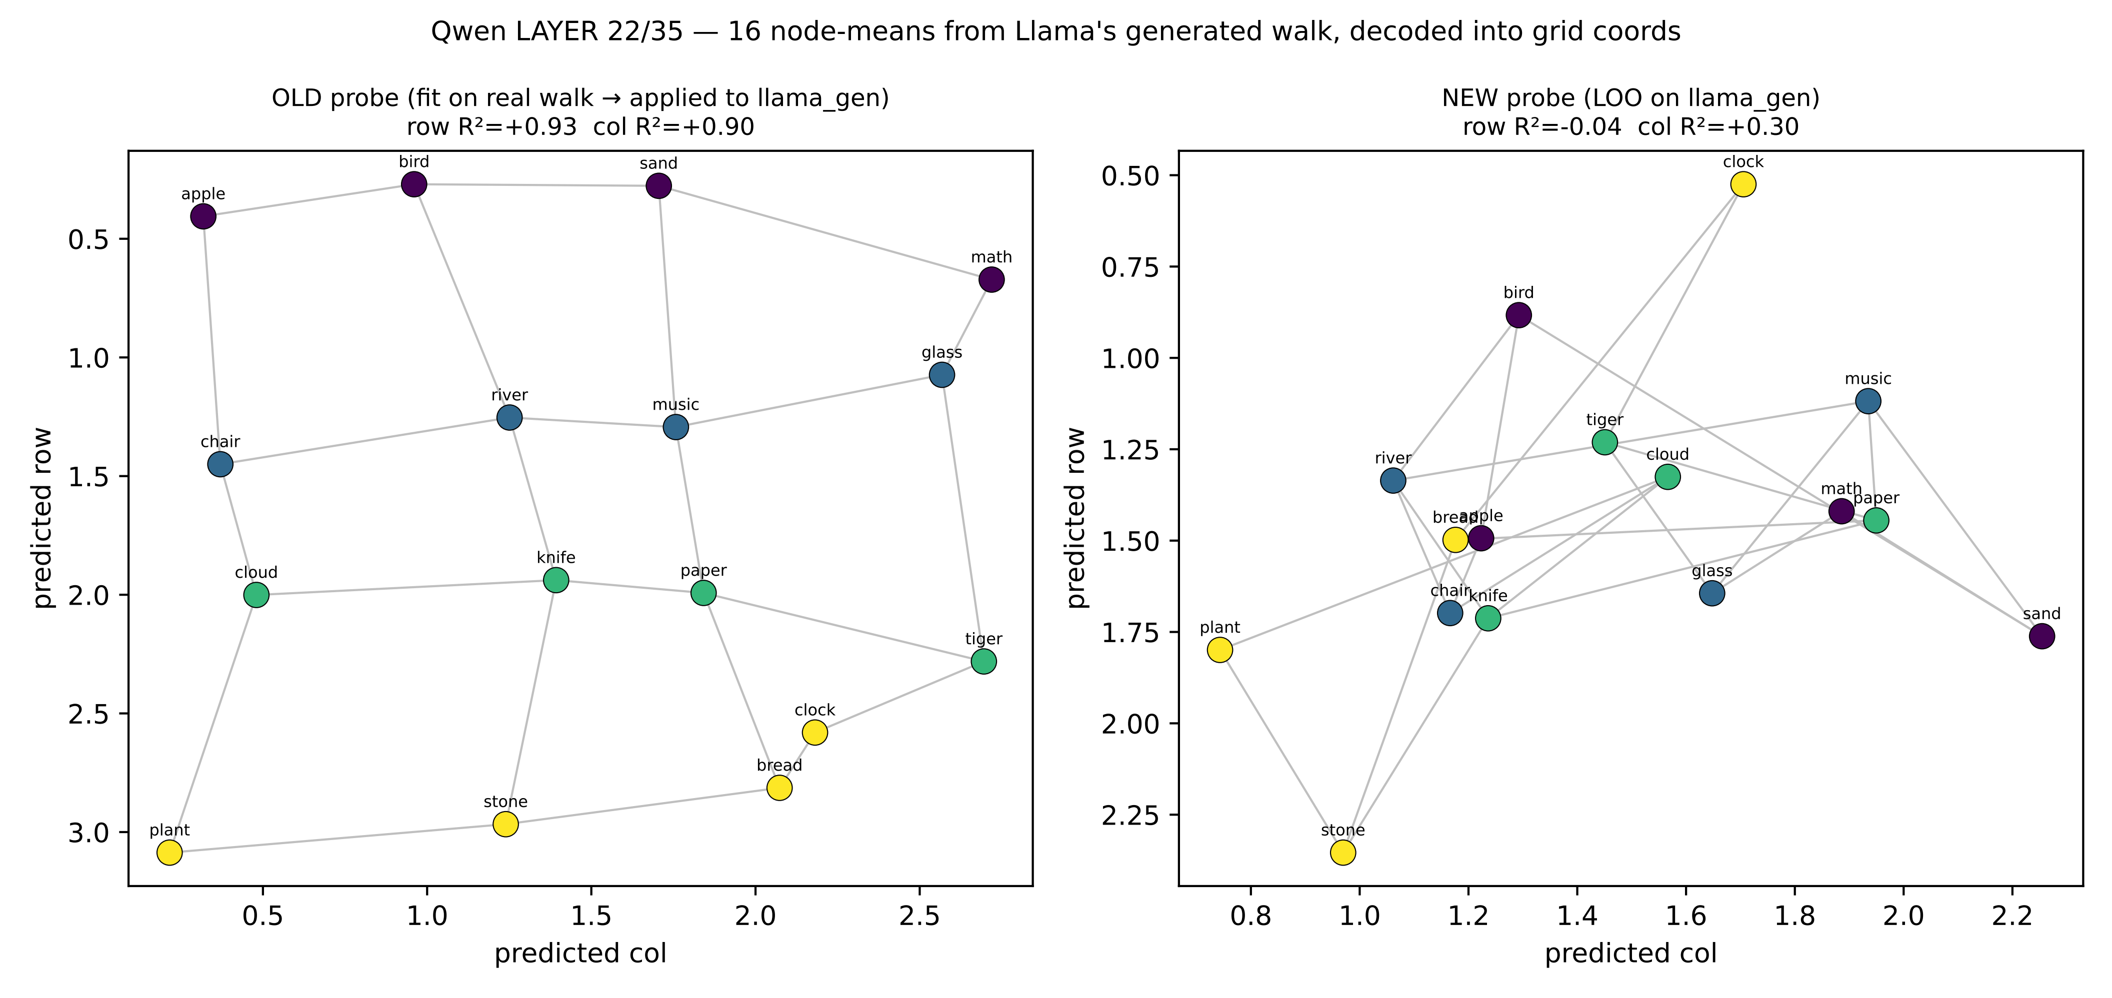

Qwen LAYER 22/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.93 col R²=+0.90
NEW probe (LOO on llama_gen)
row R²=-0.04 col R²=+0.30
predicted row
predicted col
apple
bird
sand
math
glass
river
music
chair
cloud
knife
paper
tiger
clock
bread
stone
plant

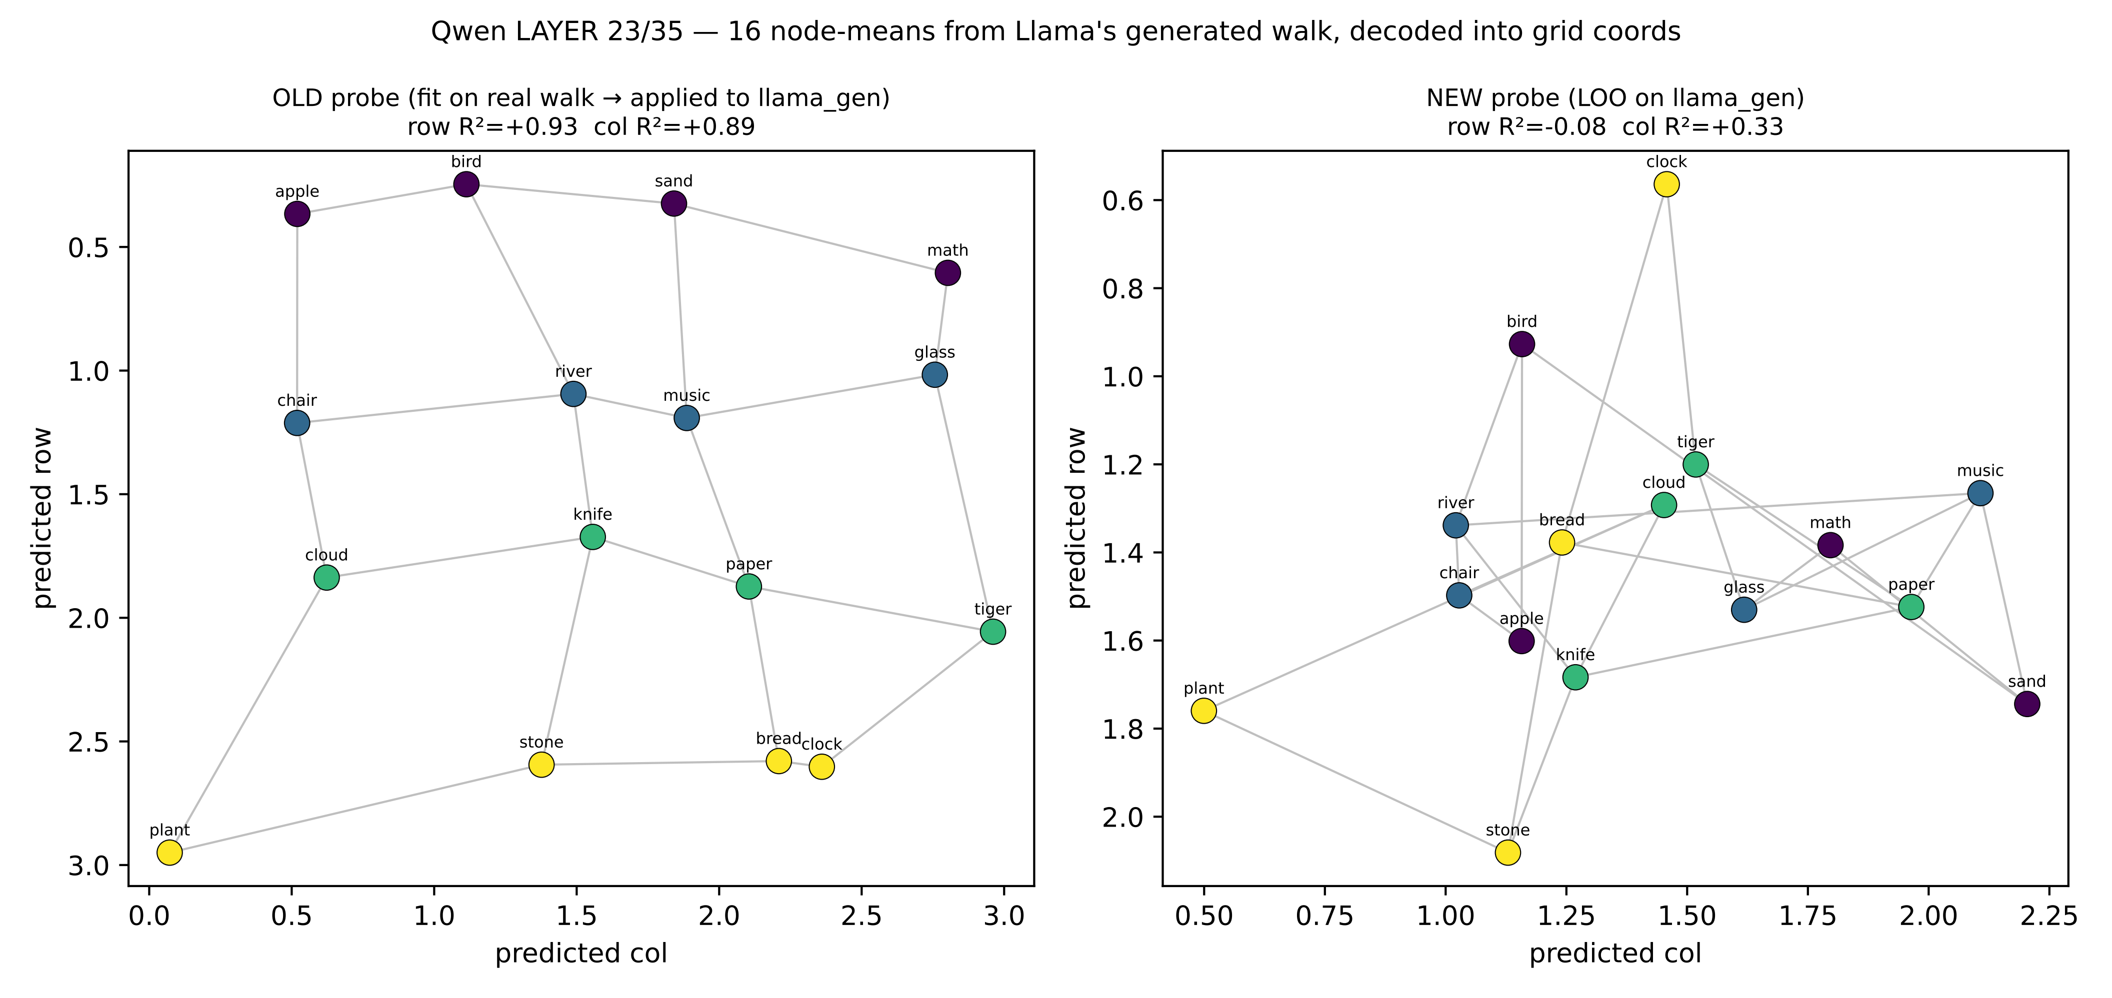

Qwen LAYER 23/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.93 col R²=+0.89
predicted row
predicted col
apple
bird
sand
math
glass
river
music
chair
knife
cloud
paper
tiger
stone
bread
clock
plant
NEW probe (LOO on llama_gen)
row R²=-0.08 col R²=+0.33
predicted row
predicted col
clock
bird
tiger
music
cloud
river
bread
math
chair
glass
paper
apple
knife
sand
plant
stone

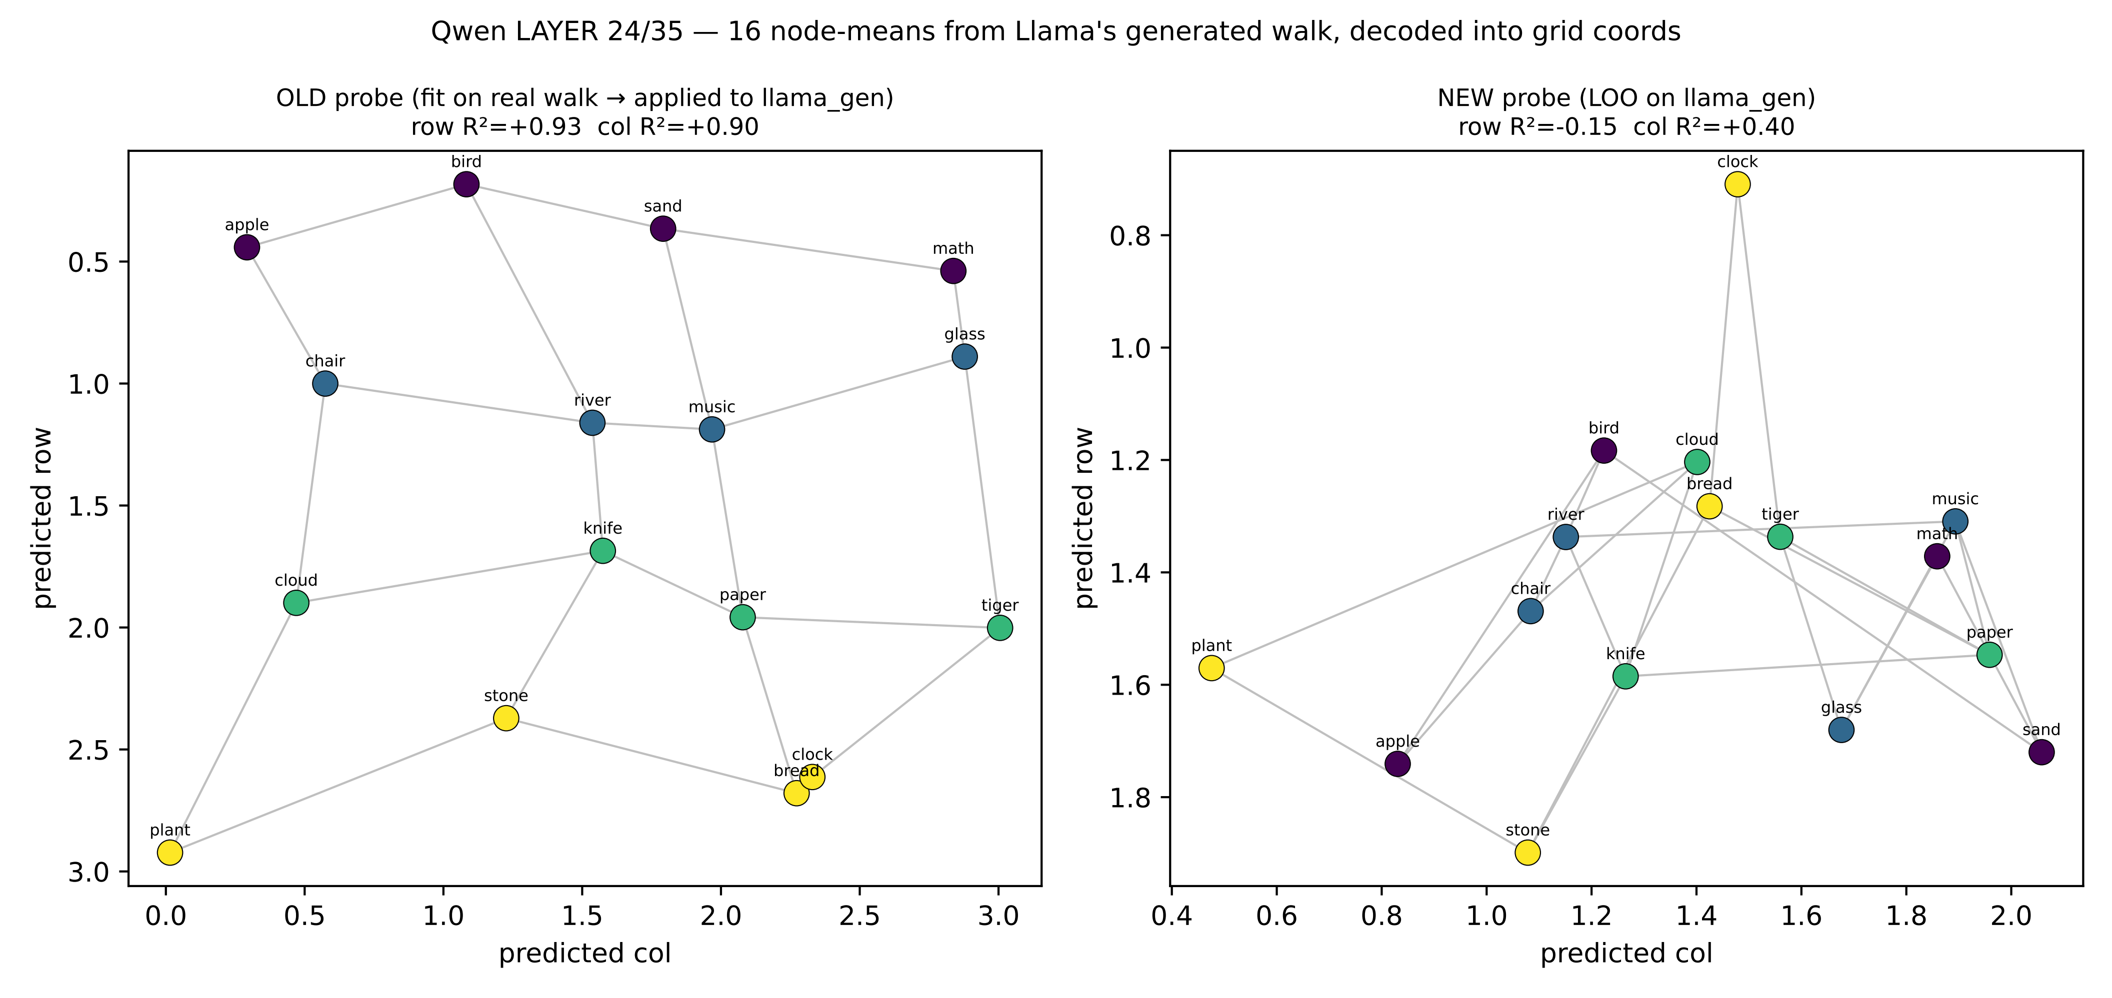

Qwen LAYER 24/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.93 col R²=+0.90
NEW probe (LOO on llama_gen)
row R²=-0.15 col R²=+0.40
predicted row
predicted col
bird
apple
sand
math
glass
chair
river
music
knife
cloud
paper
tiger
stone
clock
bread
plant

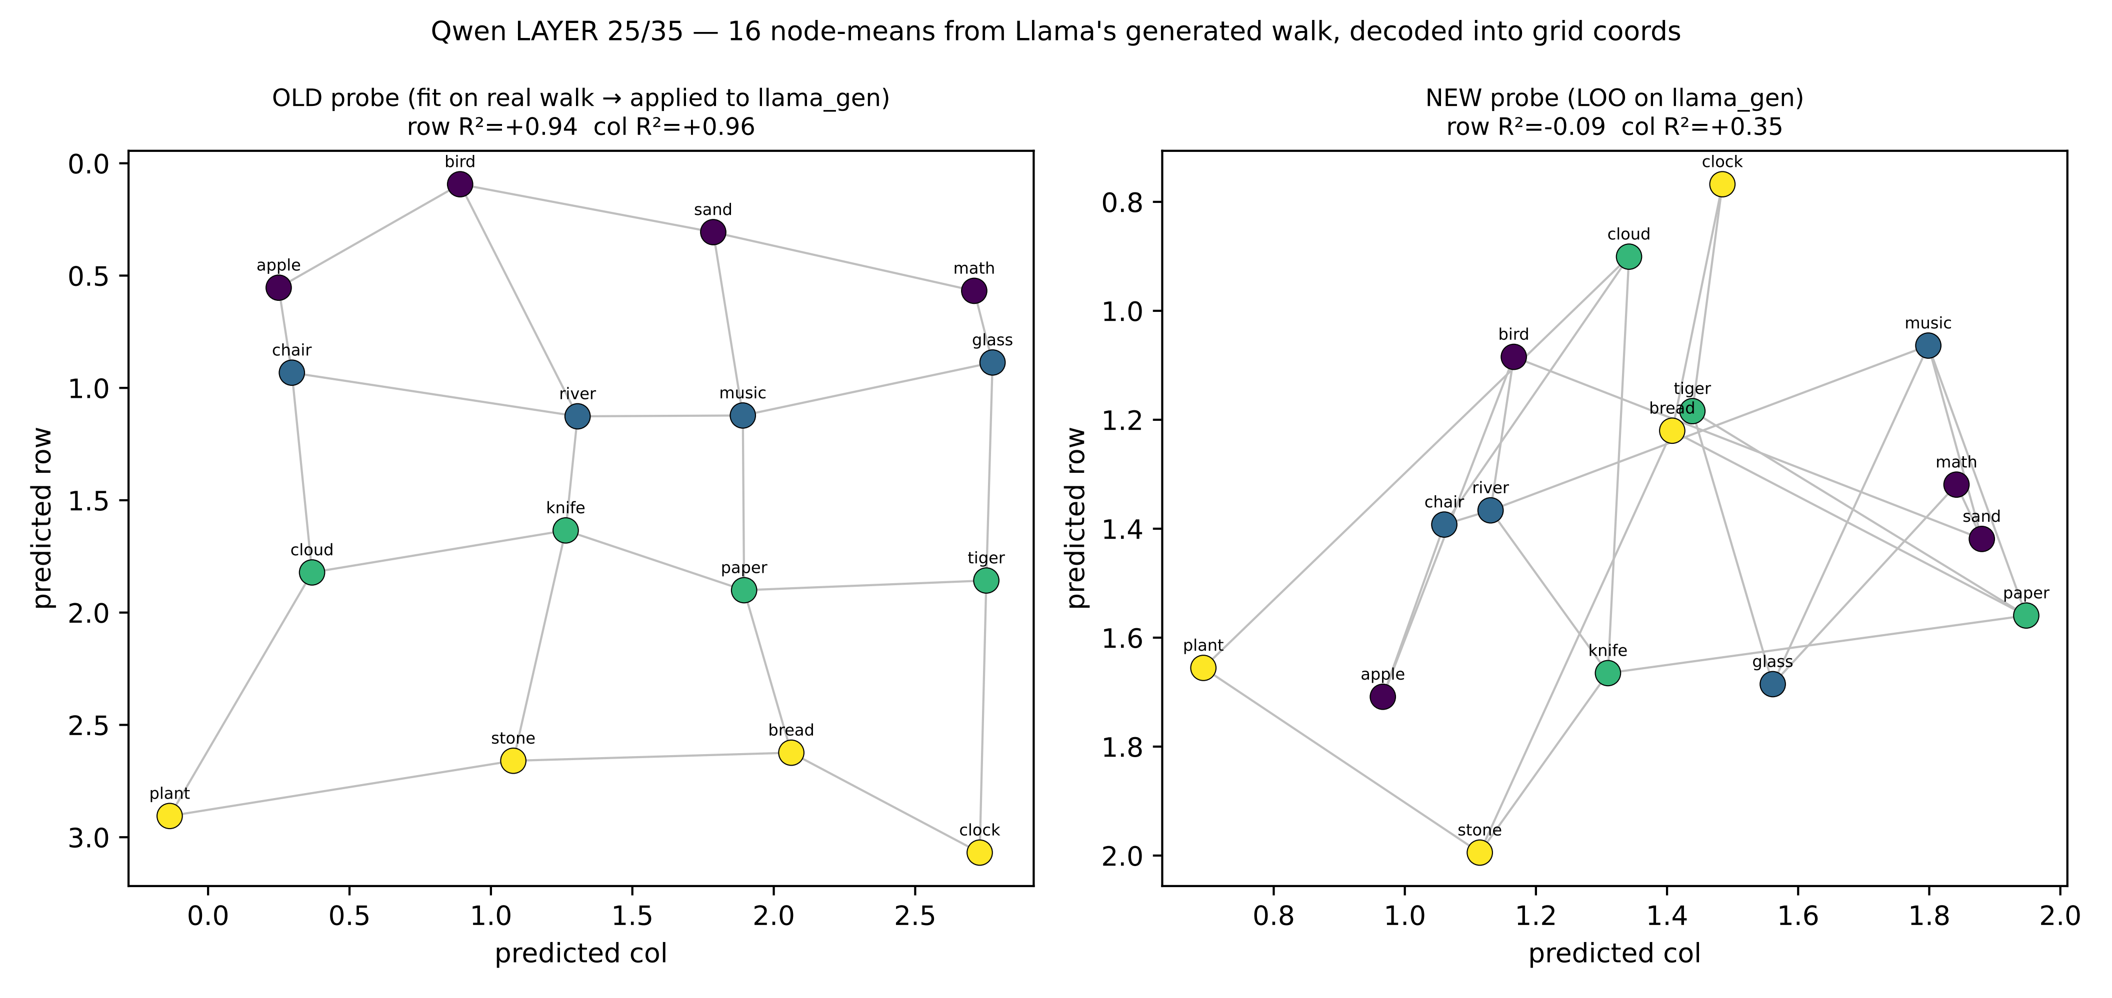
Qwen LAYER 25/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.94 col R²=+0.96
NEW probe (LOO on llama_gen)
row R²=-0.09 col R²=+0.35
predicted row
predicted col
bird
sand
apple
math
chair
glass
river
music
knife
cloud
paper
tiger
stone
bread
plant
clock

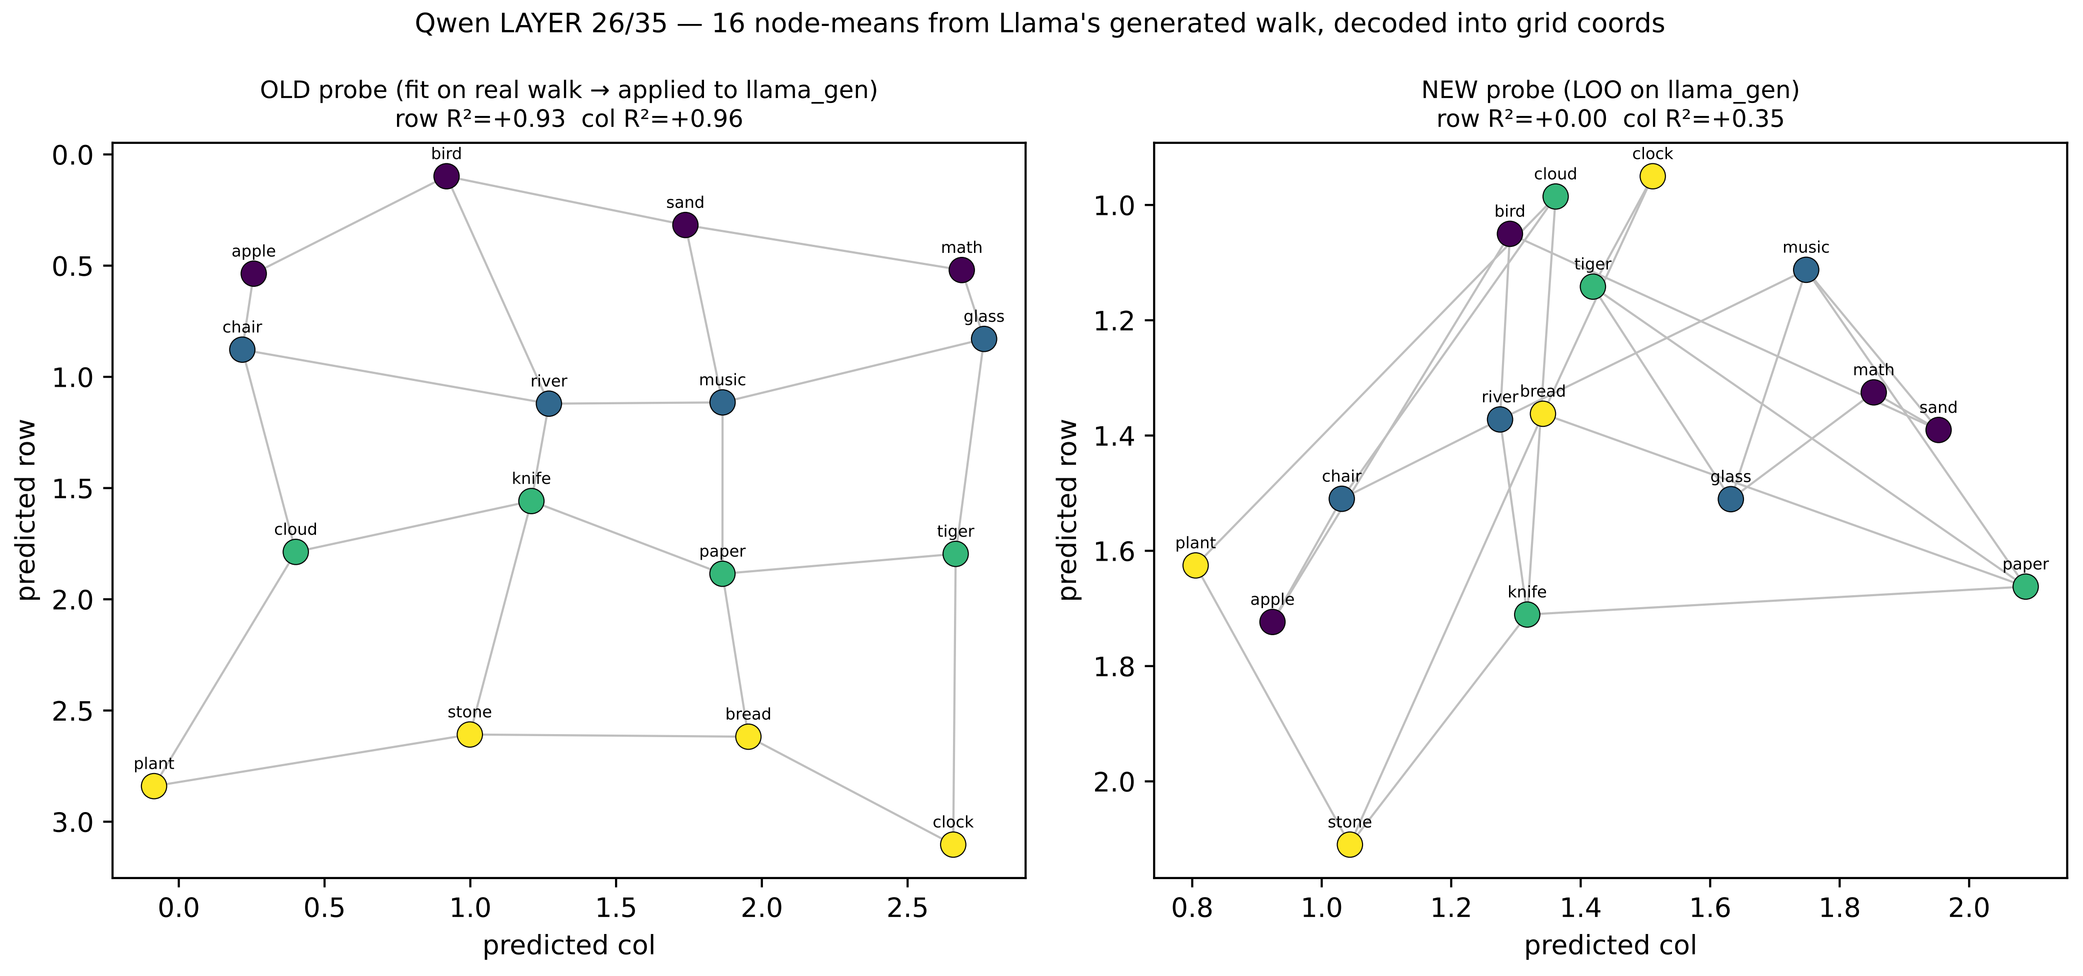

Qwen LAYER 26/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.93 col R²=+0.96
NEW probe (LOO on llama_gen)
row R²=+0.00 col R²=+0.35
predicted row
predicted col
bird
sand
apple
math
chair
glass
river
music
knife
cloud
tiger
paper
stone
bread
plant
clock
0.0
0.5
1.0
1.5
2.0
2.5
3.0
0.8
1.2
1.4
1.6
1.8

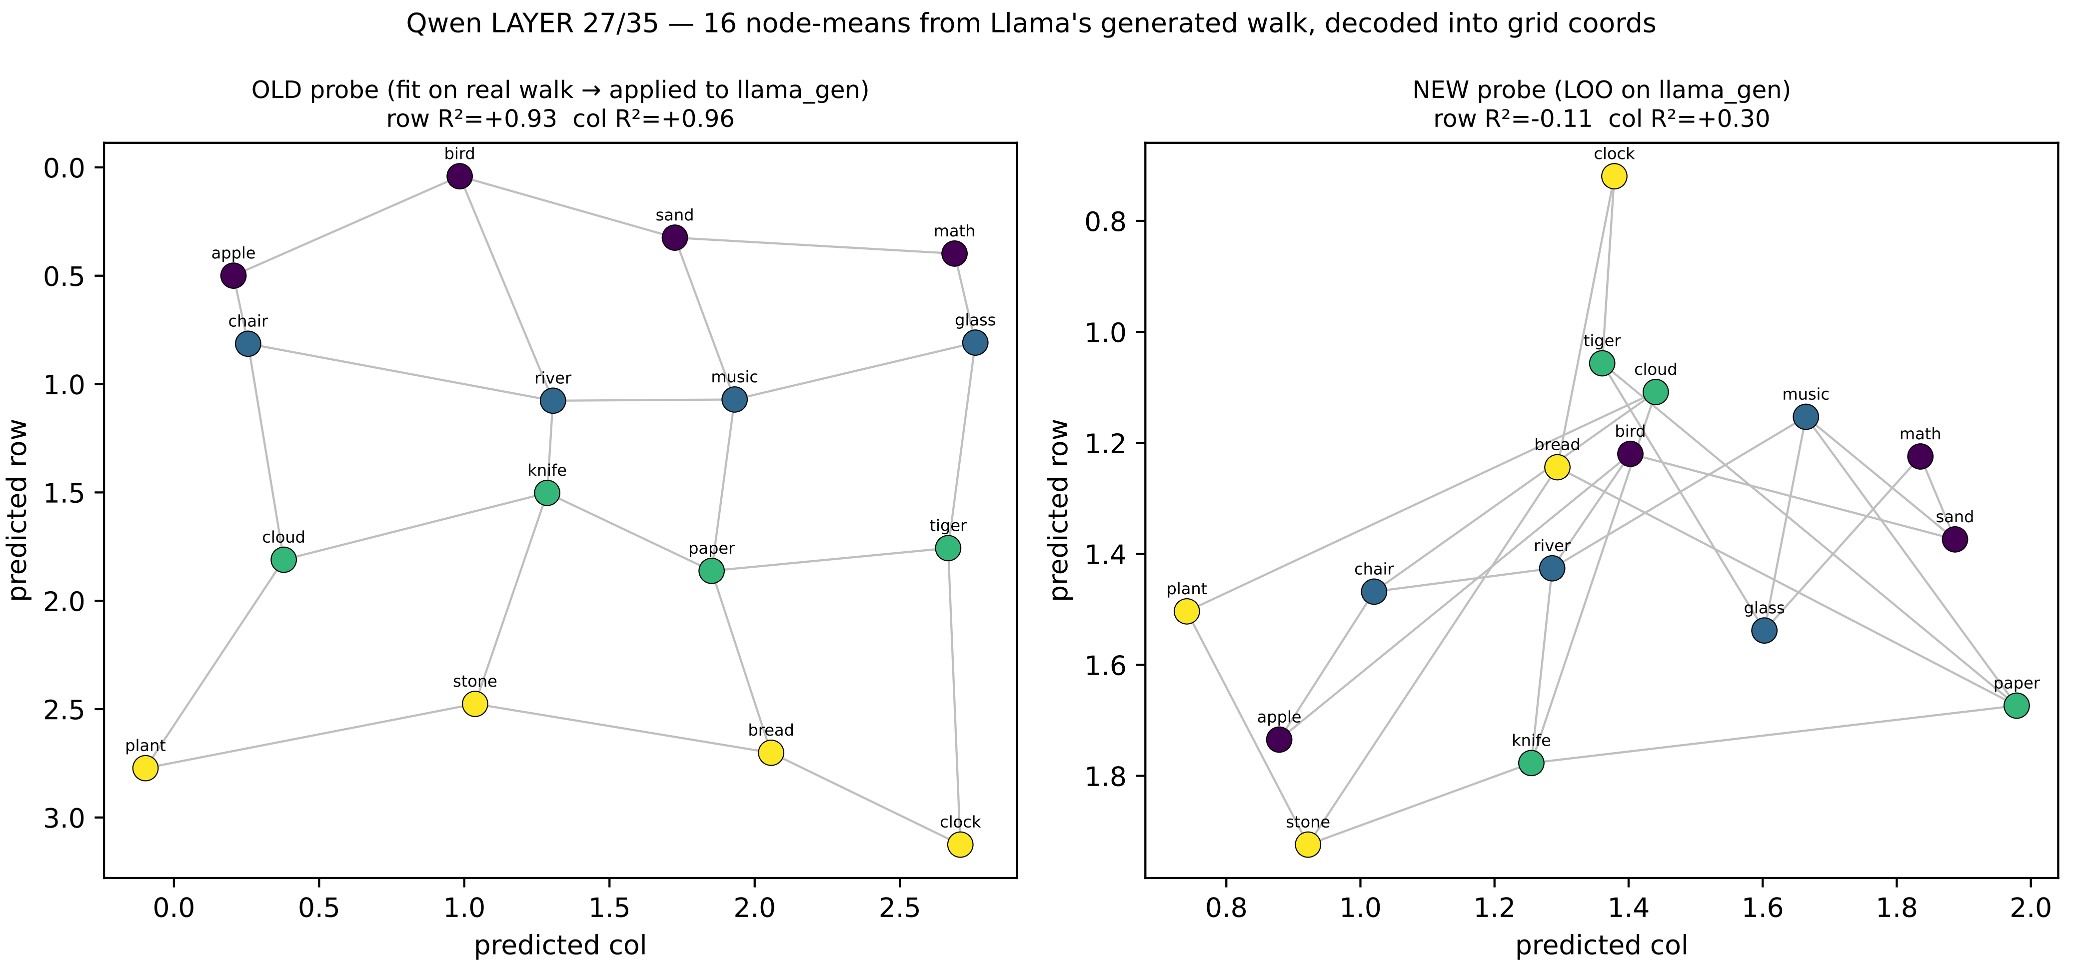

Qwen LAYER 27/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.93 col R²=+0.96
NEW probe (LOO on llama_gen)
row R²=-0.11 col R²=+0.30
predicted row
predicted col
bird
sand
math
apple
chair
glass
river
music
knife
cloud
paper
tiger
stone
bread
plant
clock

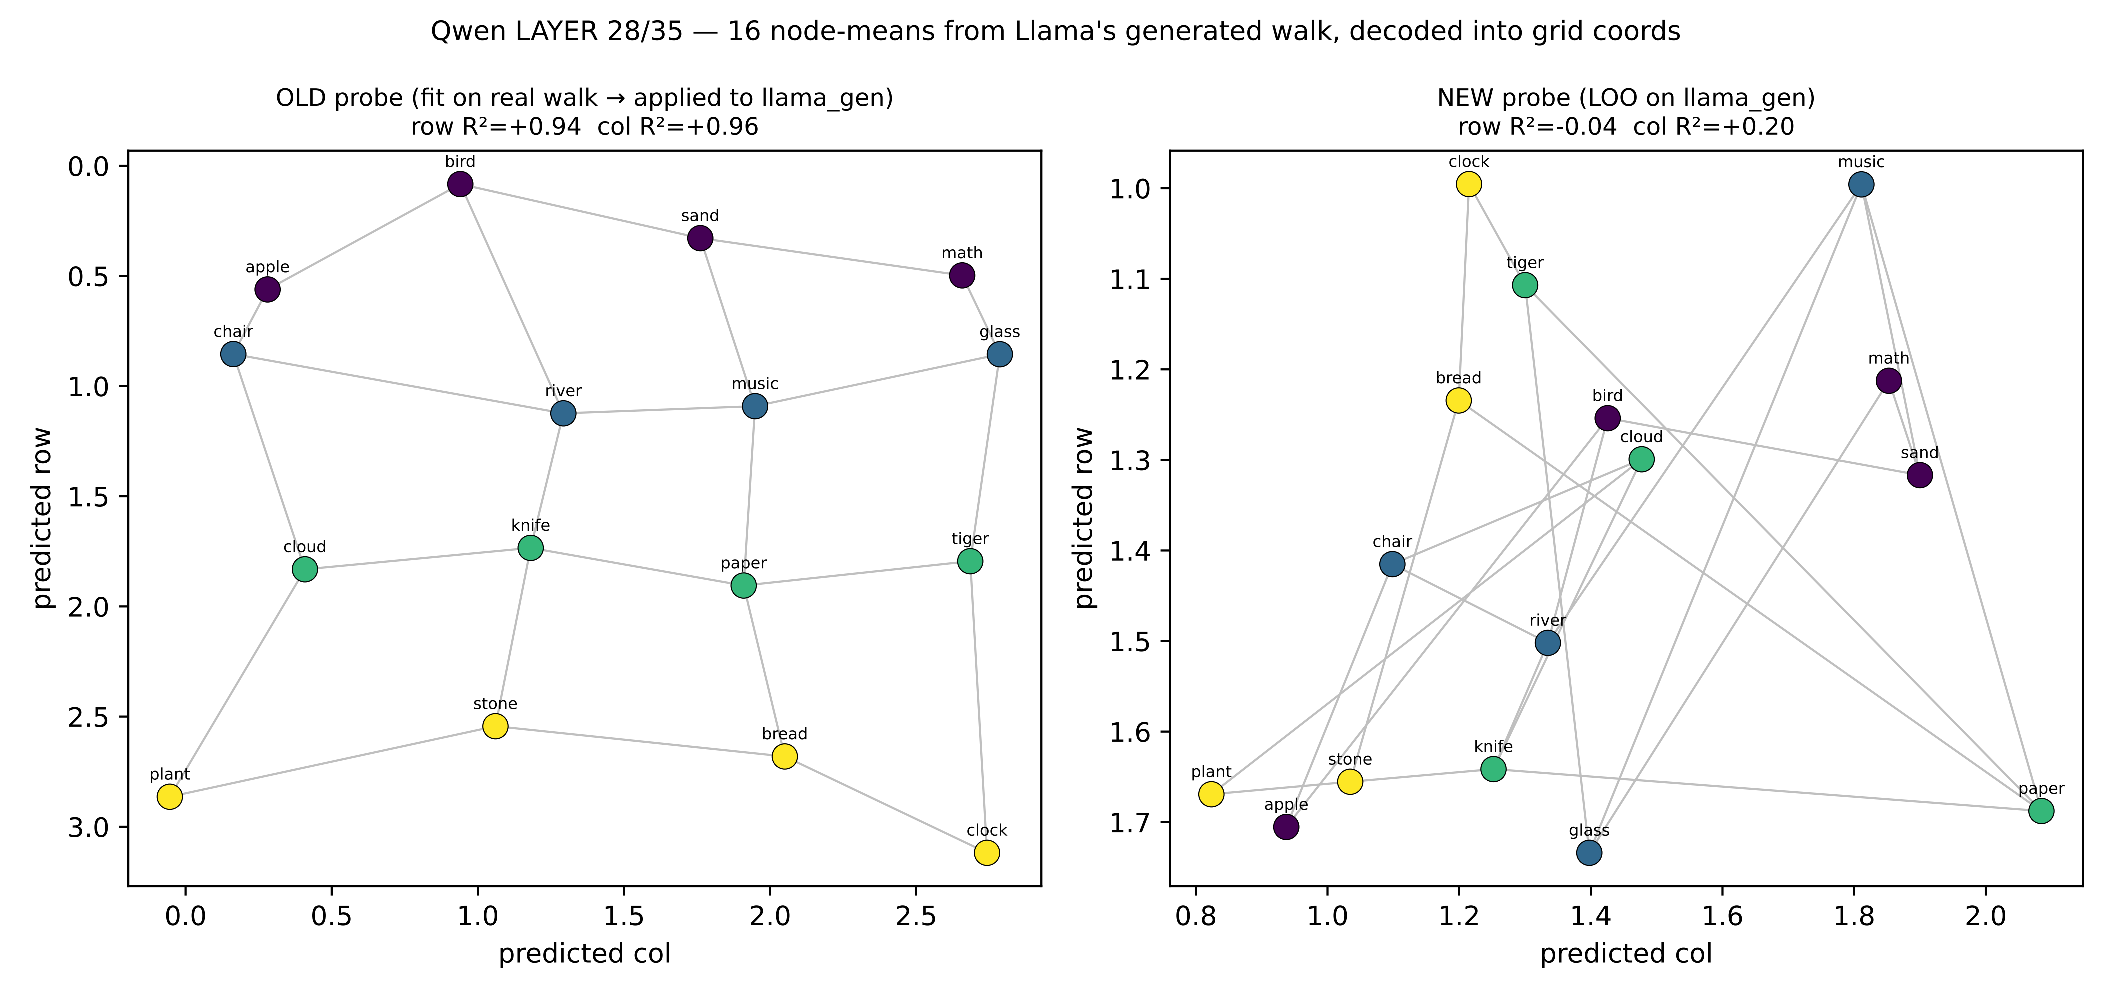
Qwen LAYER 28/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.94 col R²=+0.96
NEW probe (LOO on llama_gen)
row R²=-0.04 col R²=+0.20
predicted row
predicted col
bird
sand
math
apple
chair
glass
river
music
knife
cloud
tiger
paper
stone
bread
plant
clock

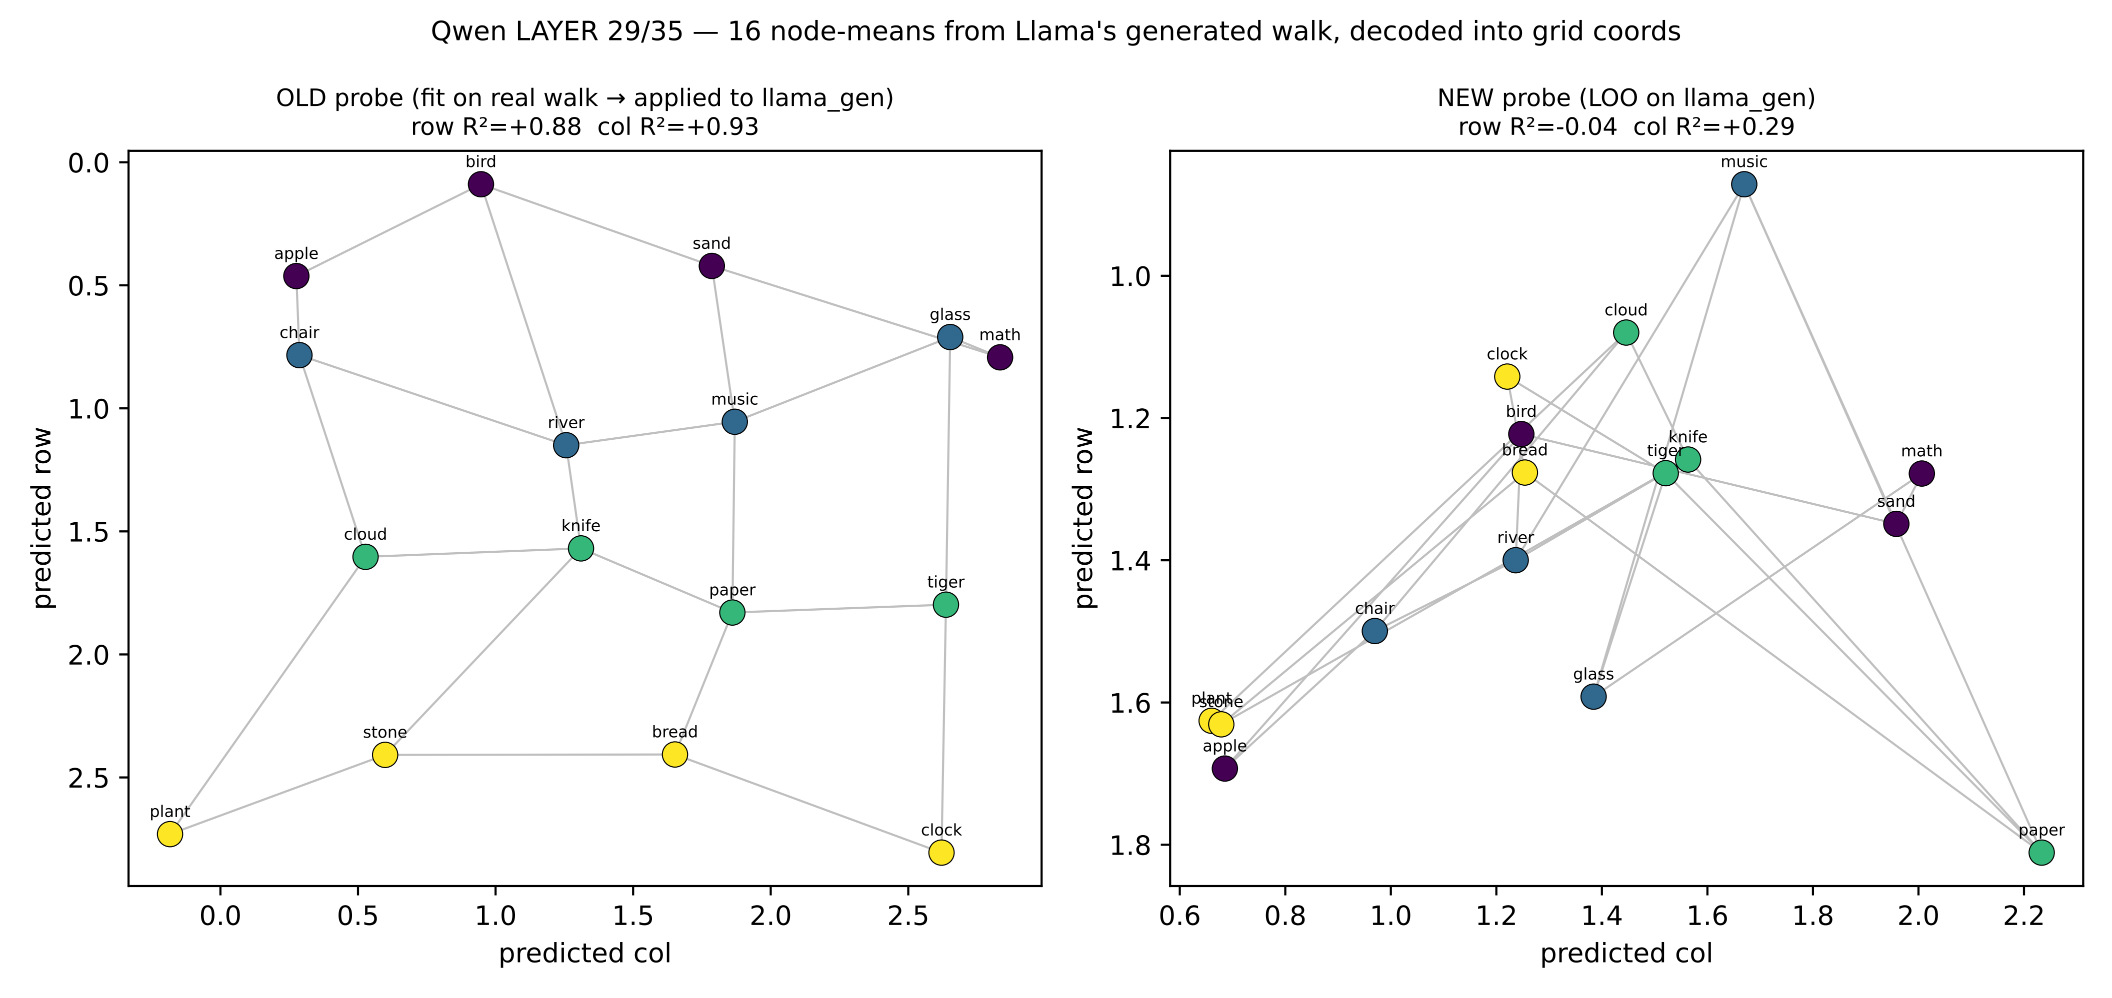

Qwen LAYER 29/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.88 col R²=+0.93
NEW probe (LOO on llama_gen)
row R²=-0.04 col R²=+0.29
predicted row
predicted col
bird
apple
sand
glass
math
chair
music
river
cloud
knife
paper
tiger
stone
bread
plant
clock

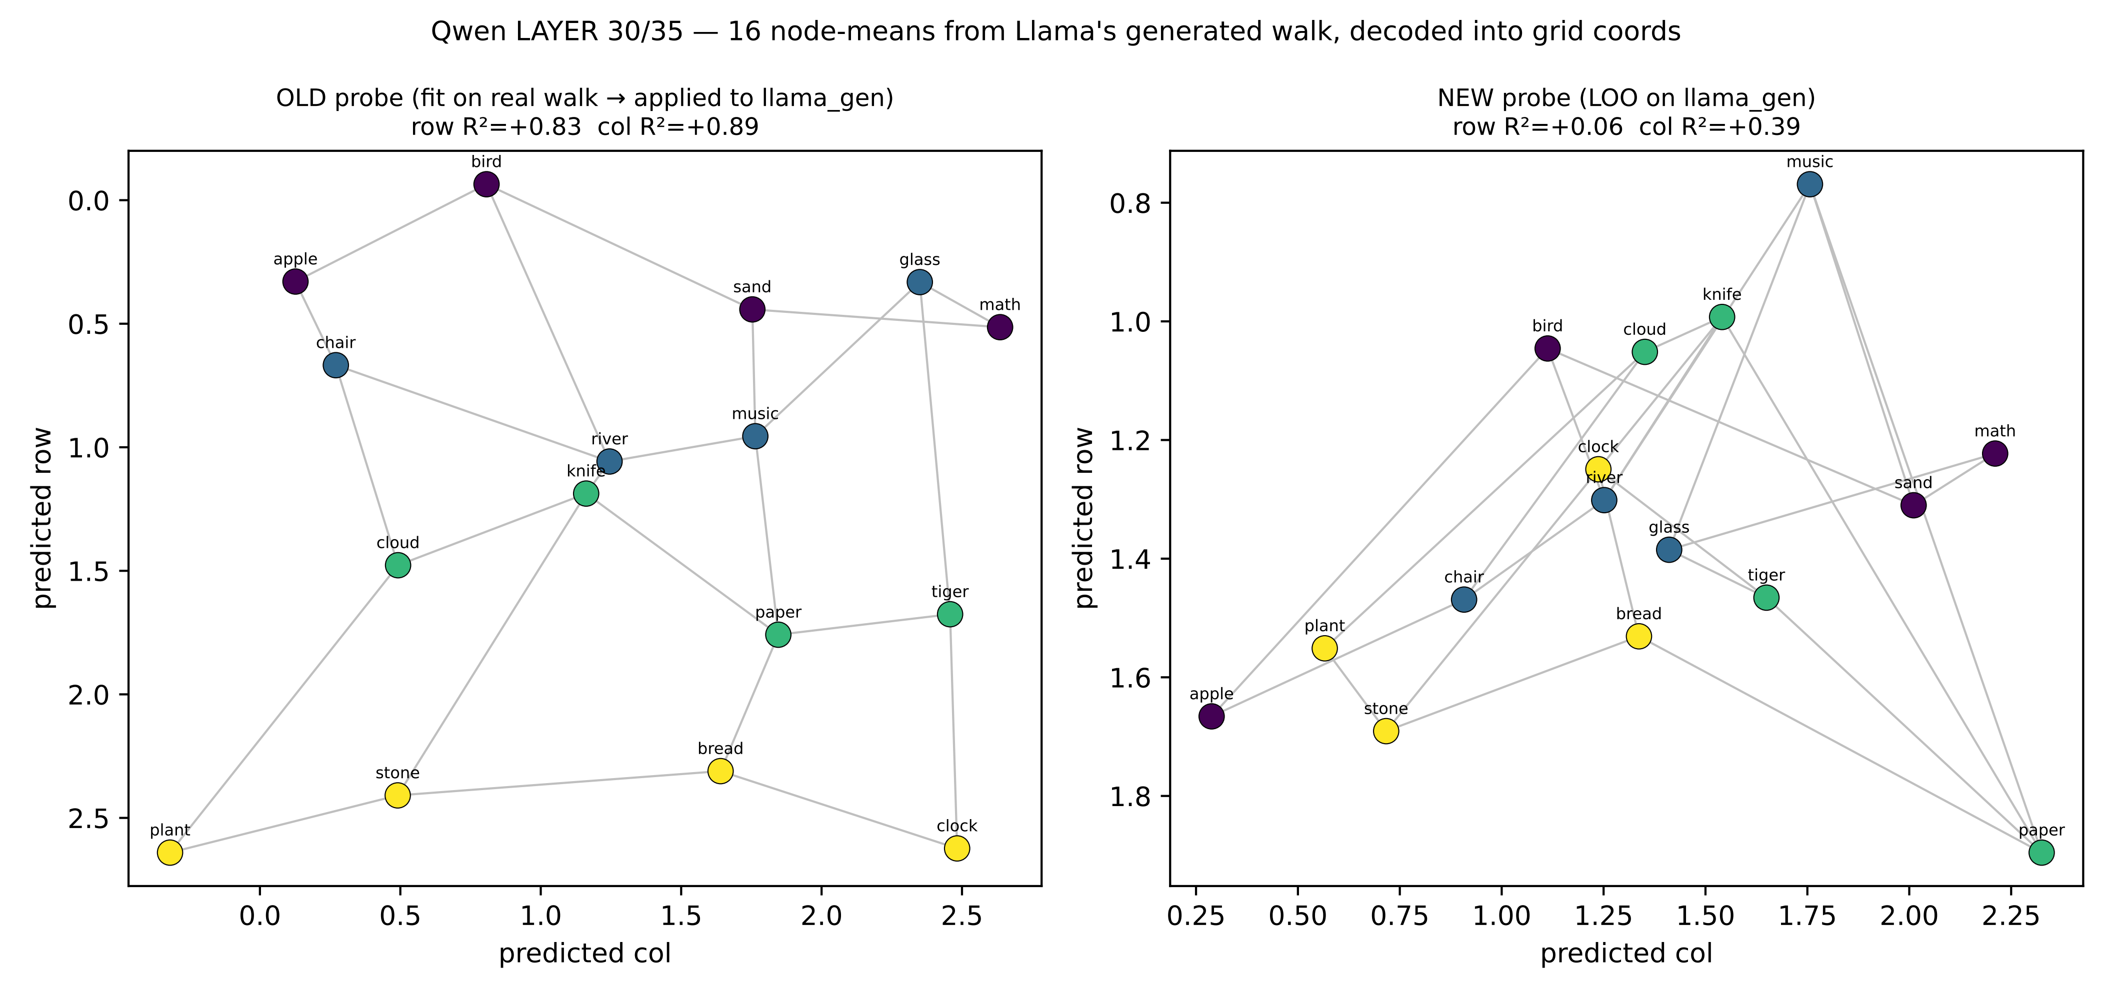
Qwen LAYER 30/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.83 col R²=+0.89
NEW probe (LOO on llama_gen)
row R²=+0.06 col R²=+0.39
predicted row
predicted col
bird
apple
sand
glass
math
chair
music
river
knife
cloud
paper
tiger
bread
stone
plant
clock

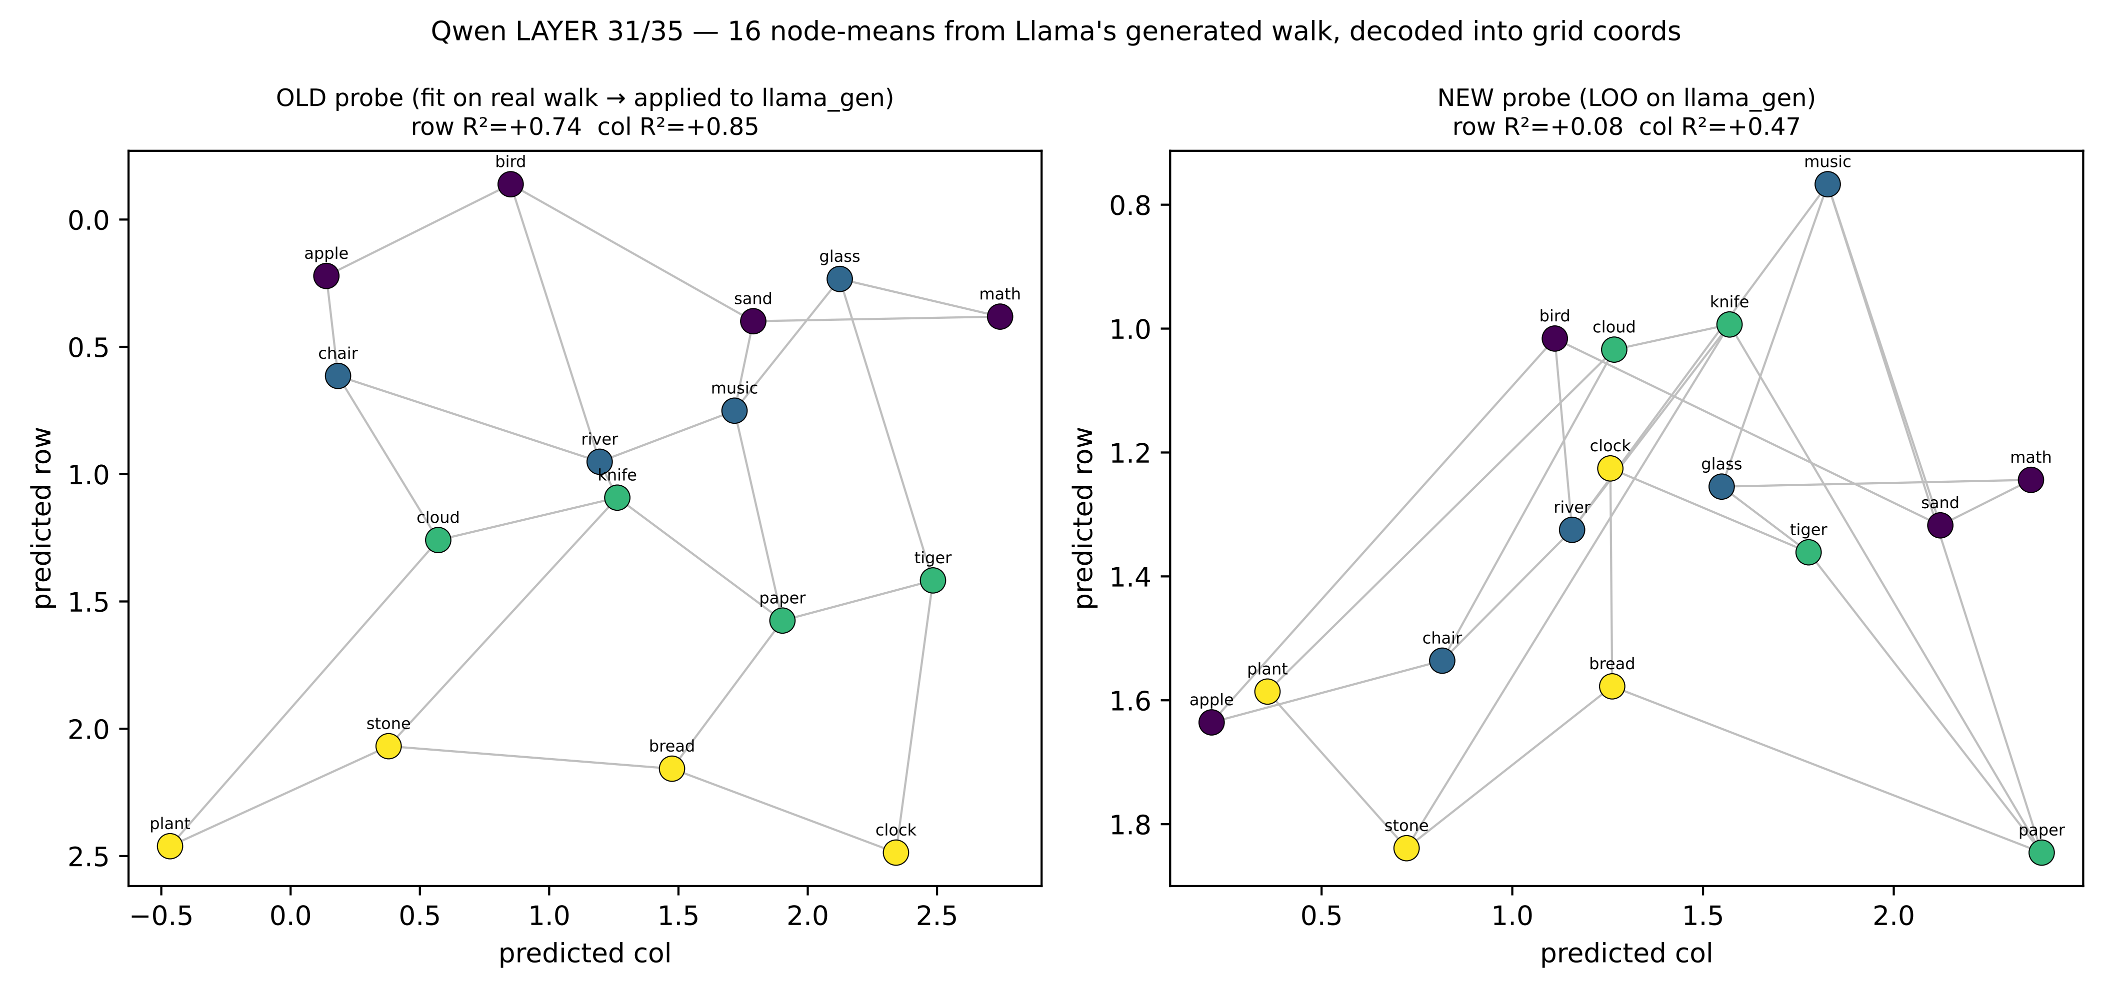

Qwen LAYER 31/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.74 col R²=+0.85
NEW probe (LOO on llama_gen)
row R²=+0.08 col R²=+0.47
predicted row
predicted col
bird
apple
glass
sand
math
chair
music
river
knife
cloud
tiger
paper
stone
bread
plant
clock

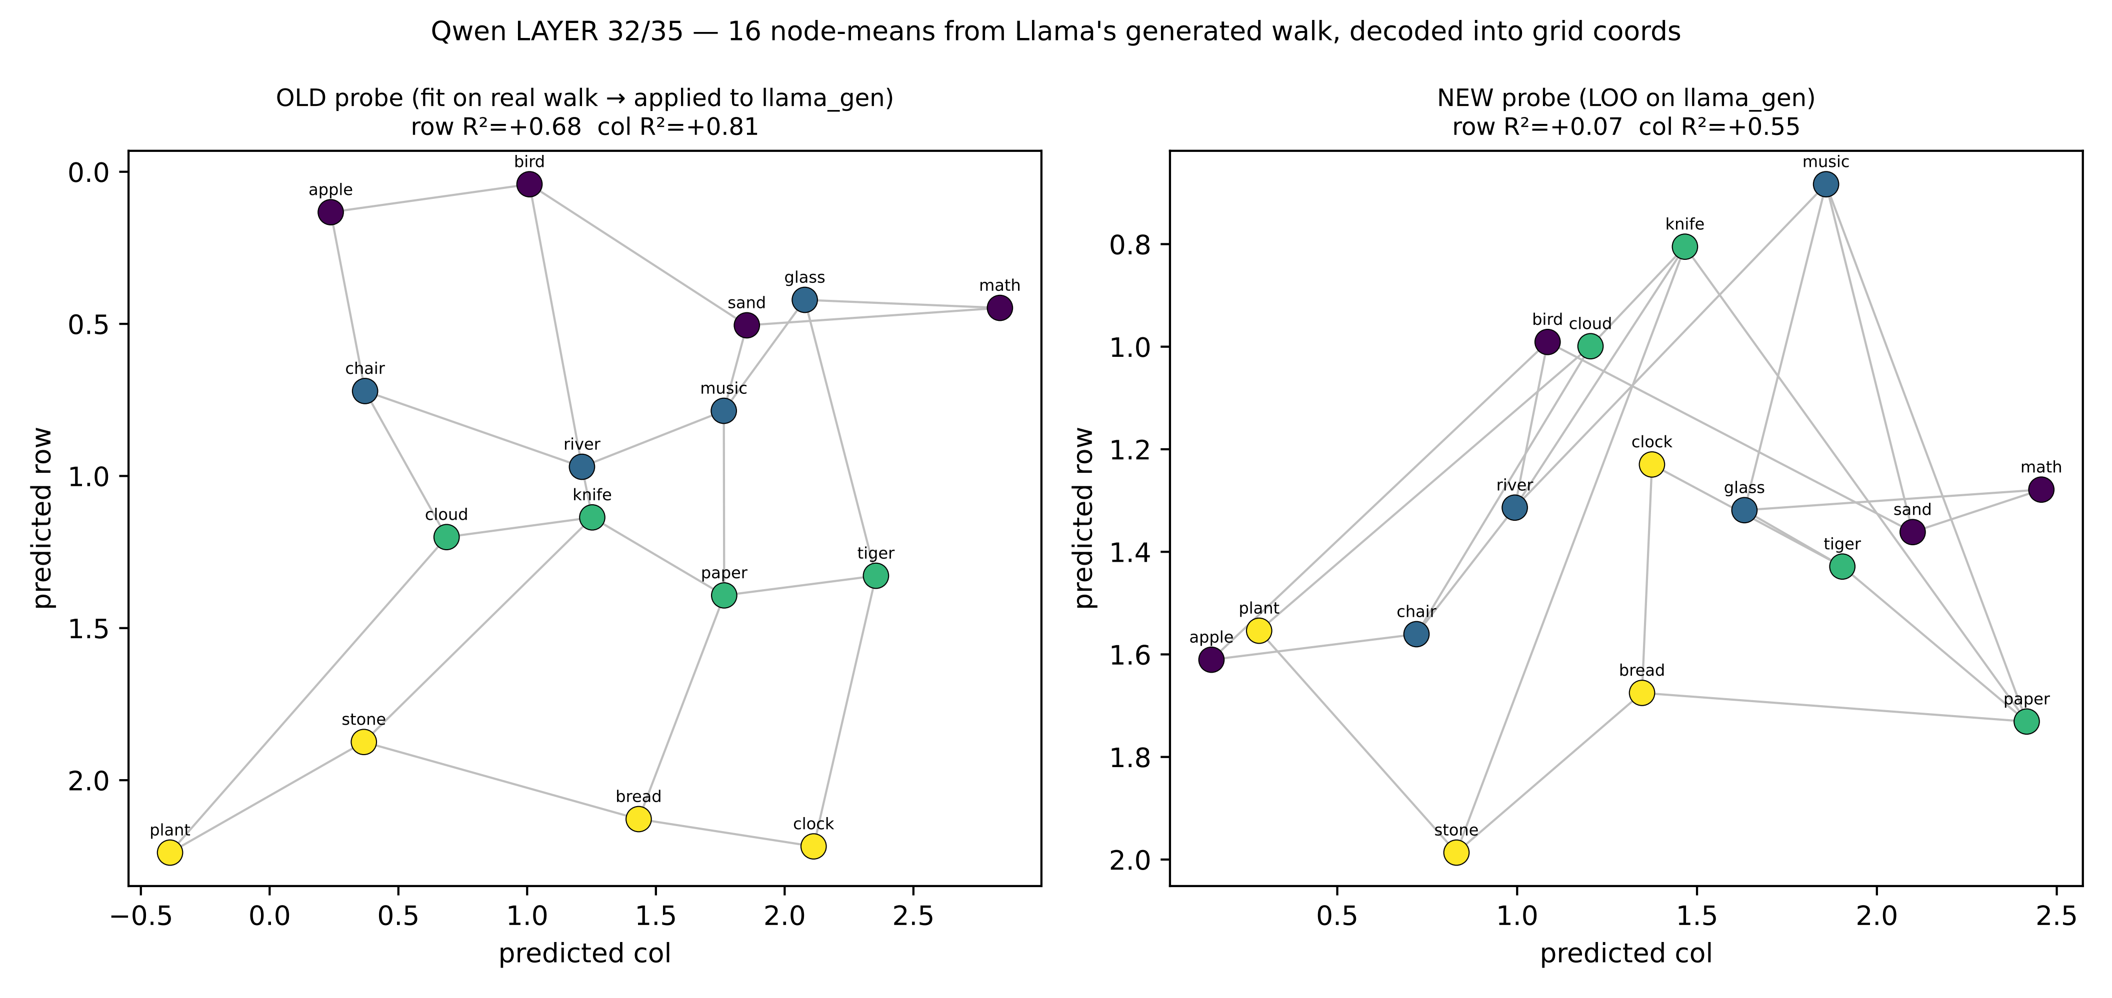
Qwen LAYER 32/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.68 col R²=+0.81
NEW probe (LOO on llama_gen)
row R²=+0.07 col R²=+0.55
predicted row
predicted col
apple
bird
glass
math
sand
chair
music
river
knife
cloud
tiger
paper
stone
bread
clock
plant

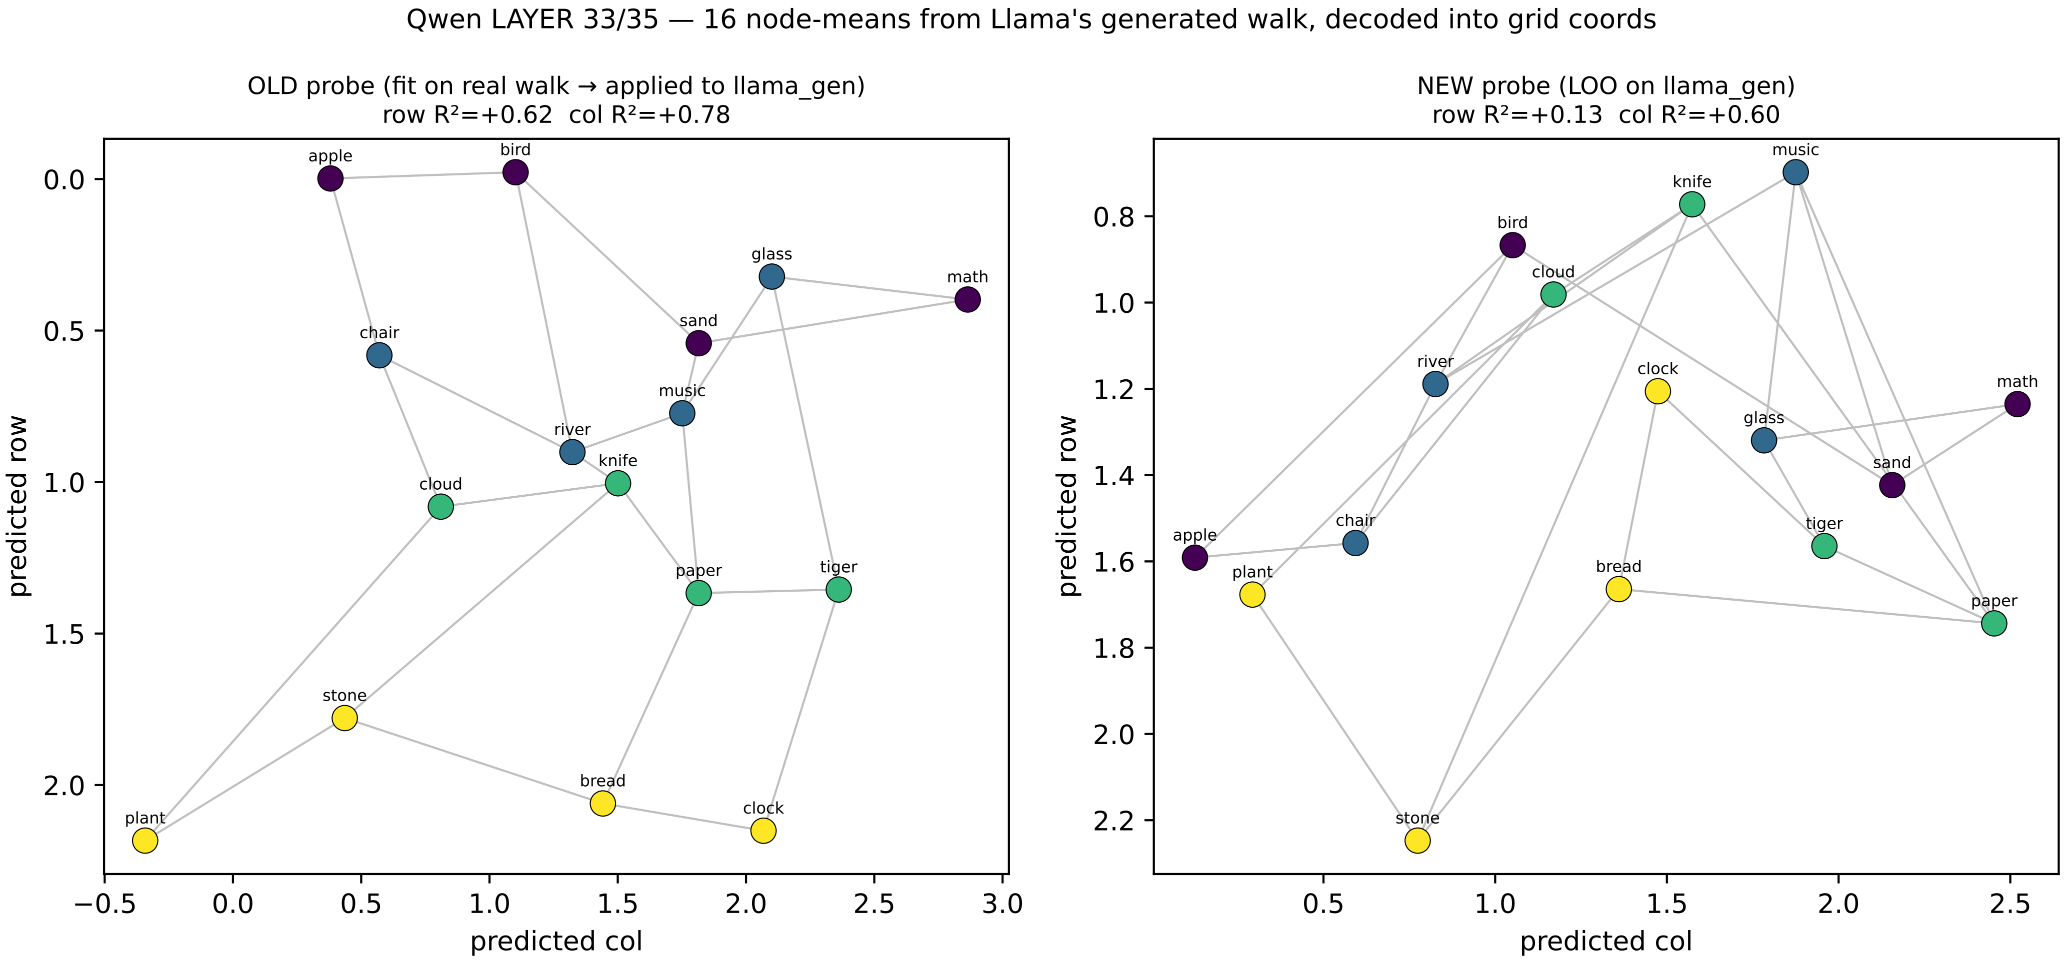

Qwen LAYER 33/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.62 col R²=+0.78
NEW probe (LOO on llama_gen)
row R²=+0.13 col R²=+0.60
predicted row
predicted col
apple
bird
glass
math
sand
chair
music
river
knife
cloud
paper
tiger
stone
bread
clock
plant

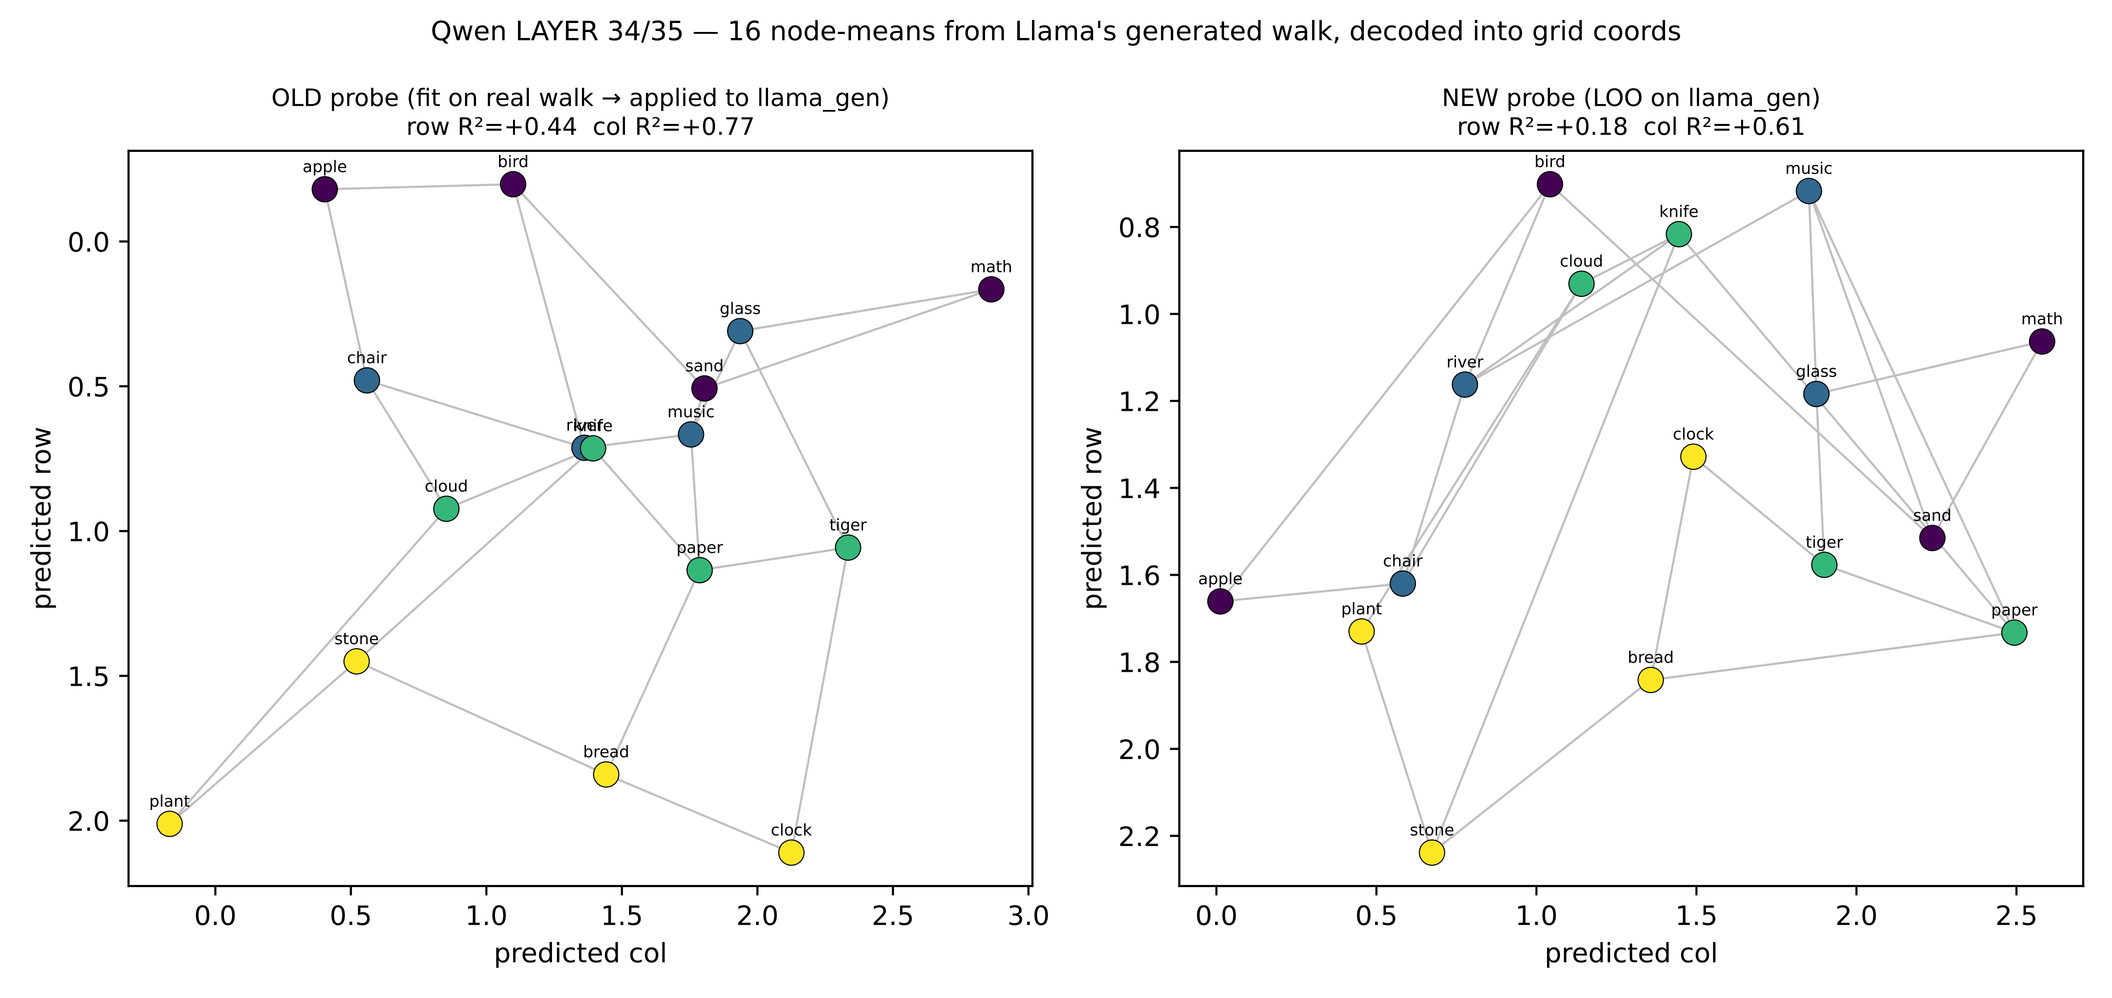
Qwen LAYER 34/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.44 col R²=+0.77
NEW probe (LOO on llama_gen)
row R²=+0.18 col R²=+0.61
predicted row
predicted col
apple
bird
math
glass
chair
sand
river
knife
music
cloud
tiger
paper
stone
bread
plant
clock
0.0
0.5
1.0
1.5
2.0
2.5
3.0
0.8
1.2
1.4
1.6
1.8
2.2

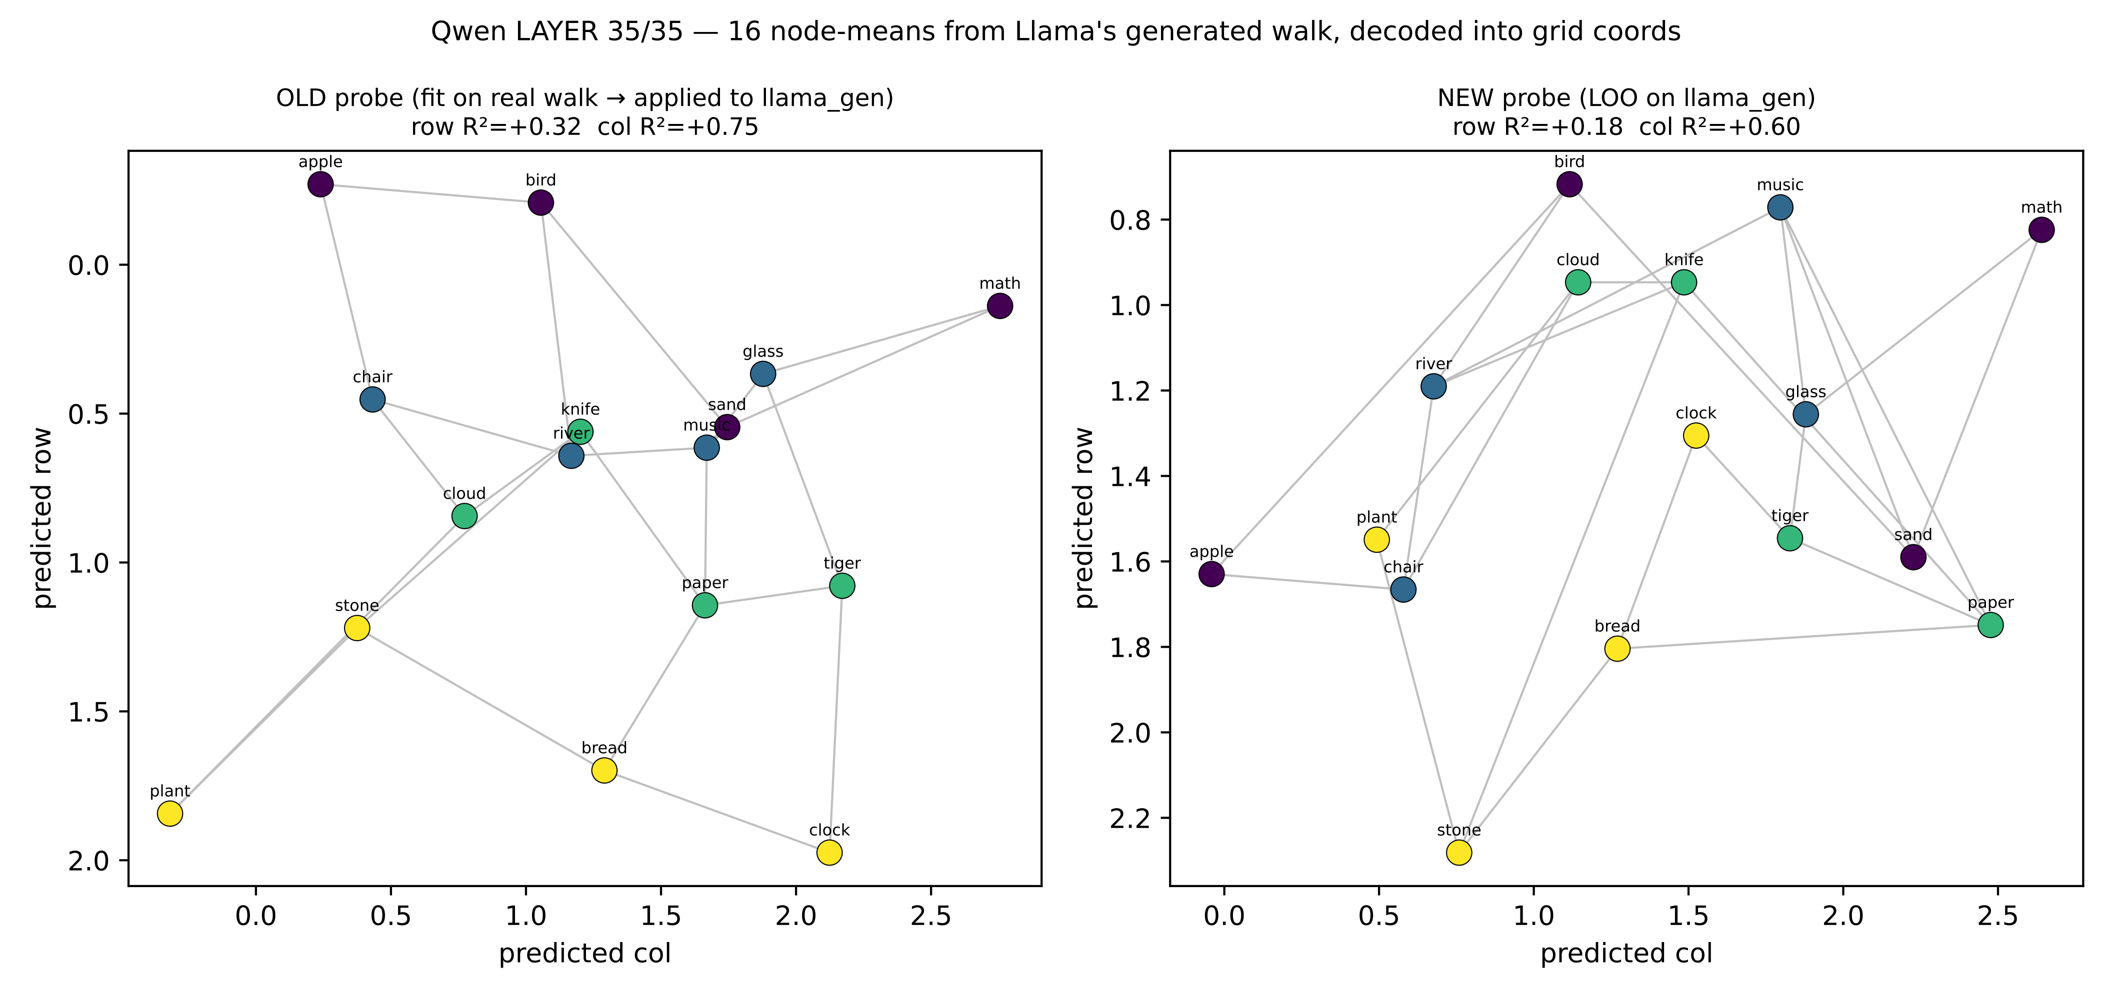

Qwen LAYER 35/35 — 16 node-means from Llama's generated walk, decoded into grid coords
OLD probe (fit on real walk → applied to llama_gen)
row R²=+0.32 col R²=+0.75
NEW probe (LOO on llama_gen)
row R²=+0.18 col R²=+0.60
predicted row
predicted col
apple
bird
math
glass
chair
knife
sand
river
music
cloud
paper
tiger
stone
bread
plant
clock