



Evidence shown across the 6 rounds ($\lambda=0.2$, $w=0.0$, $\tau=0.0$, seed=0; dotted = true θ / crowd c)


Company A




Company B




Company C ← MODEL CHOSE (rational)




Company D




Agreement with the rational (reward-maximizing) choice, by condition



What the optimal decision SHOULD weight on social vs. what the model DID




Counterfactual: causal effect of shifting the target's SOCIAL evidence on the choice




Chosen company's private vs. social estimate (points above the diagonal = social looked better than private)


$\lambda=0.2$




$\lambda=0.8$




Per-layer probe R^2 (where each signal is decodable)




Per-layer linear-probe R^2 (where each signal becomes decodable)



Patching the decision-position residual (donor→receiver): causal effect by layer



Steering the decision position by $\alpha \cdot (\text{more-social} - \text{less-social})$: dose-response



Does steering ever make the model PICK the target? (argmax = target)

