

Trust

In-context trust-learning under a reputation vs. accuracy conflict

Question: when a source's STATED REPUTATION conflicts with its OBSERVED ACCURACY,
does a frozen-weight LLM learn to trust the accurate source — and how fast vs. a rational learner?

Model: NousResearch/Meta-Llama-3.1-8B-Instruct (llama) | Source A = reputable but NOISIER, Source B = unestablished but ACCURATE

Accuracy gaps swept: $\sigma_A/\sigma_B = [1.5, 2.0, 3.0, 5.0]$ ($\sigma_B=12.0$) | Reputation strengths: ['faint', 'moderate', 'strong']

M=3 companies/round, T=16 rounds, ~20 games/condition | info-ladder: ['lean', 'self_history', 'summary', 'full']

10944 per-(round,company) observations | 2026-06-16

Method — per-round in-context loop

1. Build round-t prompt: persistent reputation header + RAW recap of past rounds (advisor record only) + this round's
2. Model outputs a numeric estimate per company (structured output) – the estimate IS the recommendation
3. Reveal the true value for each company -> accuracy evidence accumulates against the stale reputation
4. Extend history; proceed to round t+1 (one API call per round, T calls per game)

Reputation vs. accuracy CONFLICT: the header says A is trustworthy, but B's revealed errors are smaller every round.


Normative baseline (LEARNED precision): a Bayesian that does NOT know the noise levels, starts from a reputation-derived prior favoring A, and updates each source's precision from realized errors: $\pi_{\hat{}} = (\nu_0 + n) / (\nu_0 \cdot \sigma^2 + S)$.

Trust on B: $w = \pi_{\hat{}}_B / (\pi_{\hat{}}_A + \pi_{\hat{}}_B)$ -> migrates A->B as evidence accrues (oracle asymptote uses the TRUE s



Model trust readout: per round t, regress $\text{model_est} \sim b_A \cdot a + b_B \cdot b$ across companies x games; $\lambda_{\hat{}}_t = b_B / (b_A + b_B)$.

How one game runs and how the model's revealed trust ($\lambda_{\hat{}}$) and the rational trust (w) are computed.


Headline trust trajectory (gap=2.0, moderate reputation, lean info)



Both curves start anchored on the reputable-but-wrong source (low) and should climb toward the oracle. Model below Bayesian = reputation-stickiness; shaded = 95% bootstrap CI.




Stickiness = the model starts more anchored on A than warranted (left) and trails the rational A→B migration (positive over-trust, rounds-behind on the right).




THE key plot: as the accuracy gap widens (easier to tell sources apart) both flip sooner. Taller model bars than Bayes = the model needs more evidence to overcome reputation.

Reputation-STRENGTH effect (gap=2.0): how persistent authority slows the A→B flip




More persistently asserted reputation (faint→moderate→strong) should push the model's trajectory lower / later.

Asymptote — final model trust vs. the accuracy-justified weight




Gap between green (justified) and orange (model) at the end = residual deference to reputation that never resolves.



Localizes the bottleneck: if +summary (handed reliability) beats lean, it's an INFERENCE limit; if even full fails, it's DEFERENCE; if +self-history shifts it, anchoring.

Surface robustness (gap=2.0): label-swap / order / paraphrase



Trajectories should overlap — the A→B migration shouldn't depend on which letter is reputable, the listing order, or wording.

Summary — flip round (model vs Bayes), lag, final trust, verdict

condition	gap	rep	info	flip model	flip Bayes	lag	final $\bar{\lambda}$	oracle	verdict
gap1.5_moderate_le	1.5	moderate	lean	—	10	—	0.45	0.69	no flip in T
gap2.0_moderate_le	2.0	moderate	lean	14	5	+9	0.48	0.80	sticky/slow
gap3.0_moderate_le	3.0	moderate	lean	2	3	-1	0.45	0.90	learns→B
gap5.0_moderate_le	5.0	moderate	lean	2	2	+0	0.47	0.96	learns→B
gap2.0_faint_lean	2.0	faint	lean	2	5	-3	0.47	0.80	learns→B
gap2.0_strong_lean	2.0	strong	lean	—	5	—	0.39	0.80	no flip in T
gap2.0_moderate_self_h	2.0	moderate	self_history	14	5	+9	0.48	0.80	sticky/slow
gap2.0_moderate_summ	2.0	moderate	summary	14	5	+9	0.45	0.80	sticky/slow
gap2.0_moderate_full	2.0	moderate	full	—	5	—	0.46	0.80	no flip in T
gap2.0_robust_swapla	2.0	moderate	lean	1	5	-4	0.51	0.80	learns→B
gap2.0_robust_order	2.0	moderate	lean	11	5	+6	0.51	0.80	sticky/slow
gap2.0_robust_paraph	2.0	moderate	lean	14	5	+9	0.49	0.80	sticky/slow
gap2.0_moderate_no_ad	2.0	moderate	no_advisor	-	-	-	-	-	RMSE=150 (necessity)

lag = model flip round minus Bayesian flip round (positive = behind). final $\bar{\lambda}$ vs oracle shows residual deference.

Takeaways

- Does it learn to trust the accurate source? YES — at gap=2.0 the model's revealed trust crosses 0.5 at round 14 (Bayes: 5).
- Reputation-stickiness: starts anchored at $\hat{\lambda}_1=0.45$ (prior favors A) and over-trusts A by +0.11 on average vs. the rational curve; flip lag = 9 rounds.
- Where the transition sits vs. Bayes across gaps (model vs Bayes flip round): gap1.5:None/10, gap2.0:14/5, gap3.0:2/3, gap5.0:2/2.
- Asymptote: final $\hat{\lambda}=0.48$ vs accuracy-justified 0.80 — residual deference remains.
- Inference vs. deference: compare lean vs. +summary on S8 — if handing the model running accuracy helps, the bottleneck is inference; if not, it's deference.
- Sanity: no-advisor RMSE=150 (advisors must add info); comprehension correct in 25% of probes.
- Caveats: scoped run (~20 games/condition → wider CIs than the 1000-game target); thinking disabled; single model; synthetic Gaussian world.